

Lecture 7

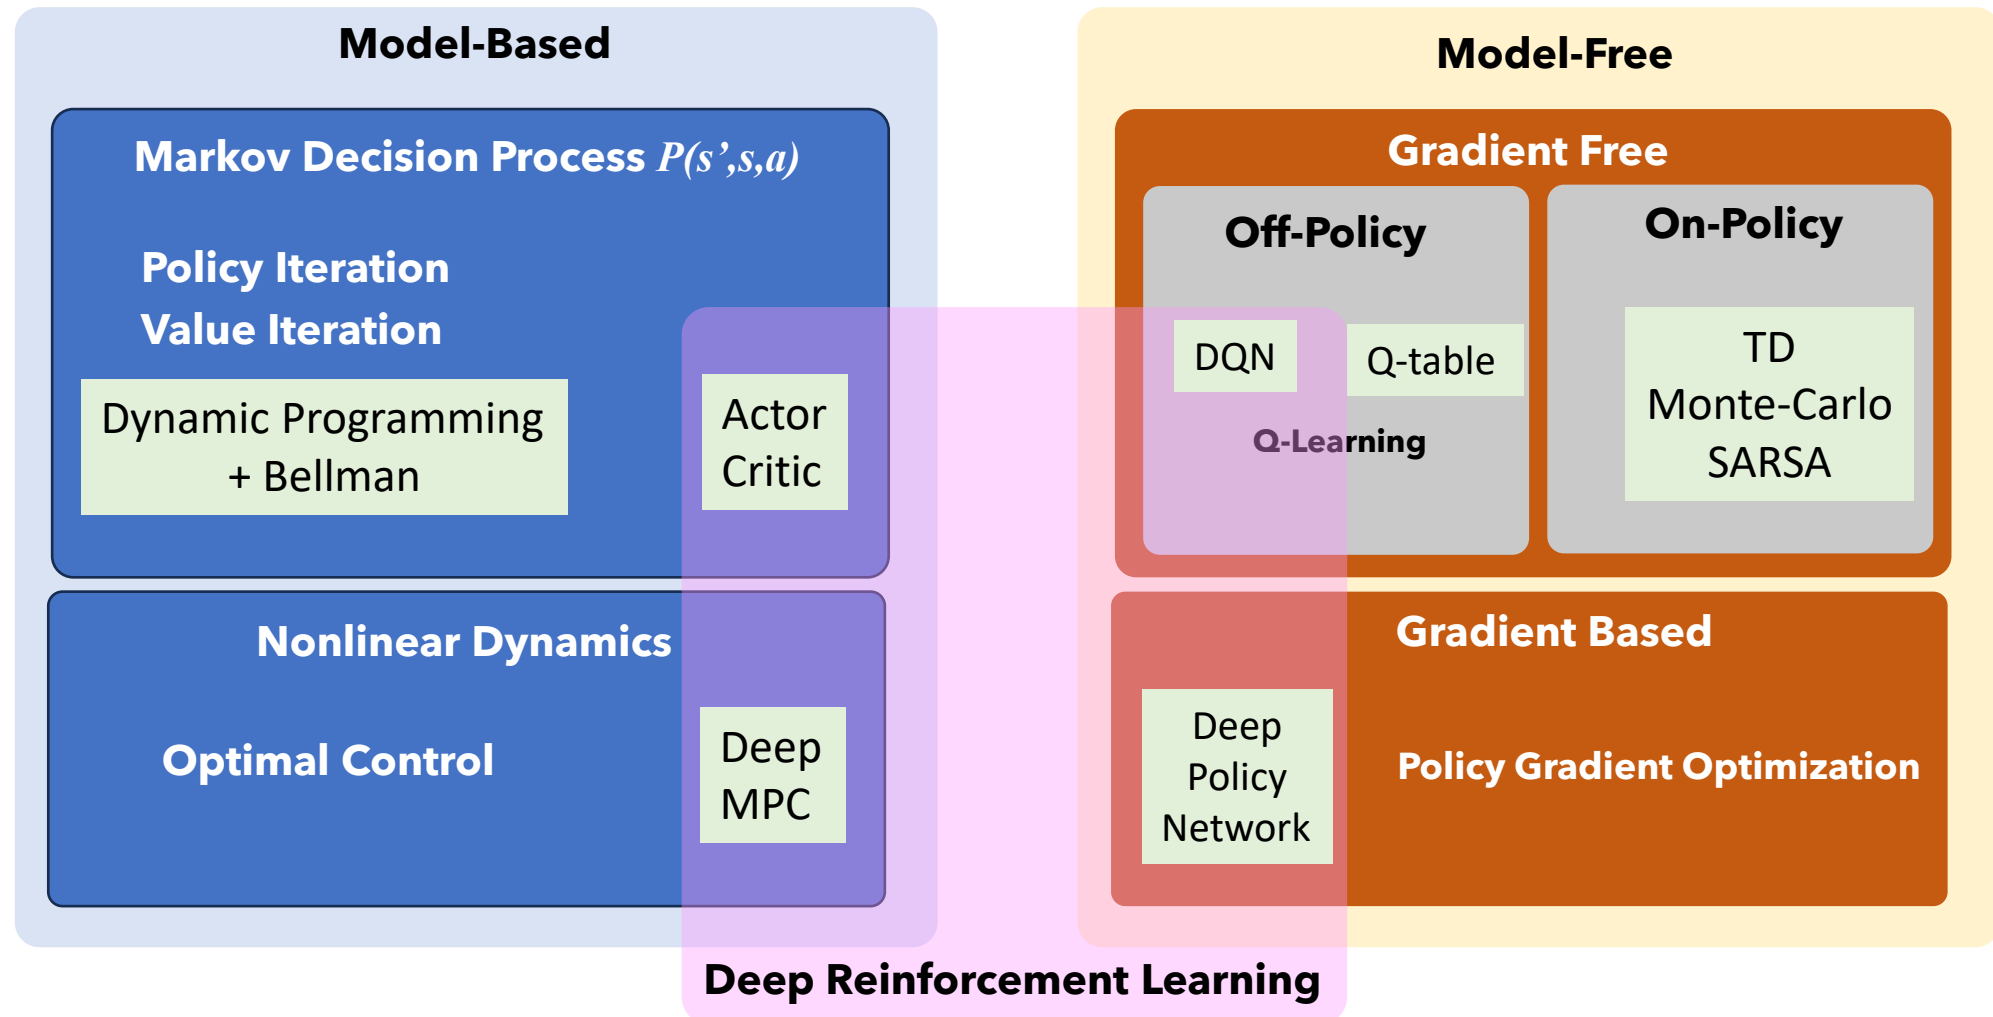
SARSA / Q-Learning

jmanero@faculty.ie.edu

MBD-EN2024ELECTIVOS-MBDMCSBT_37E89_467614

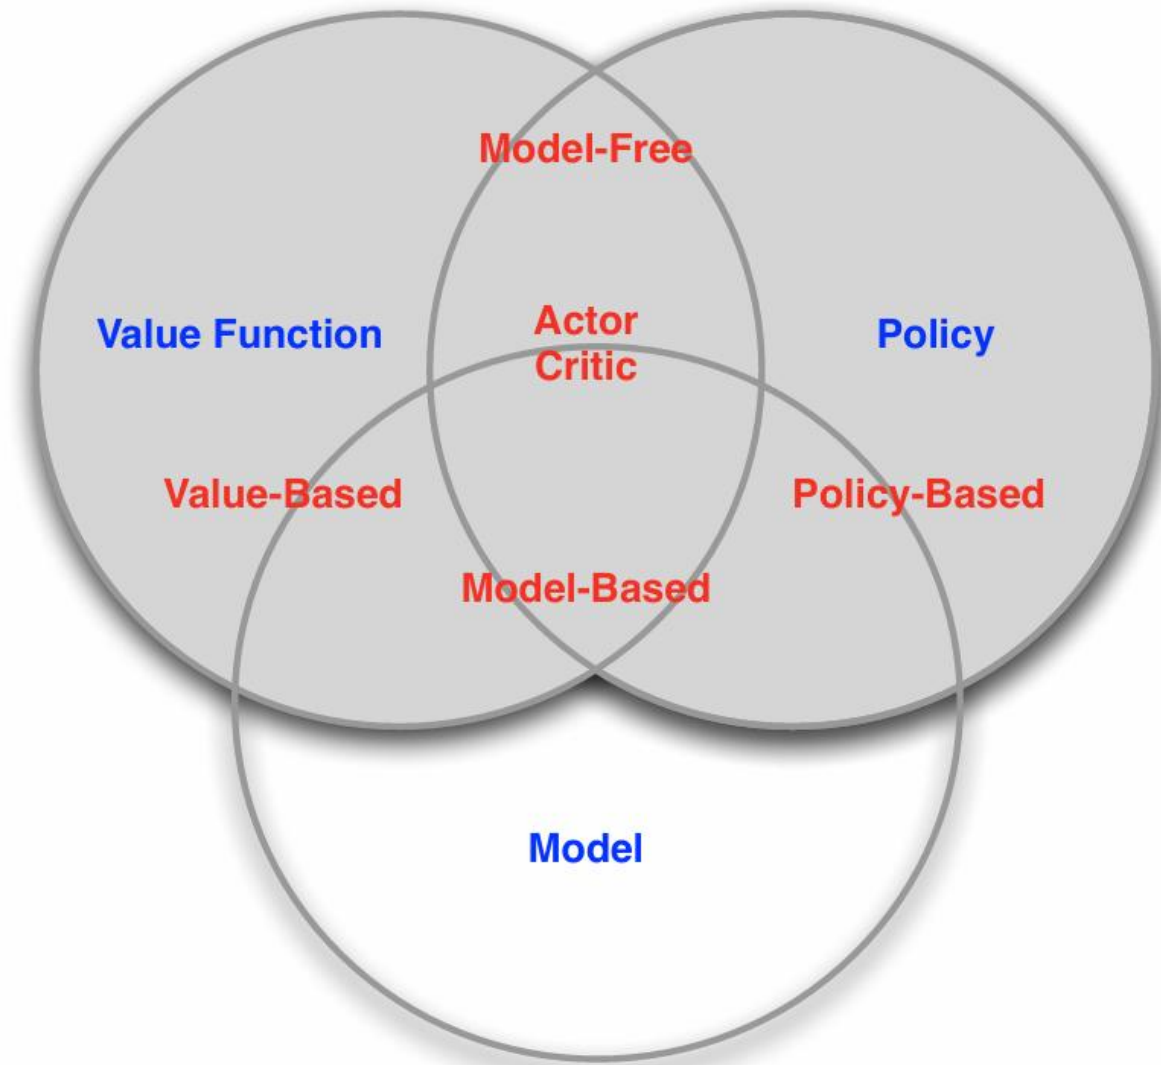
Where we are

Classification of RL Methods



Where we are

Taxonomy of RL Agents



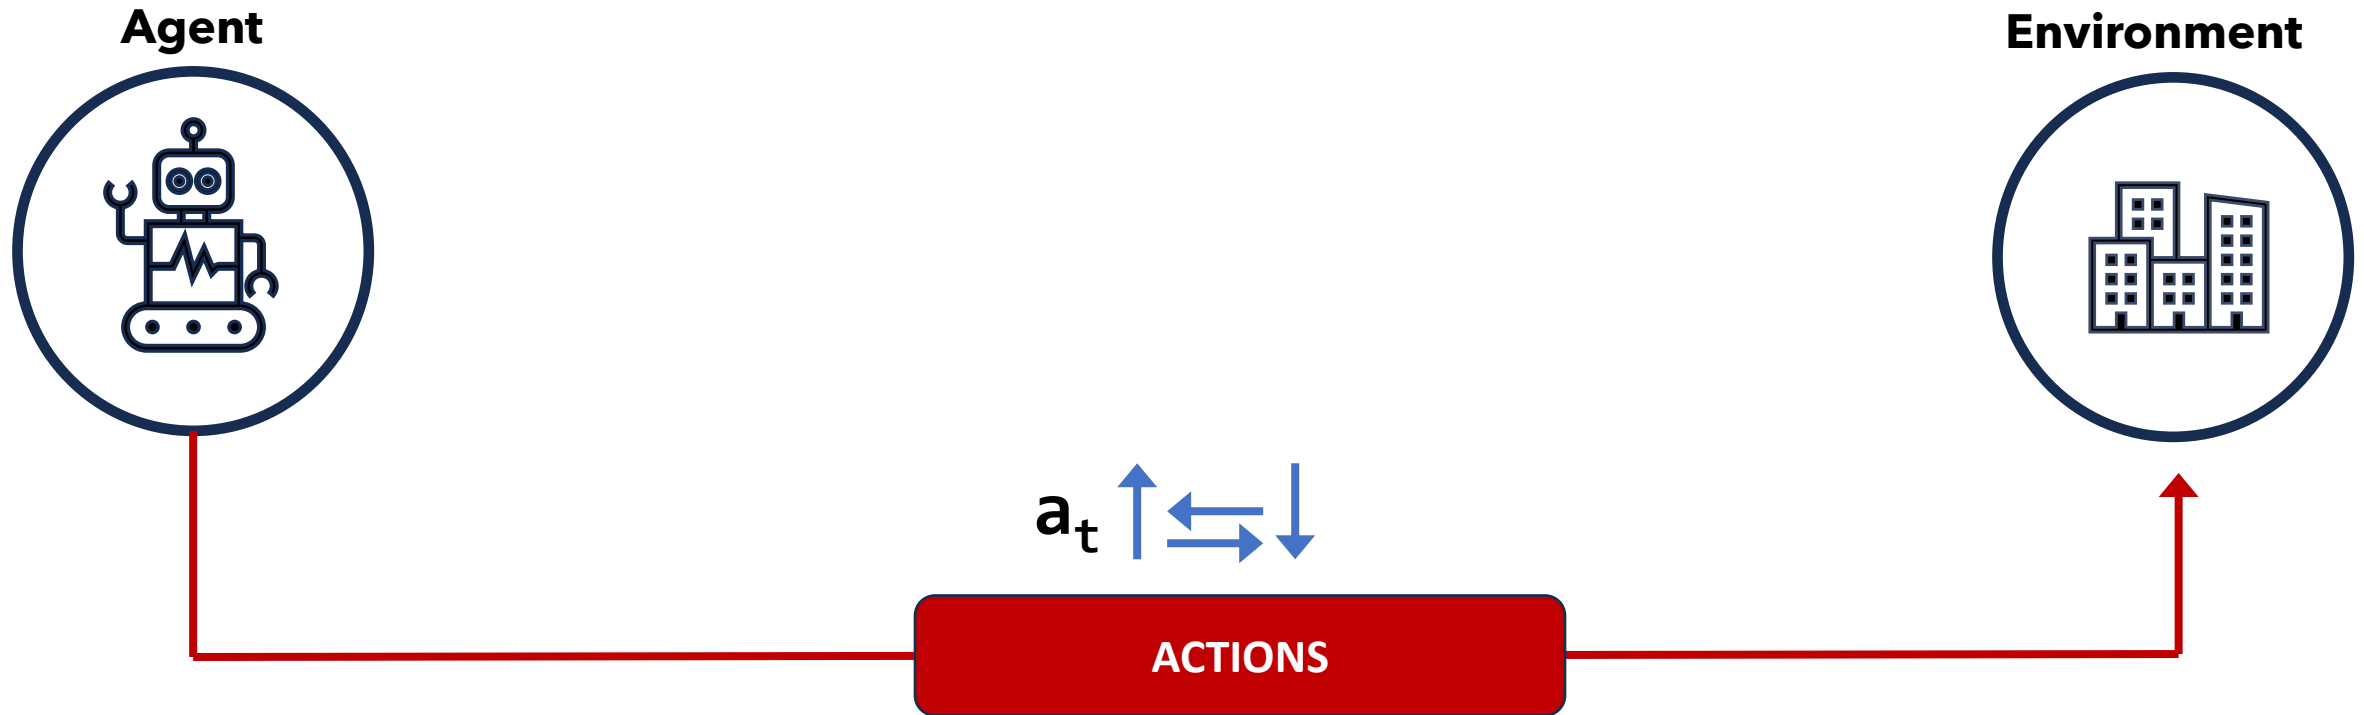
Taxonomy of RL Agents

(*) From David Silver RL Course UCL

- **Review Concepts**
- **SARSA**
- **Q-learning**

Key Concepts

Key Concepts: Actions



ACTION DEFINITION

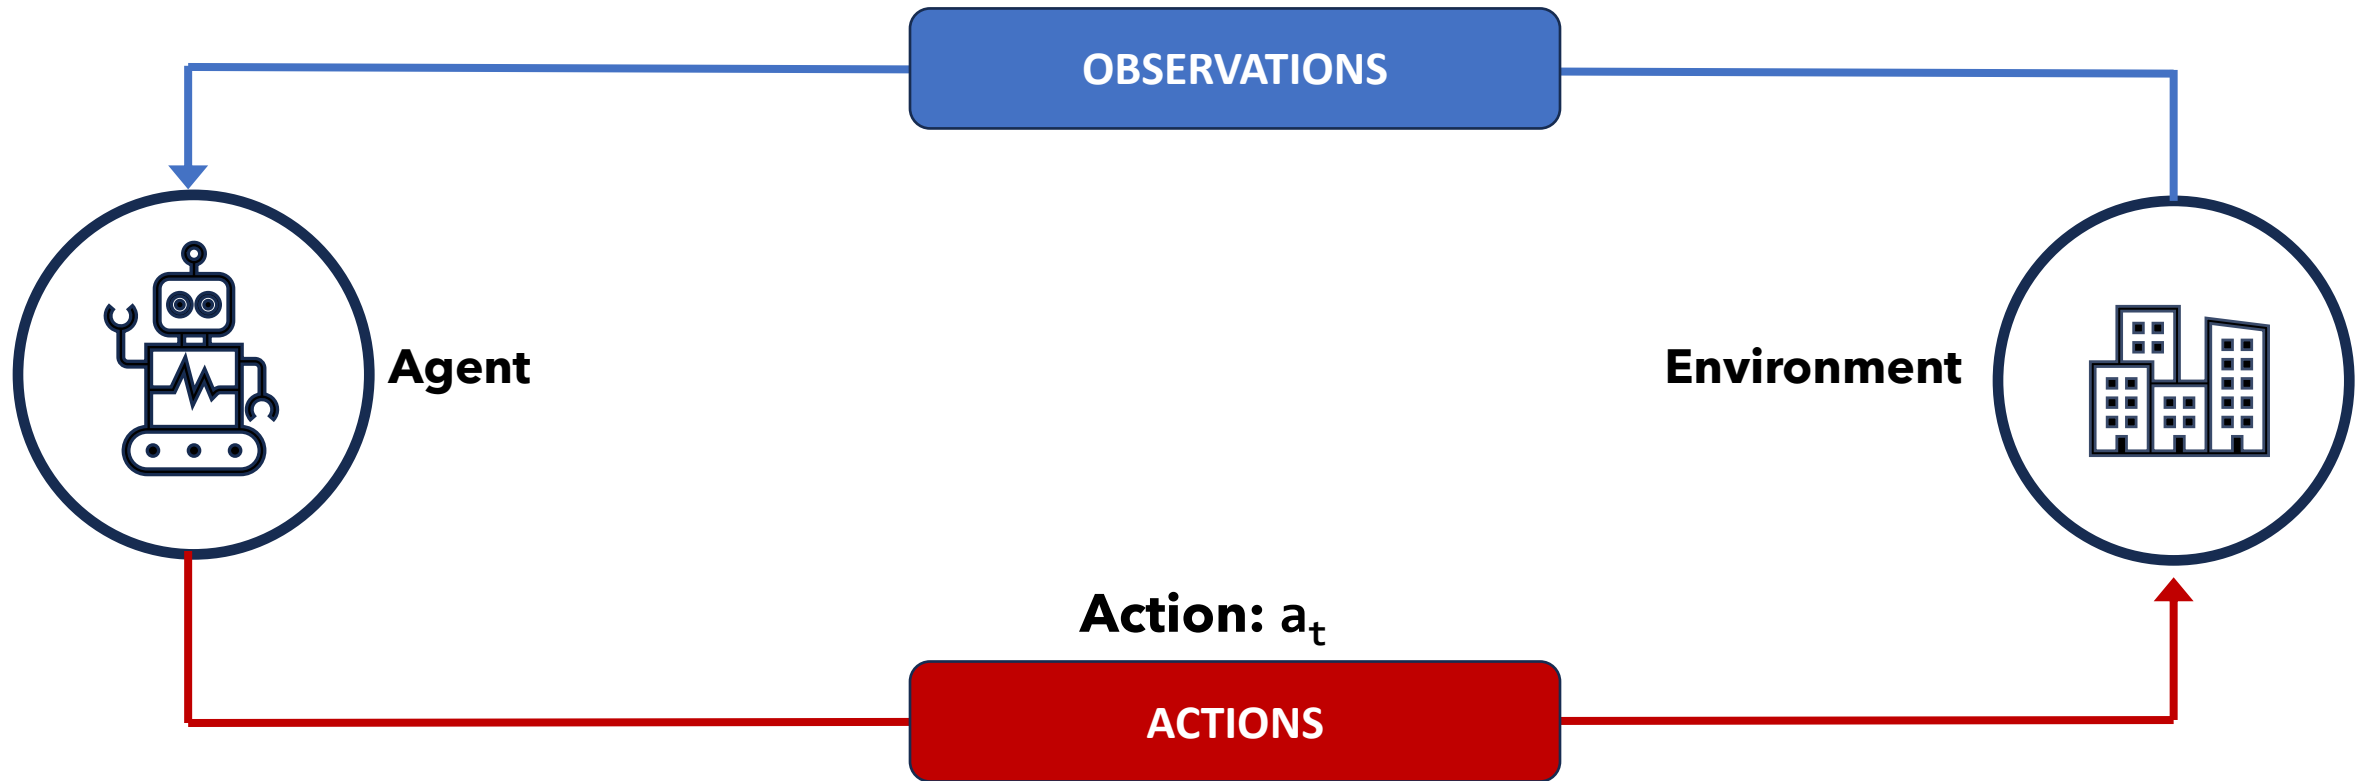
A move the agent can make in the Environment

Action Space **A**: The set of possible actions an agent can perform in the environment

$$\mathbf{A} = \{a_1, a_2, \dots, a_n\}$$

Key Concepts

Key Concepts: Observations

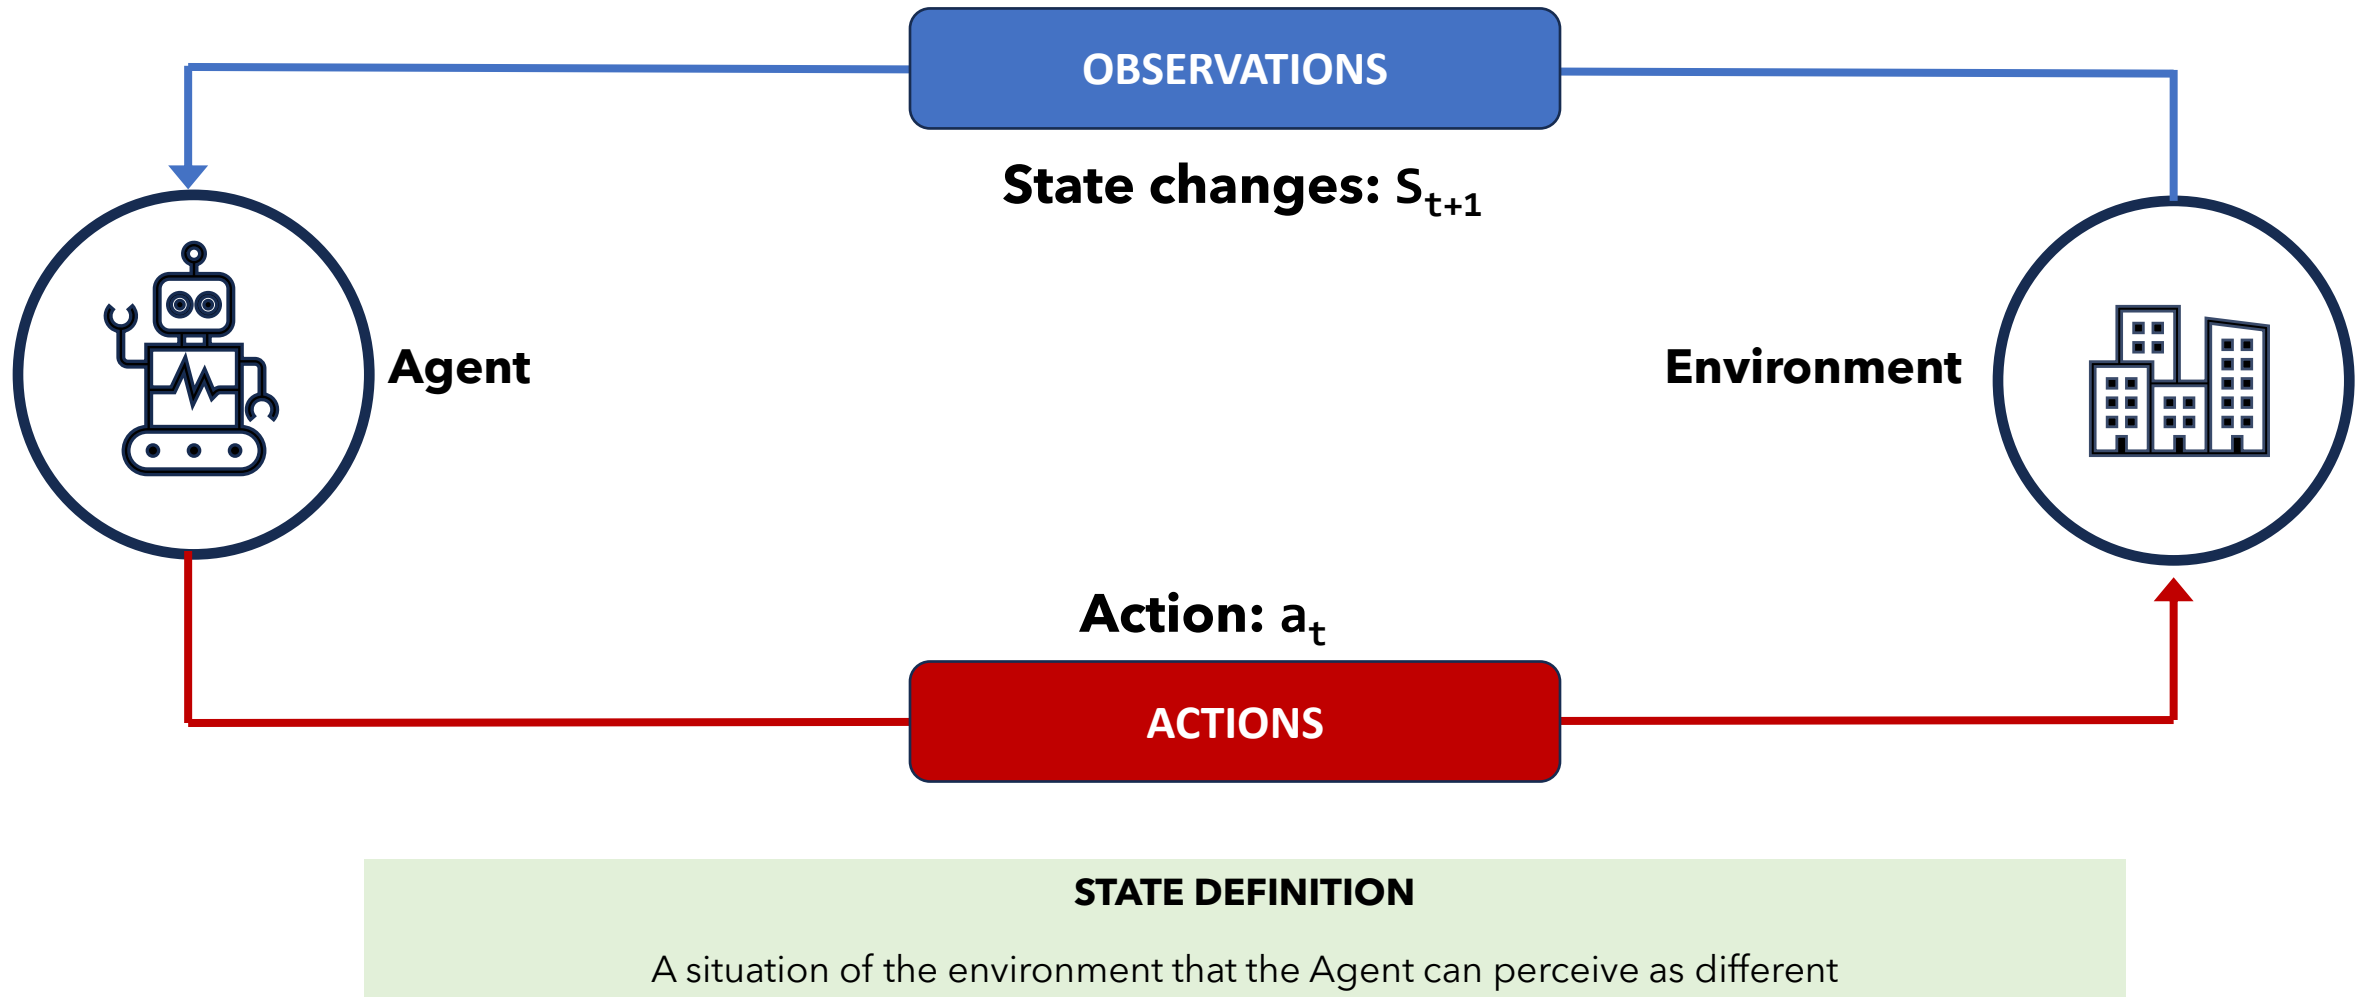


OBSERVATION DEFINITION

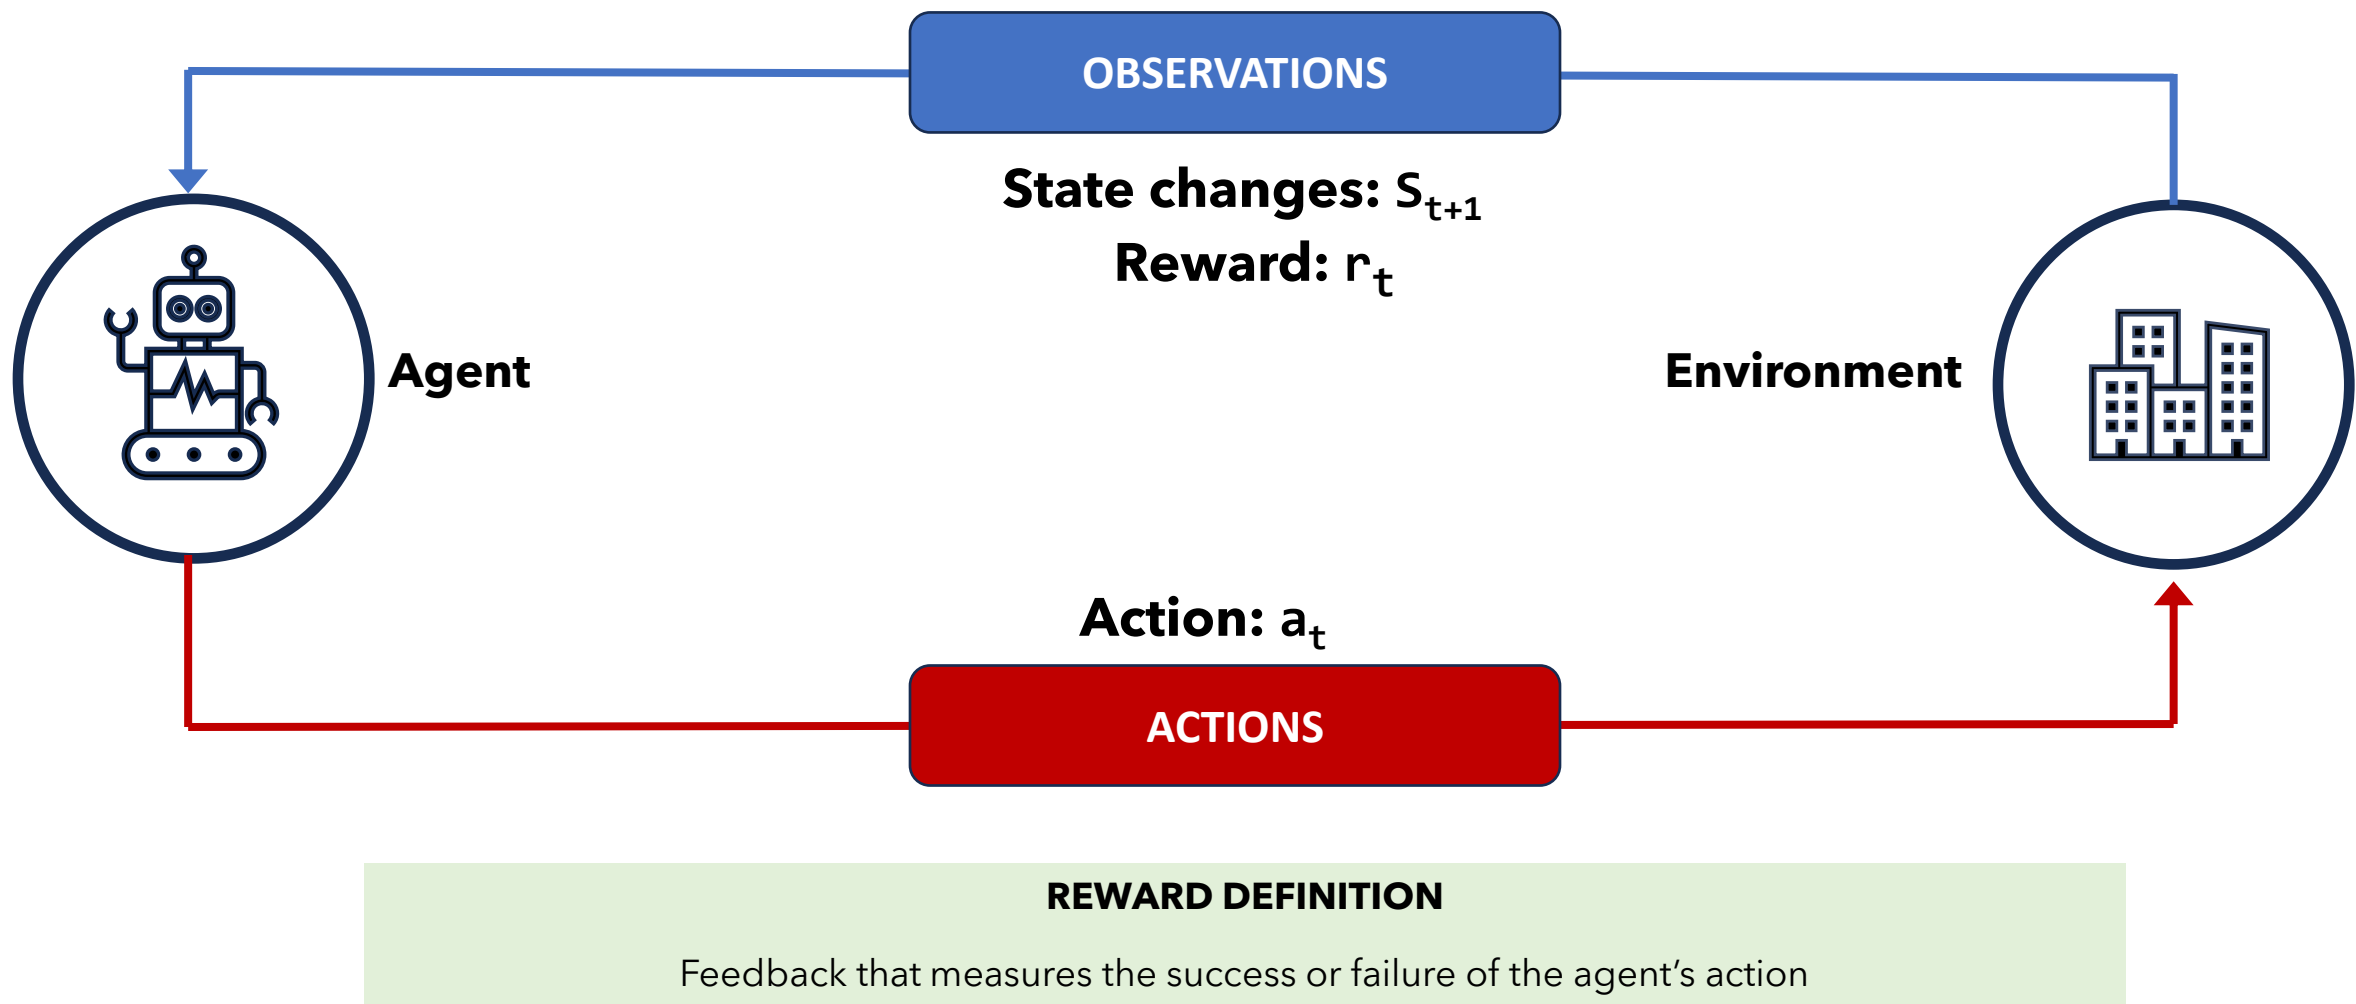
Understand the environment after taking actions
(What has changed from last observation, what is new, ...)

Key Concepts

Key Concepts: State and State change




Key Concepts
Key Concepts: Reward



Key Concepts

Key Concepts: Total Reward



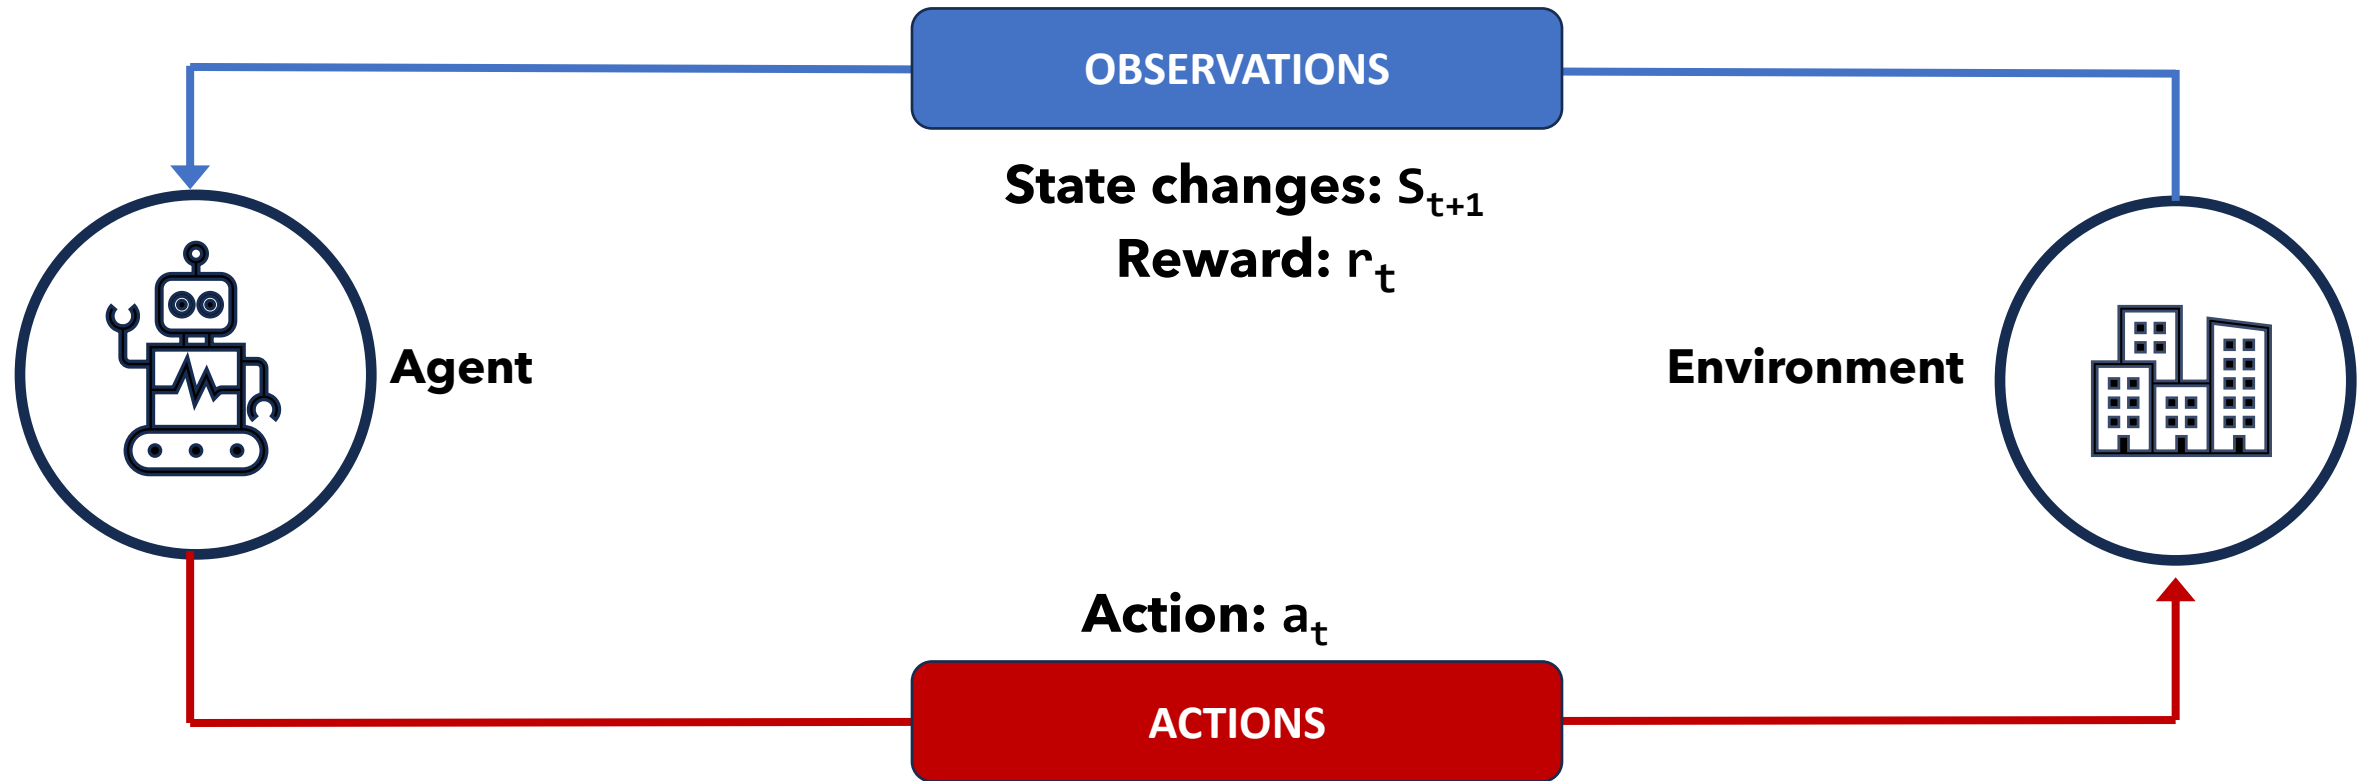
TOTAL REWARD (Return) DEFINITION

Is the Summarization of the total rewards pending until the end of the movement in the universe

$$R_t = \sum_{i=t}^{\infty} r_i$$

Key Concepts

Key Concepts: Total Reward decomposition



DEFINITION

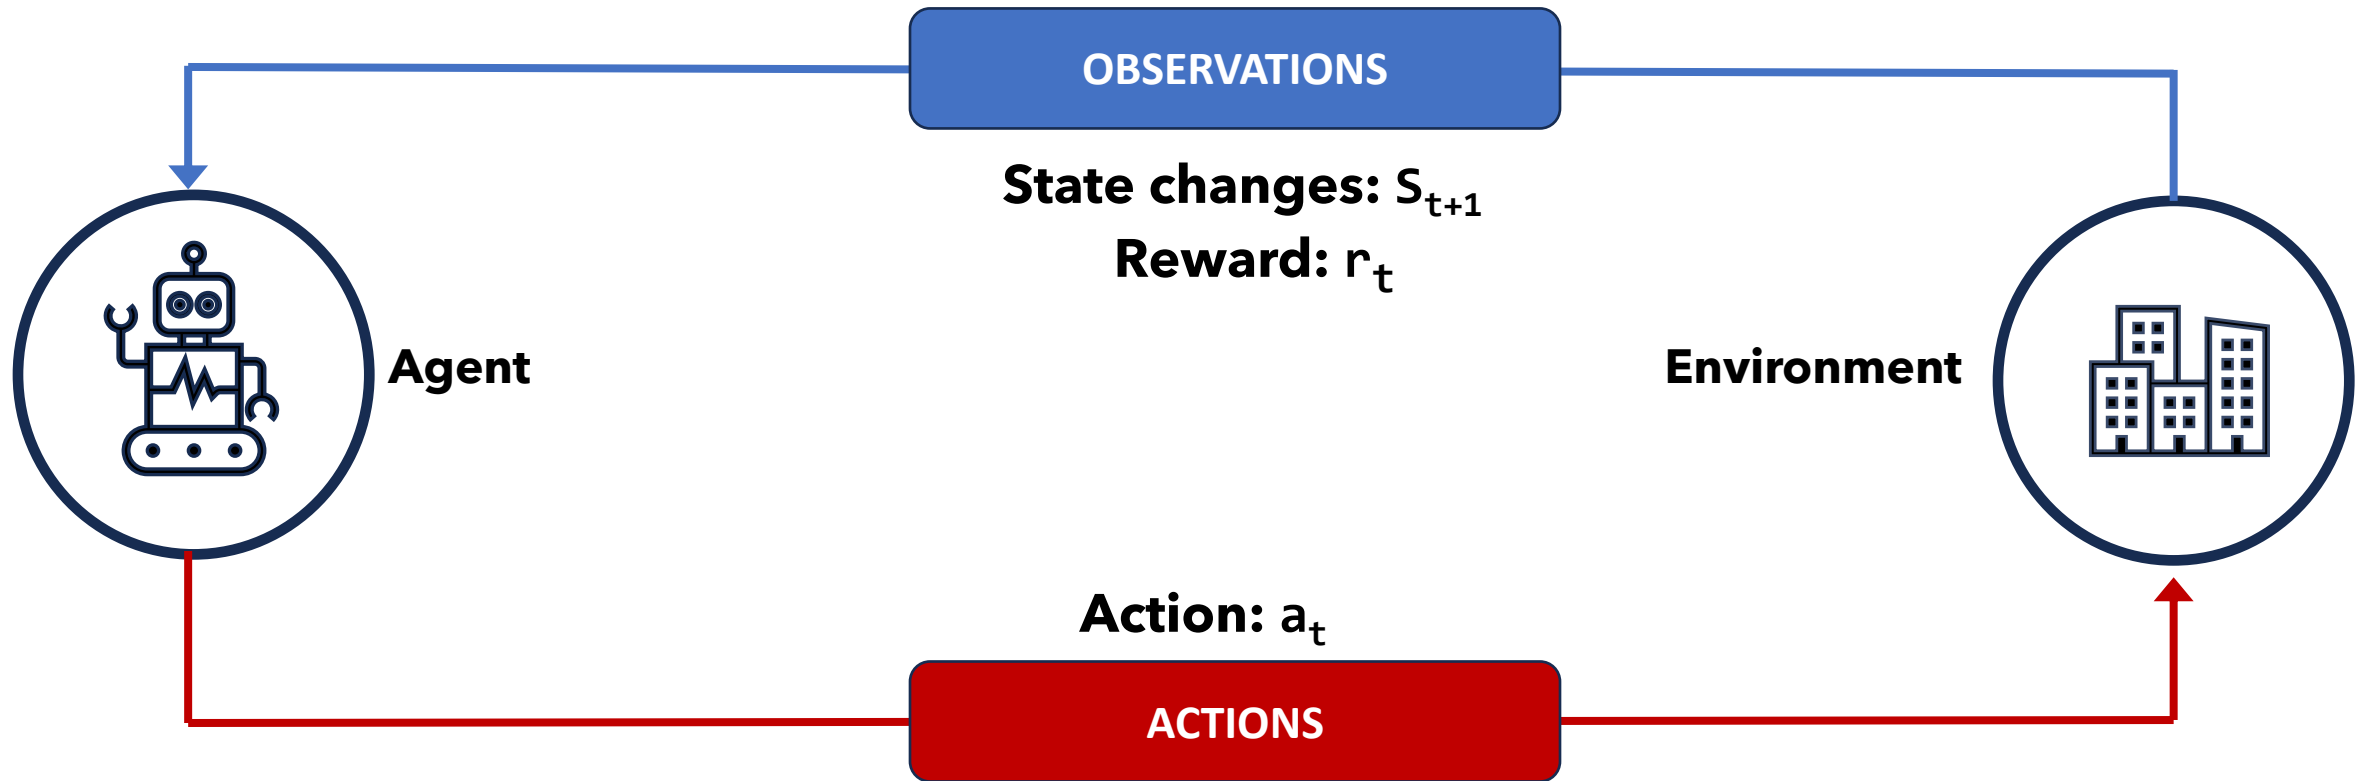
TOTAL REWARD (Return)

Is the Summarization of the total rewards pending until the end of the movement in the universe

$$R_t = \sum_{i=t}^{\infty} r_i = r_t + r_{t+1} + \dots + r_{t+n} + \dots$$

Key Concepts

Key Concepts: Discounted Total Reward



DEFINITION

Discount is a value between 0 and 1

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i$$

Review Concepts

Exploration vs exploitation

TD(0) Prediction - Exploitation-Exploration

Input, the policy π to be evaluated

$\epsilon \leftarrow$ Exploration Rate

Initialize $V(s)$ arbitrarily (e.g. $V(s) = 0, \forall s \in S^+$)

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

if $random < \epsilon$ **then**

$A \leftarrow$ Random possible action

else

$A \leftarrow$ action given by π for S

$A \leftarrow$ action given by π for S

Take action A ; observe reward, R and next state S'

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

$\epsilon \leftarrow$ decay

until S is terminal

Review Concepts

Exploration vs exploitation

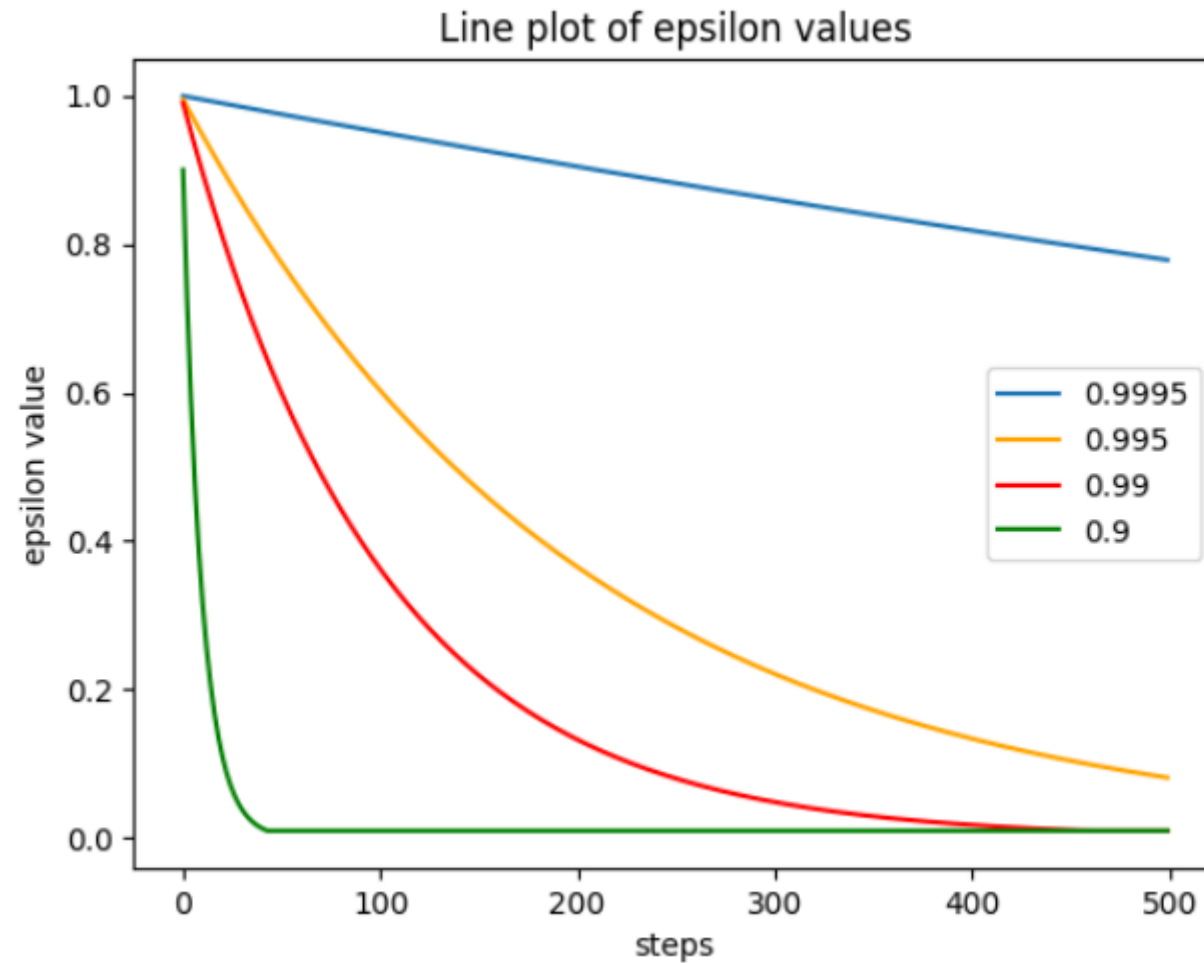
- Simplest idea for ensuring continual exploration
- All m actions are tried with non-zero probability
- With probability $1 - \epsilon$ choose the greedy action
- With probability ϵ choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a = \arg \max_{a' \in \mathcal{A}} Q(s, a') \\ \epsilon/m, & \text{otherwise} \end{cases}$$

where $m = |\mathcal{A}(s)|$

Review Concepts

Epsilon is a sensible parameter



00_Epsilon_decay_visualization.ipynb

Reinforcement Learning

Policy vs Value learning

Policy Learning

- The policy is modeled and updated directly without consulting a value function.
- Policy gradient methods are commonly used to optimize the policy by estimating which direction improves returns

Value Learning

- In value-based RL, the focus is on learning the optimal value function, denoted Q^* or V^* . The value function estimates expected cumulative future rewards for being in a given state and following the current policy thereafter.
- Key notes on value-based methods:
- Finding the optimal value function allows deriving the optimal policy.
- Temporal-difference learning is commonly used to update value estimates.
- Algorithms Monte-Carlo, Q-Learning, SARSA, Actor-critic

Reinforcement Learning Policy

“Right action for each state”

Policy Format (4x4):

S → ↓ ←
↓ X ↓ X
→ ↓ ↓ X
X → → G



The Policy function tells us the best action to take on each state

Reinforcement Learning

Value Function

Value Function Format (4x4):

S	0.656	0.729	0.656
0.656	X	0.810	X
0.729	0.810	0.900	X
X	0.900	1.000	G




The Value function gives us an estimate of value for each possible movement from our actual state

Reinforcement Learning


Environment examples



- Action space : 4
- Observation space: 16 – 4x4
- Ice patches: (3,0), (1,1), (1, 3), (2,3)
- Goal: (3,3)
- Termination: Reaches goal



- Action space : 6
- Observation space: 500 (5x5) There are 500 discrete states since there are 25 taxi positions, 5 possible locations of the passenger (including the case when the passenger is in the taxi), and 4 destination locations.
- Starting state: 300 starting states
- Termination: drops passenger




Stochastic Environment

Stochastic Environments

Definition

- A stochastic environment in Reinforcement Learning (RL) is one in which the outcome of an agent's action is not deterministic – that is, taking the same action in the same state may lead to different results with some probability.



Stochastic Effect: If you try to move **right**, there's a chance you'll slip **down** or **stay in place**.

If **is_slippery=True** the player will move in intended direction with probability of $1/3$ else will move in either perpendicular direction with equal probability of $1/3$ in both directions.

For example, if action is left and **is_slippery** is True, then:

- $P(\text{move left}) = 1/3$
- $P(\text{move up}) = 1/3$
- $P(\text{move down}) = 1/3$

Stochastic Environments

Solving Stochastic environments

- Agents must **learn expectations**, not absolute rules.
- Algorithms like SARSA or Q-learning estimate **expected value functions**

$$Q(s, a) = E[R]$$

- **Exploration** is crucial to properly model the stochasticity.
- In stochastic environments, we don't aim for perfect prediction — we aim for optimal expectation.

Stochastic Environments

Real life examples of stochastic environments

Real-world Examples:

Scenario	Why It's Stochastic
Robot on slippery floor	Wheels may slip; direction changes unpredictably
Stock trading agent	Market responds probabilistically to actions
Game with dice (e.g., Monopoly)	Player's move depends on dice outcome
Medical treatment agent	Same treatment may affect patients differently

SARSA

State-action-reward-state-action

- Updates the action-value function based on experiences in an environment
- **Action-value function is the $Q(s,a)$ function**
- Requires a decaying exploration strategy (like gradually reducing ϵ in ϵ -greedy policies)

SARSA

Definition

- SARSA is the acronym of State-action-reward-state-action
- Is a model-free on-policy method
- Idea: Maximize the cumulative reward obtained by the agent over time
- It is basically a TD algorithm on-policy
- The goal is to update the action-value function $Q(s,a)$
- Operates by estimating the Q -values iteratively based on the experiences
- The main update rule in the process is :

$$Q(s, a) \leftarrow Q(s, a) + \alpha[R + \gamma Q(s', a') - Q(s, a)]$$

SARSA: On-Policy TD Control

Initialize $Q(s, a), \forall s \in S, a \in A(s)$ arbitrarily and $Q(\text{terminal_state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A ; observe reward, R and next state S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A'$

 until S is terminal

- Uses the Q-value of the actual next action a' chosen by the policy.
- This makes SARSA an on-policy algorithm because it updates the Q-values using the same policy that the agent follows during learning.

SARSA

ϵ -greedy by definition

- The learning process in SARSA involves the agent interacting with the environment in a sequence of episodes.
- In each episode:
 - the agent starts in an initial state and selects actions based on its current policy (e.g., ϵ -greedy policy).
 - The ϵ -greedy policy balances exploration and exploitation by selecting the action with the highest estimated Q-value with probability $1-\epsilon$ and selecting a random action with probability ϵ .
- This allows the agent to explore the environment and potentially discover better actions while still exploiting its current knowledge.

SARSA

Why is SARSA on-policy

- Updates the values based the actions taken by the actual policy
- Continuously refines its policy
- The Q-values are intimately related to the previous experience of the agent and improve as the experience increases. The update to the policy is based on the experience
- Very important SARSA is on-policy while Q-learning is off-policy

SARSA

Why is SARSA good?

- Fairly straightforward to implement
- Can handle discrete and continuous state and action spaces
- It converges most of the time
- But
 - Selection learning rate (α) and discount factor (γ) impacts algorithm efficiency and convergence
 - If learning rate is too high – maybe does not converge
 - If learning is too low – slow learning
 - Not very efficient for large action-state spaces

SARSA

SARSA vs Q-learning

- SARSA shares similarities with Q-learning
- Both algorithms aim to learn the optimal action-value function and use similar update rules. However, the key difference lies in their learning strategies.
- While Q-learning is an off-policy algorithm that updates the Q-values based on the maximum expected future reward, regardless of the action taken, SARSA updates the Q-values based on the actual action taken by the current policy. This difference can lead to different learning dynamics and performance characteristics in certain problems.

SARSA

SARSA vs Q-learning



<https://www.youtube.com/watch?v=ma8MgPB0iXc>

<https://www.youtube.com/watch?v=vhuQLxvflxg>

26:28 – 27:...

Q-learning

Q-learning

Q-learning is off-policy

- While Q-learning is an off-policy algorithm that updates the Q-values based on the maximum expected future reward, regardless of the action taken, SARSA updates the Q-values based on the actual action taken by the current policy. This difference can lead to different learning dynamics and performance characteristics in certain scenarios.
- It is off-policy!!

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

Q-learning

The algorithm

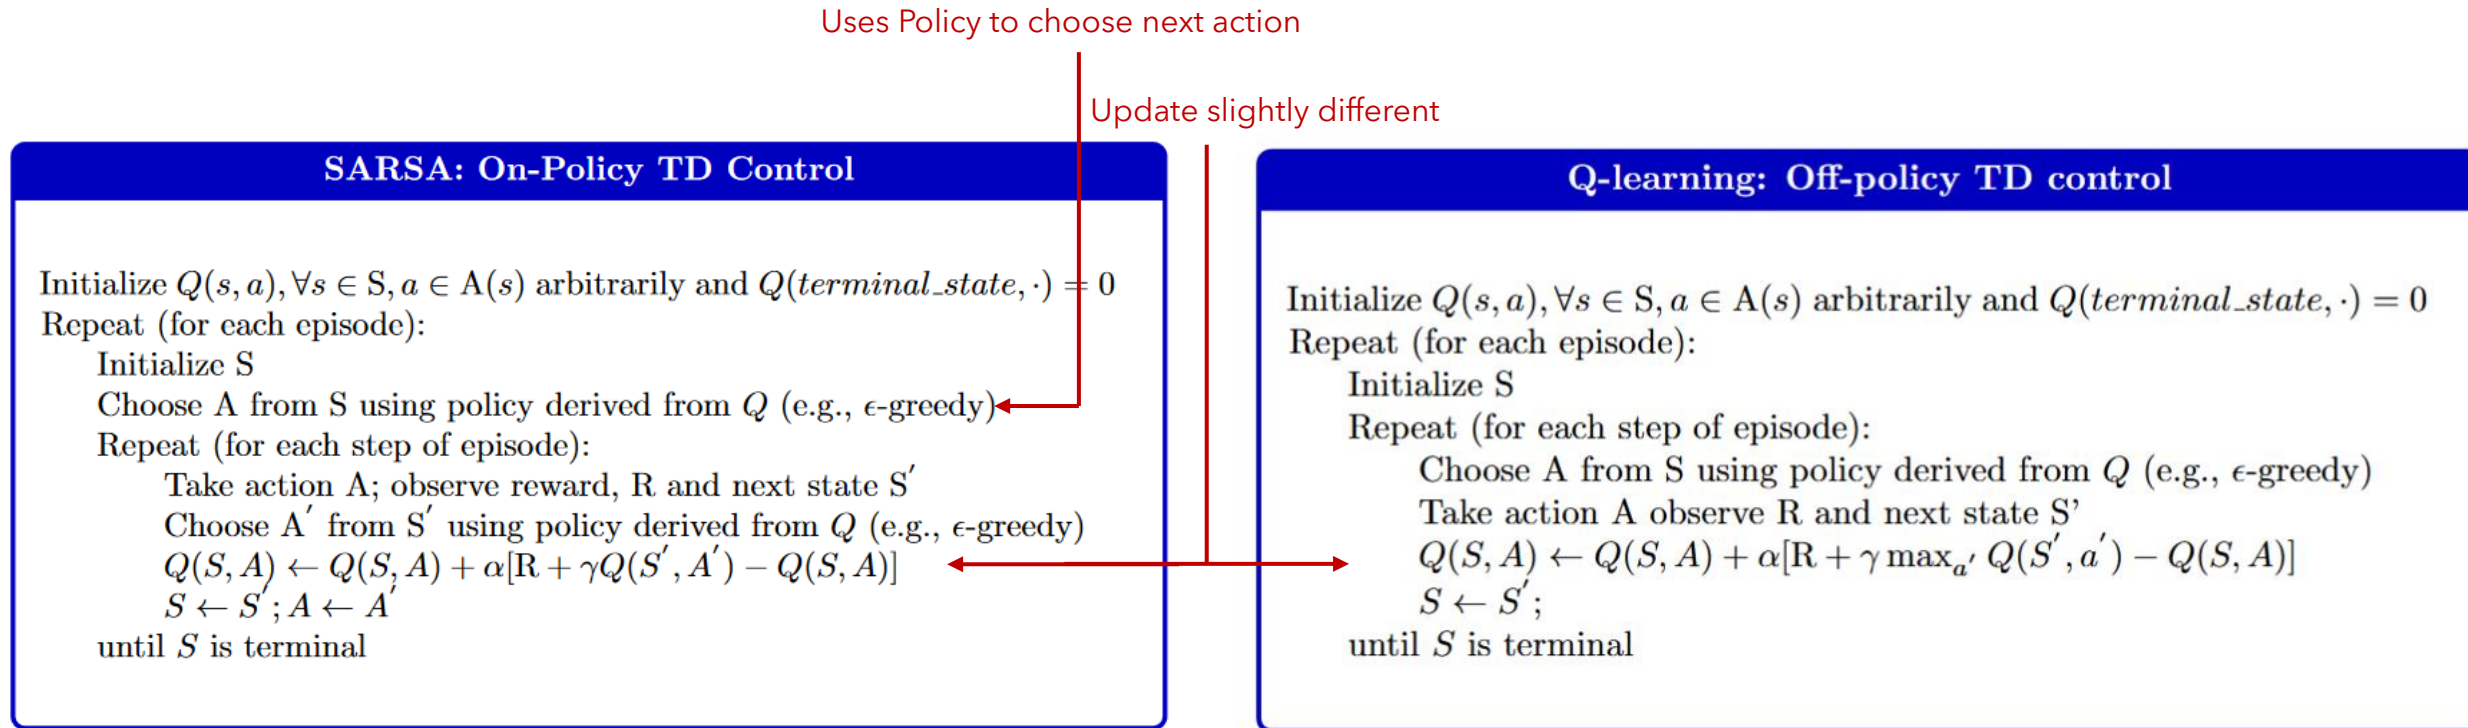
Q-learning: Off-policy TD control

```
Initialize  $Q(s, a), \forall s \in S, a \in A(s)$  arbitrarily and  $Q(\text{terminal\_state}, \cdot) = 0$ 
Repeat (for each episode):
  Initialize  $S$ 
  Repeat (for each step of episode):
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
    Take action  $A$  observe  $R$  and next state  $S'$ 
     $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_{a'} Q(S', a') - Q(S, A)]$ 
     $S \leftarrow S'$ ;
  until  $S$  is terminal
```

- Uses the maximum possible future *Q-value* for the next state s' , regardless of the agent's actual next action
- This makes Q-Learning an off-policy algorithm because it learns from the best possible actions rather than the actions actually taken by the current policy.

Q-learning

See the Differences



1. The most important difference between the two is how Q is updated after each action. SARSA uses the \mathbf{Q}' following a ϵ -greedy policy exactly, as A' is drawn from it. In contrast, Q-learning uses the maximum \mathbf{Q}' over all possible actions \mathbf{a}' for the next step. This makes it look like following a greedy policy with $\epsilon=0$, i.e. NO exploration in this part.
2. However, when actually taking an action, Q-learning still uses the action taken from a ϵ -greedy policy. This is why "Choose A ..." is inside the repeat loop.
3. Following the loop logic in Q-learning, A' is still from the ϵ -greedy policy.

Q-learning

See the Differences

	SARSA	Q-learning
Choosing A'	π	π
Updating Q	π	μ

Where π is a ϵ -greedy policy (e.g. $\epsilon > 0$ with exploration), and μ is a greedy policy (e.g. $\epsilon == 0$, NO exploration).

1. Given that Q-learning is using different policies for choosing next action A' and updating Q. In other words, it is trying to evaluate π while following another policy μ , so it's an off-policy algorithm.
2. In contrast, SARSA uses π all the time, hence it is an on-policy algorithm.

<https://stackoverflow.com/questions/6848828/what-is-the-difference-between-q-learning-and-sarsa>

Example

Q-learning

See the Differences

```
# Q-Learning algorithm - Commens changes from SARSA example
for episode in range(episodes):
    state, _ = env.reset()
    done = False
    while not done:
        action = choose_action(state)  # Action selection moved here for Q-learning

        next_state, reward, done, _, _ = env.step(action)

        # Q-learning update
        # The key difference is using max Q-value of next state instead of Q-value of next action
        Q[state, action] += alpha * (reward + gamma * np.max(Q[next_state, :]) - Q[state, action])

        state = next_state
        # Removed: action = next_action (Q-learning is off-policy, so we don't need to track the next action)
```

```
# SARSA algorithm
for episode in range(episodes):
    state, _ = env.reset()
    action = choose_action(state)
    done = False
    while not done:

        next_state, reward, done, _, _ = env.step(action)

        next_action = choose_action(next_state)

        # SARSA update
        Q[state, action] += alpha * (reward + gamma * Q[next_state, next_action] - Q[state, action])

        state = next_state
        action = next_action
```

Q-learning

Exploration vs Exploitation

```
def choose_action(state):  
    if np.random.uniform(0, 1) < epsilon:  
        return env.action_space.sample() # Explore  
    else:  
        return np.argmax(Q[state, :]) # Exploit
```

Wrap-up

SARSA & Q-Learning

Policy Type

- SARSA is an on-policy algorithm: learns Q-values based on actions actually taken using current policy
- Q-Learning is off-policy: learns about optimal policy regardless of actions taken

Action Selection for Updates

- SARSA considers the next action that will actually be taken (including exploration)
- Q-Learning assumes optimal future actions regardless of exploration policy, using max Q-value

Convergence Properties

- SARSA converges to the optimal policy when using an epsilon-greedy policy with epsilon decreasing to zero
- Q-Learning converges to the optimal Q-function regardless of the exploration policy (as long as all state-action pairs are visited infinitely often)

Safety and Conservative Behavior

- SARSA tends to learn safer paths as it considers exploration in its value estimates
- Q-Learning can learn optimal paths that might be risky during training (because it assumes optimal future actions)

Exploration-Exploitation Trade-off

- SARSA directly incorporates exploration into its learning, making it more conservative
- Q-Learning separates exploration policy from learning, potentially leading to more optimal but riskier policies Example scenario:

SARSA

SARSA vs Q-learning



<https://www.youtube.com/watch?v=ma8MgPB0iXc>

<https://www.youtube.com/watch?v=vhuQLxvflxg>

26:28 – 27:...

Wrap-up

Q-learning

Summary differences SARSA and Q-Learning

Aspect	Q-Learning	SARSA
Policy Type	Off-policy	On-policy
Next Q-value Estimation	Uses $\max Q(s', a')$ (greedy action)	Uses $Q(s', a')$ of the actual action taken
Exploration vs. Exploitation	Encourages more exploitation (greedy updates)	Encourages more exploration (policy-following updates)
Risk Sensitivity	More aggressive, assumes optimal future actions	More conservative, considers the actual exploration strategy
Convergence Speed	Typically, faster but less stable in stochastic environments	More stable but slower

END

Session 7

