

Deep Learning: Session 4

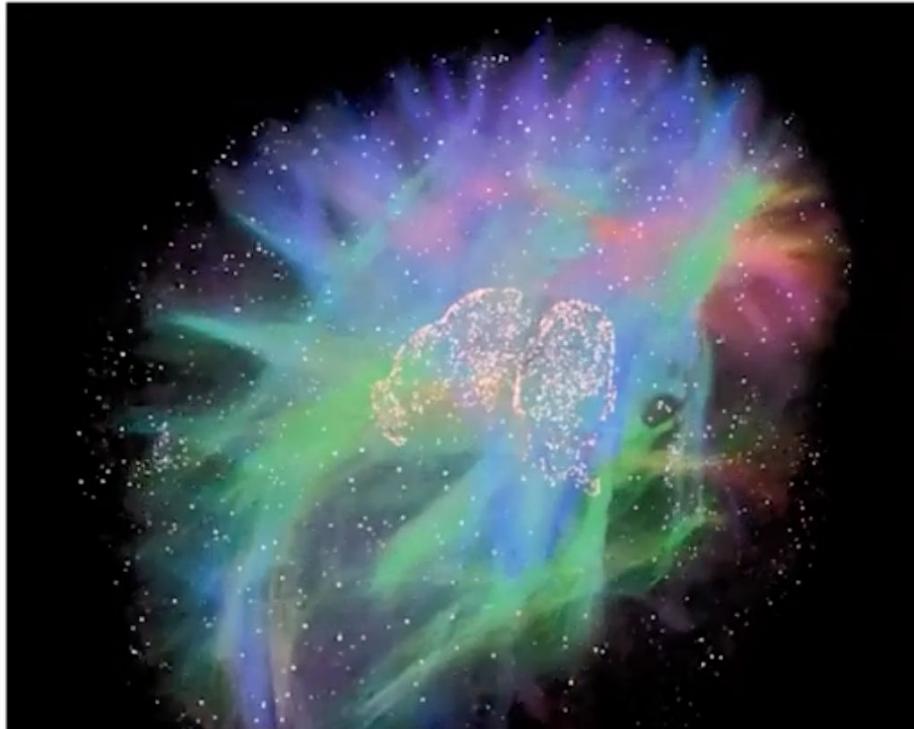
Deep Neural Nets

Outline



1. Recap
2. Deep L-Layers
3. Divide and Conquer
4. The Neuron: Pattern Recognition
5. DNNs

In previous chapters...



Human Brain

- Thalamocortical system:
3 million neurons
476 million synapses
- Full brain:
100 billion neurons
1,000 trillion synapses

Artificial Neural Network

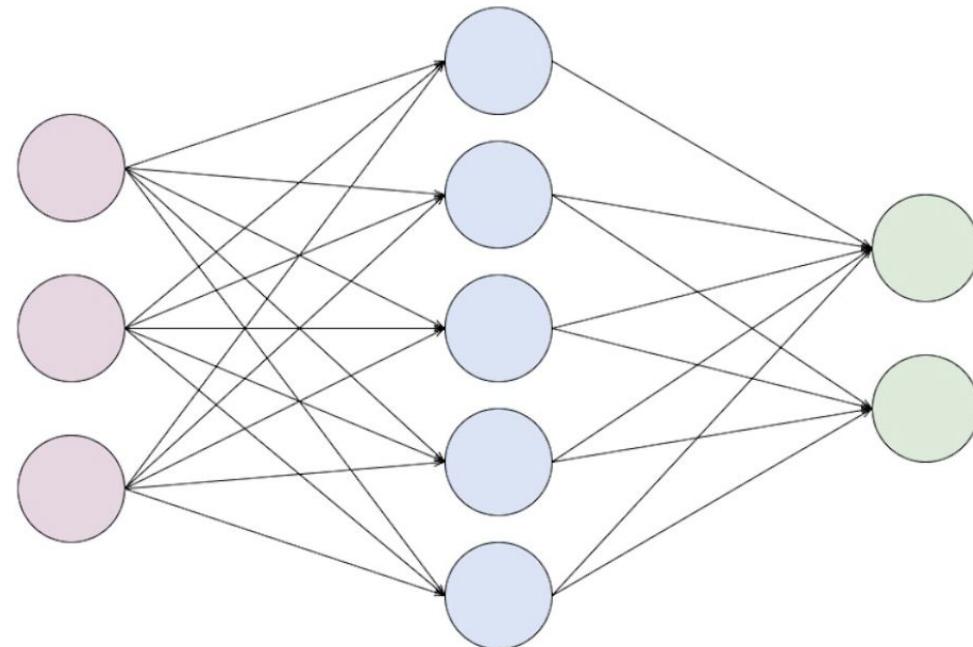
- ResNet-152:
60 million synapses

Human brains have ~10,000,000 times synapses than artificial neural networks.

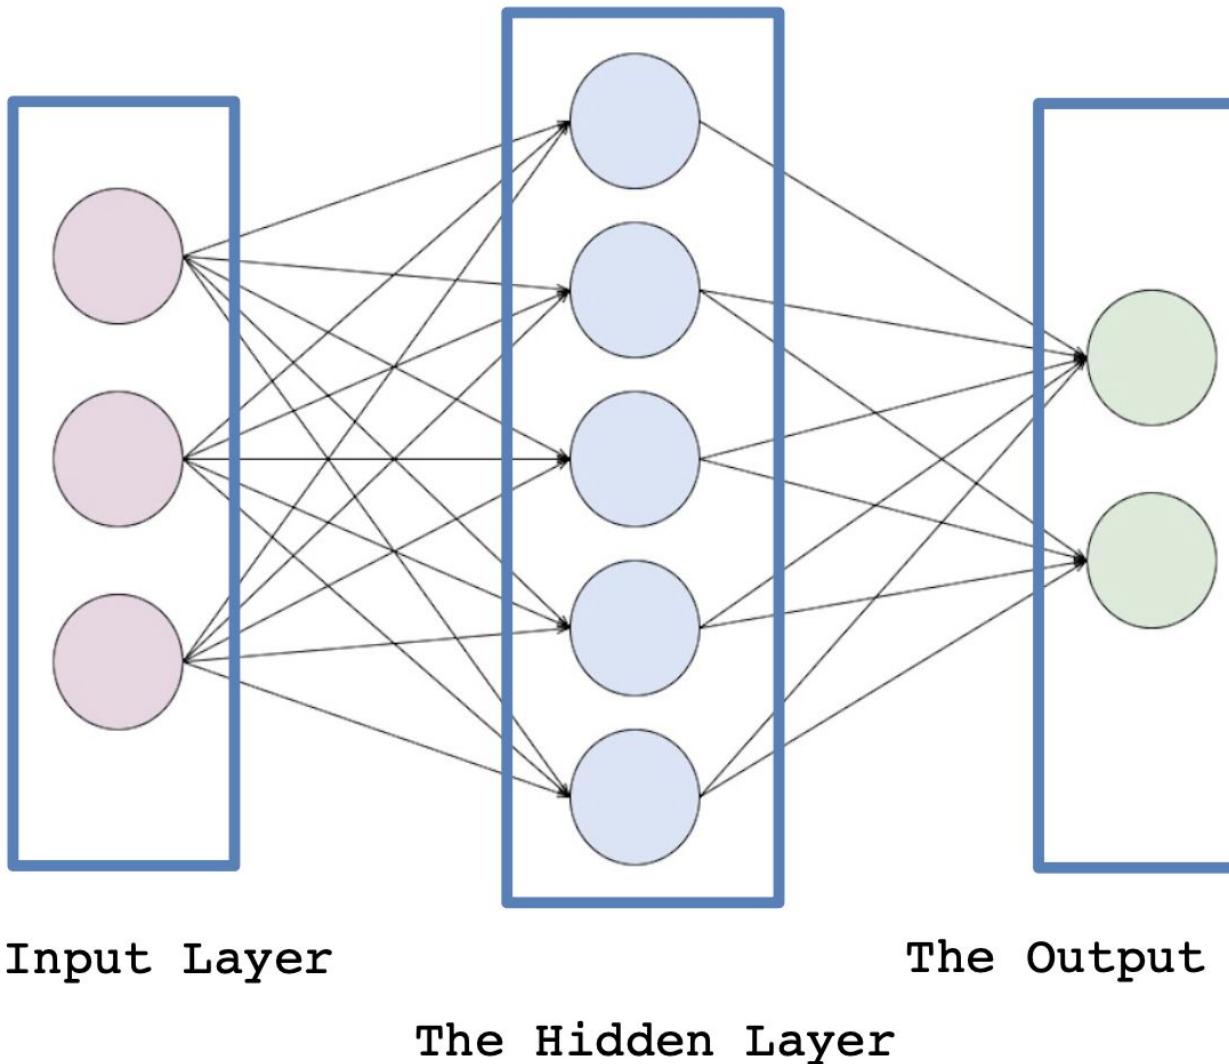
In previous chapters...



A net of information

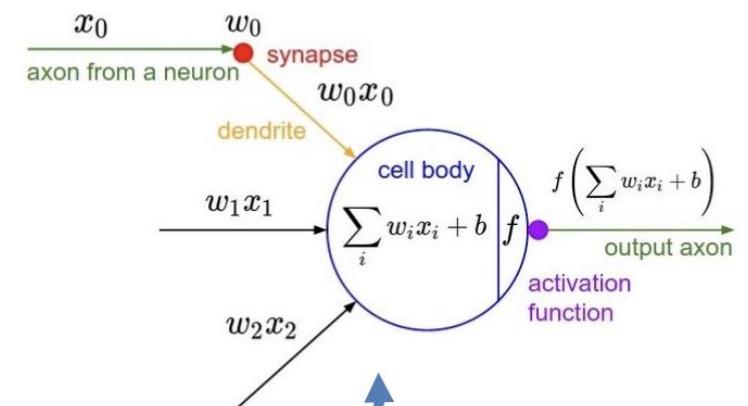
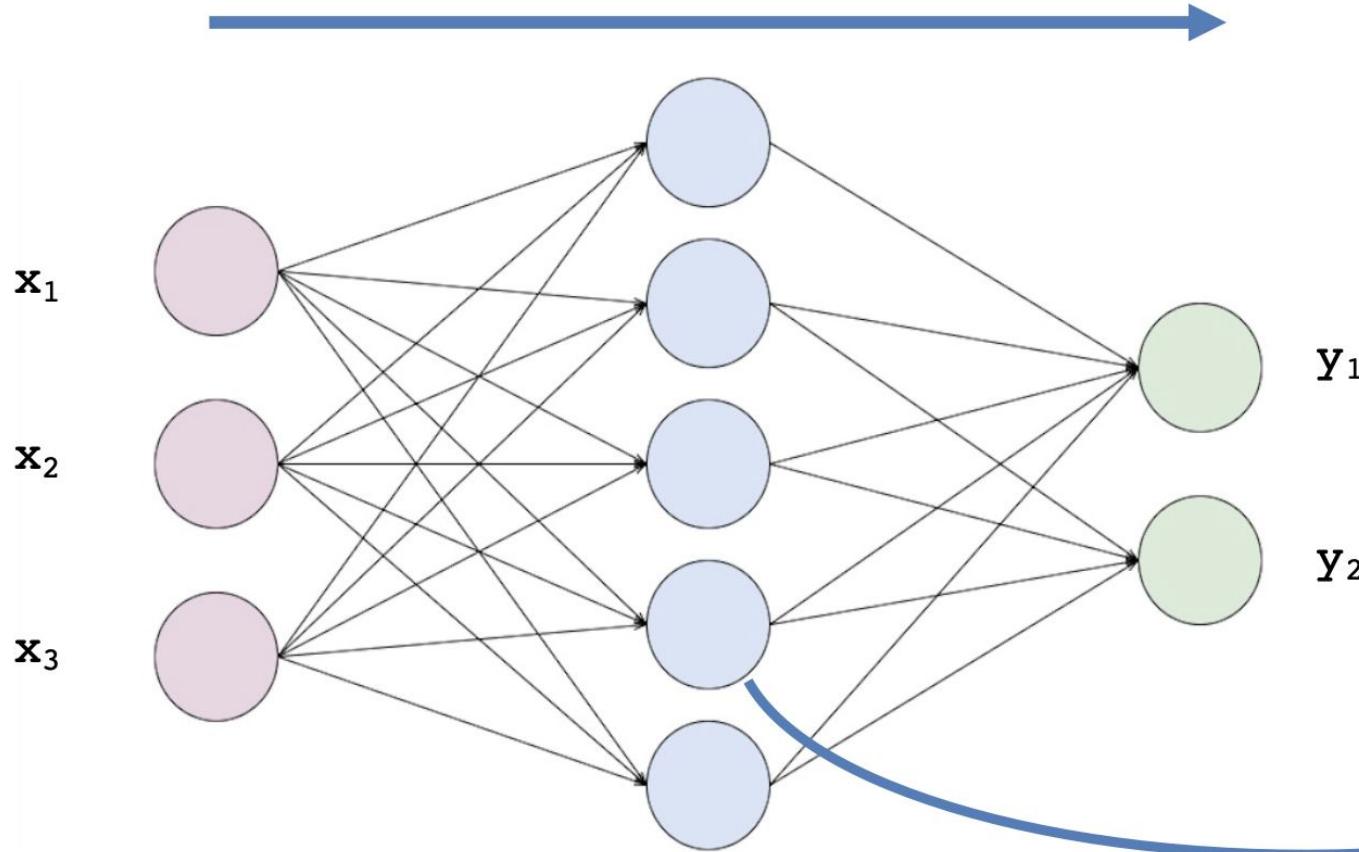


In previous chapters...



In previous chapters...

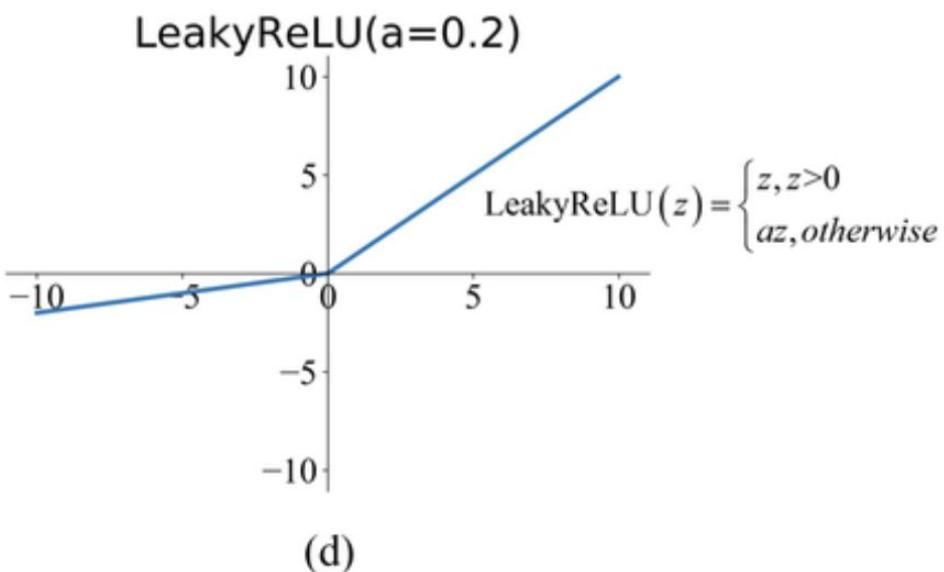
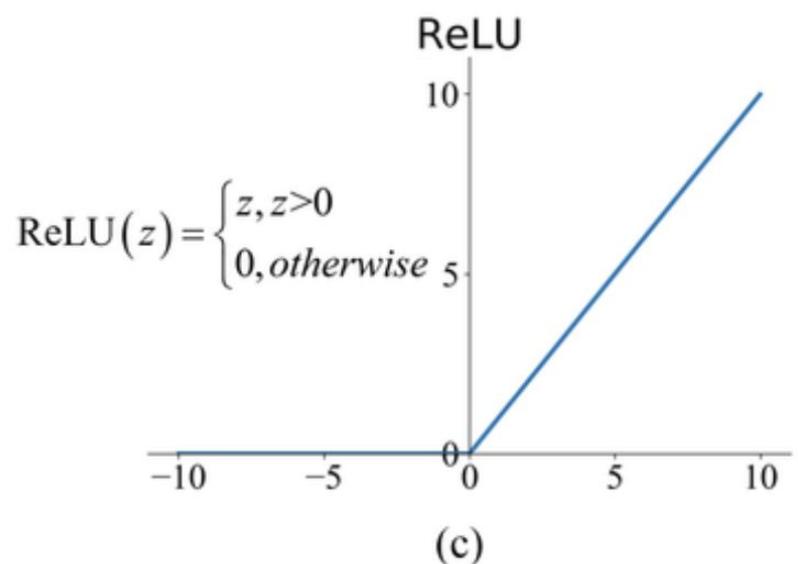
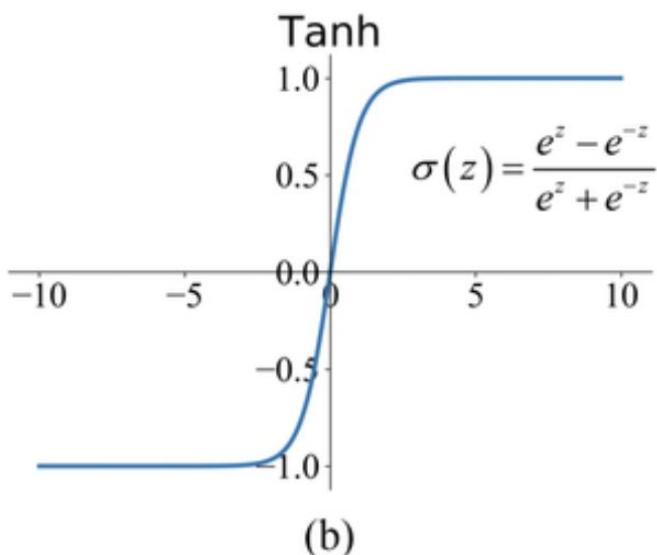
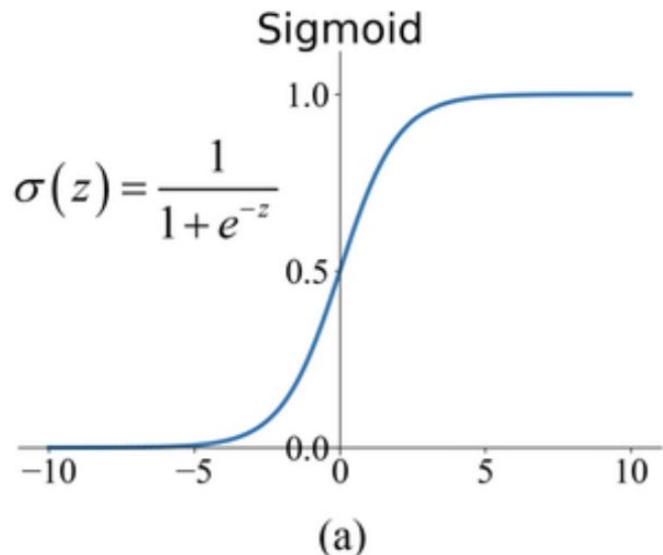
$$Y = F(x)$$



In previous chapters...

```
my_first_neural_net = tf.keras.Sequential([
    layers.Input(shape=(3, ), name="input"),
    layers.Dense(units=5, activation="sigmoid", name="hidden"),
    layers.Dense(units=2, activation="softmax", name="output")
])
```

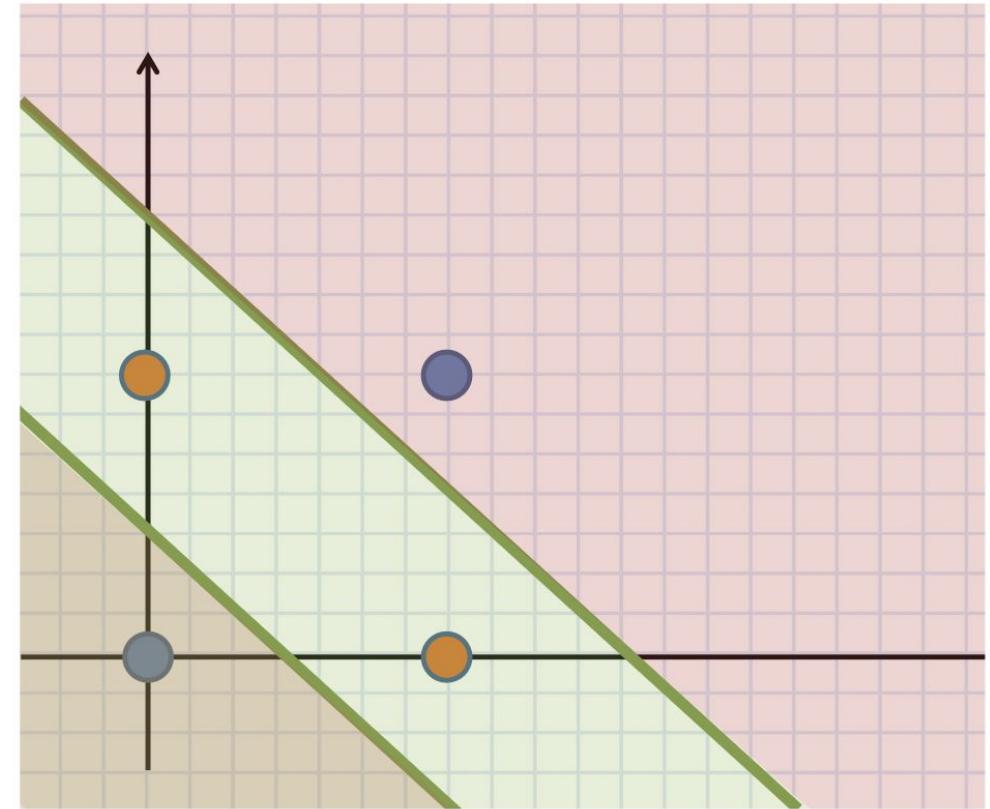
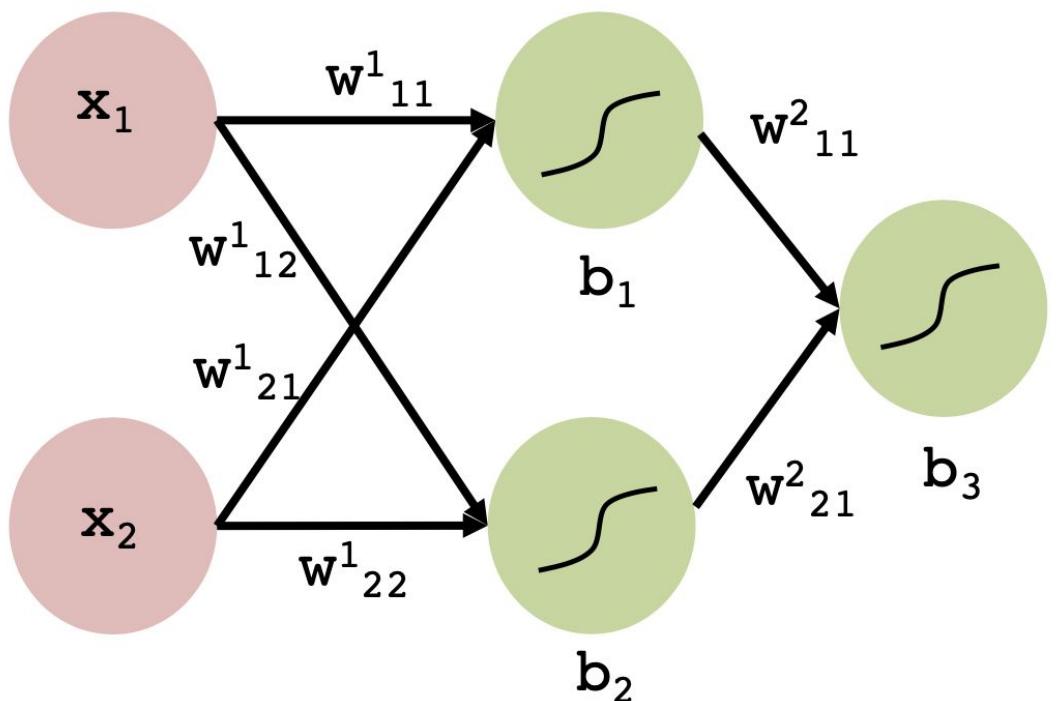
In previous chapters...



In previous chapters...

- The combination of linear functions is just another linear function. No matter the number of hidden layers, if we use linear functions as activations, we could just deal with outputs that are a linear combination of our inputs... not interesting.
- Non-linear activations functions are key in neural nets to represent non-linear mapping functions.
- **In the output layer**, at general: use sigmoid for binary classification, softmax for multiclass classification and linear for regression.
 - You can use softmax for binary classification using two outputs and in the case of a problem where data that can belong to multiple classes, then use sigmoid. [i.e., a picture of Audi belongs to the classes' car and vehicle]
- **In hidden layers:**
 - Tanh used to work better than sigmoid as it centers data.
 - ReLu is faster and it works very well too. It is the most used.

In previous chapters...



$$y = s(s(w^1_{11}x_1 + w^1_{21}x_2 + b_1) + s(w^1_{21}x_1 + w^1_{22}x_2 + b_2) + b_3)$$

Outline

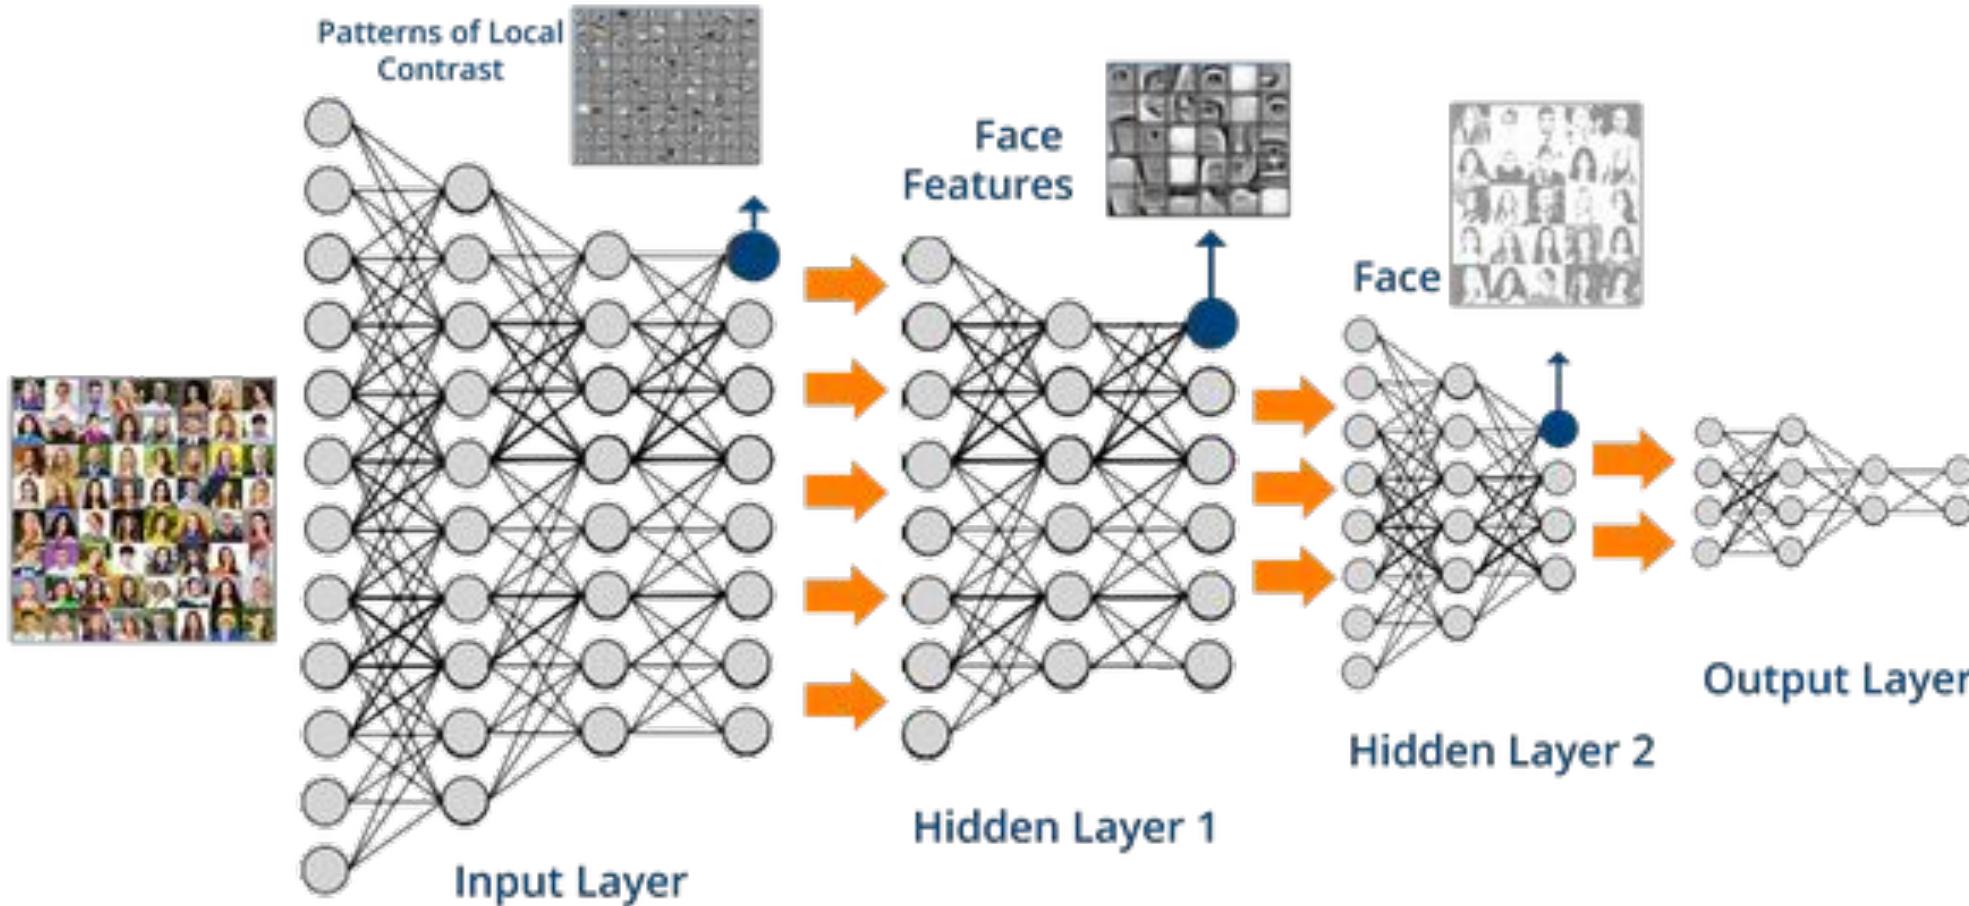


1. Recap
2. Deep L-Layers
3. Divide and Conquer
4. The Neuron: Pattern Recognition
5. DNNs

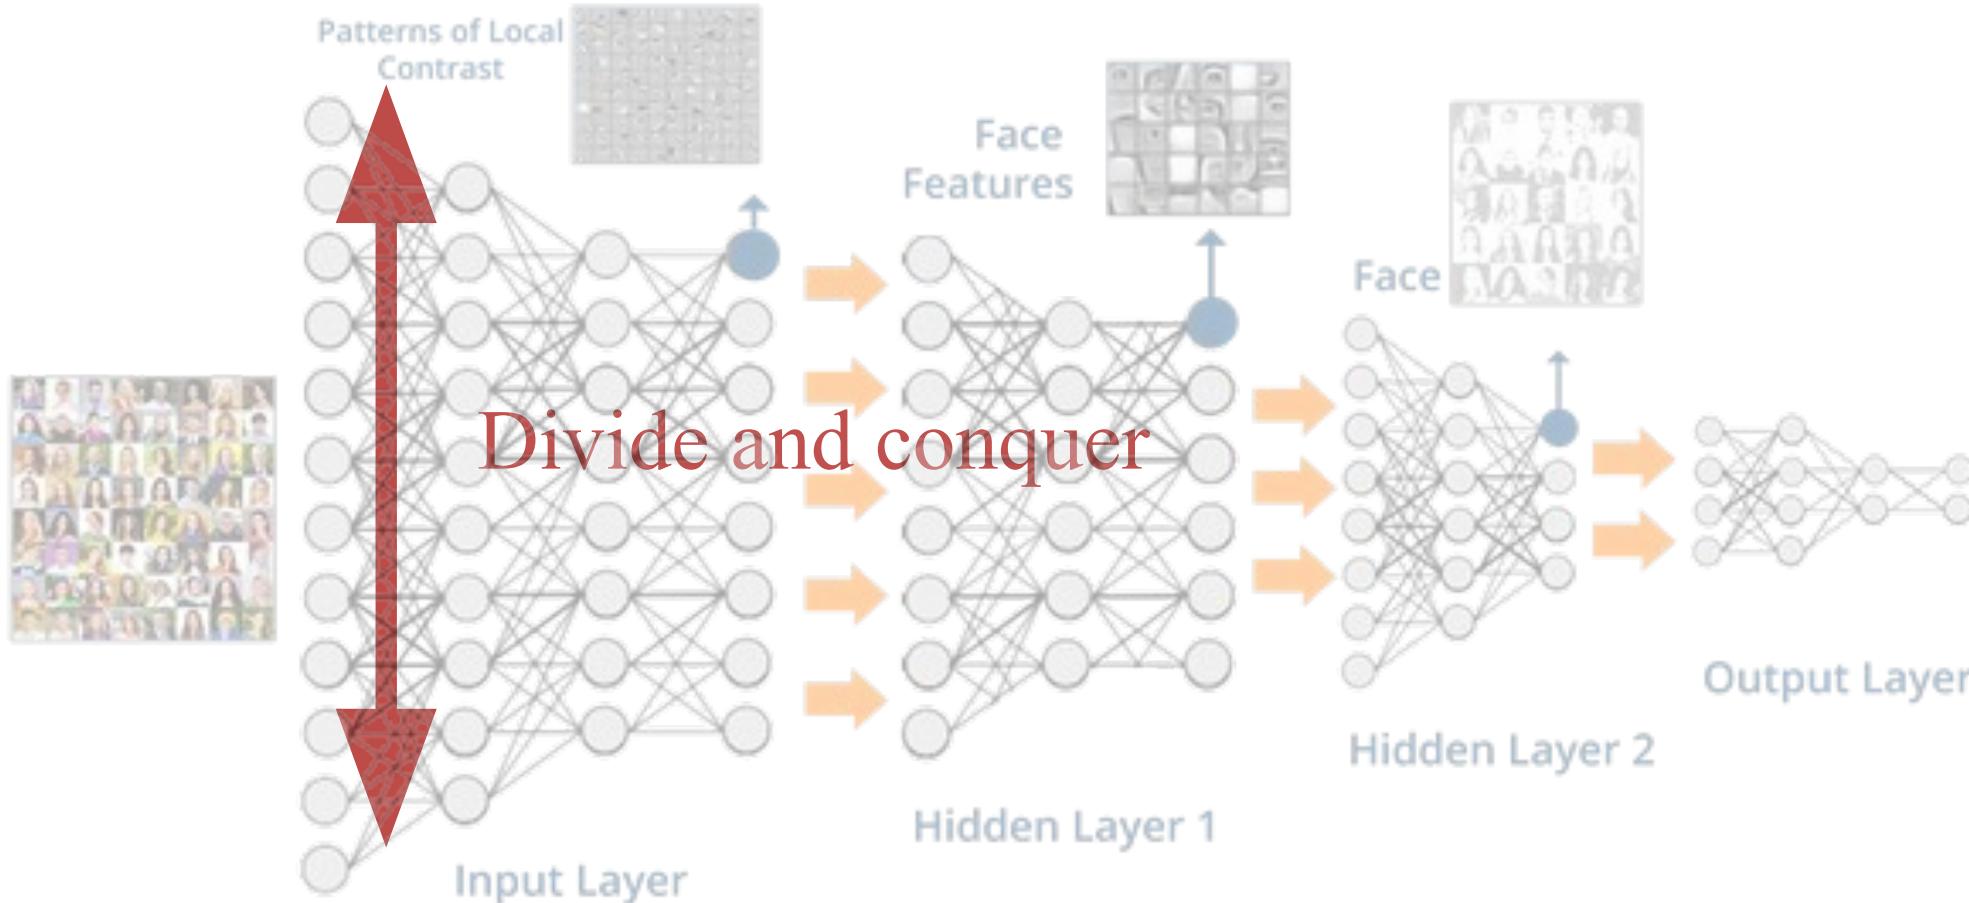
Why Deep L-layer Neural Nets

- From circuit theory:
 - We know that for representing some functions we could need many components in shallow architectures.
 - ... Functions that we can easily represent if we include more layers.
- The Divide & Conquer strategy is powerful, and it is efficient.

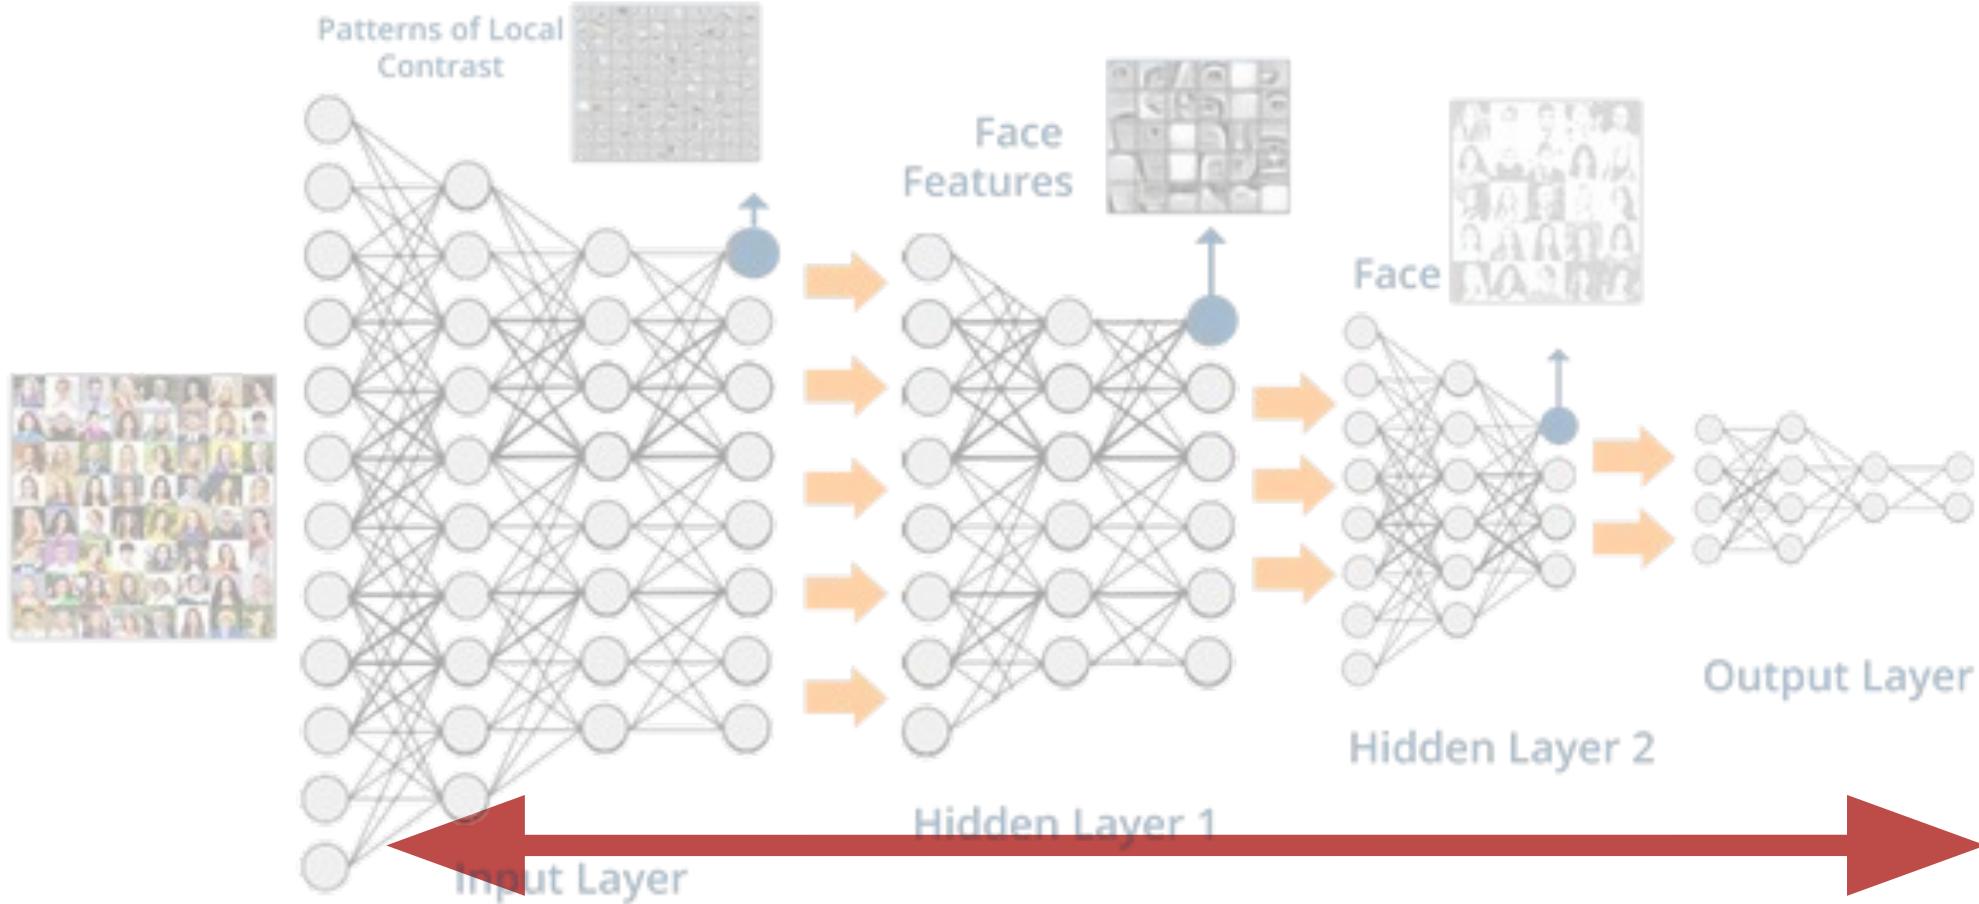
Deep L-Layers



Deep L-Layers

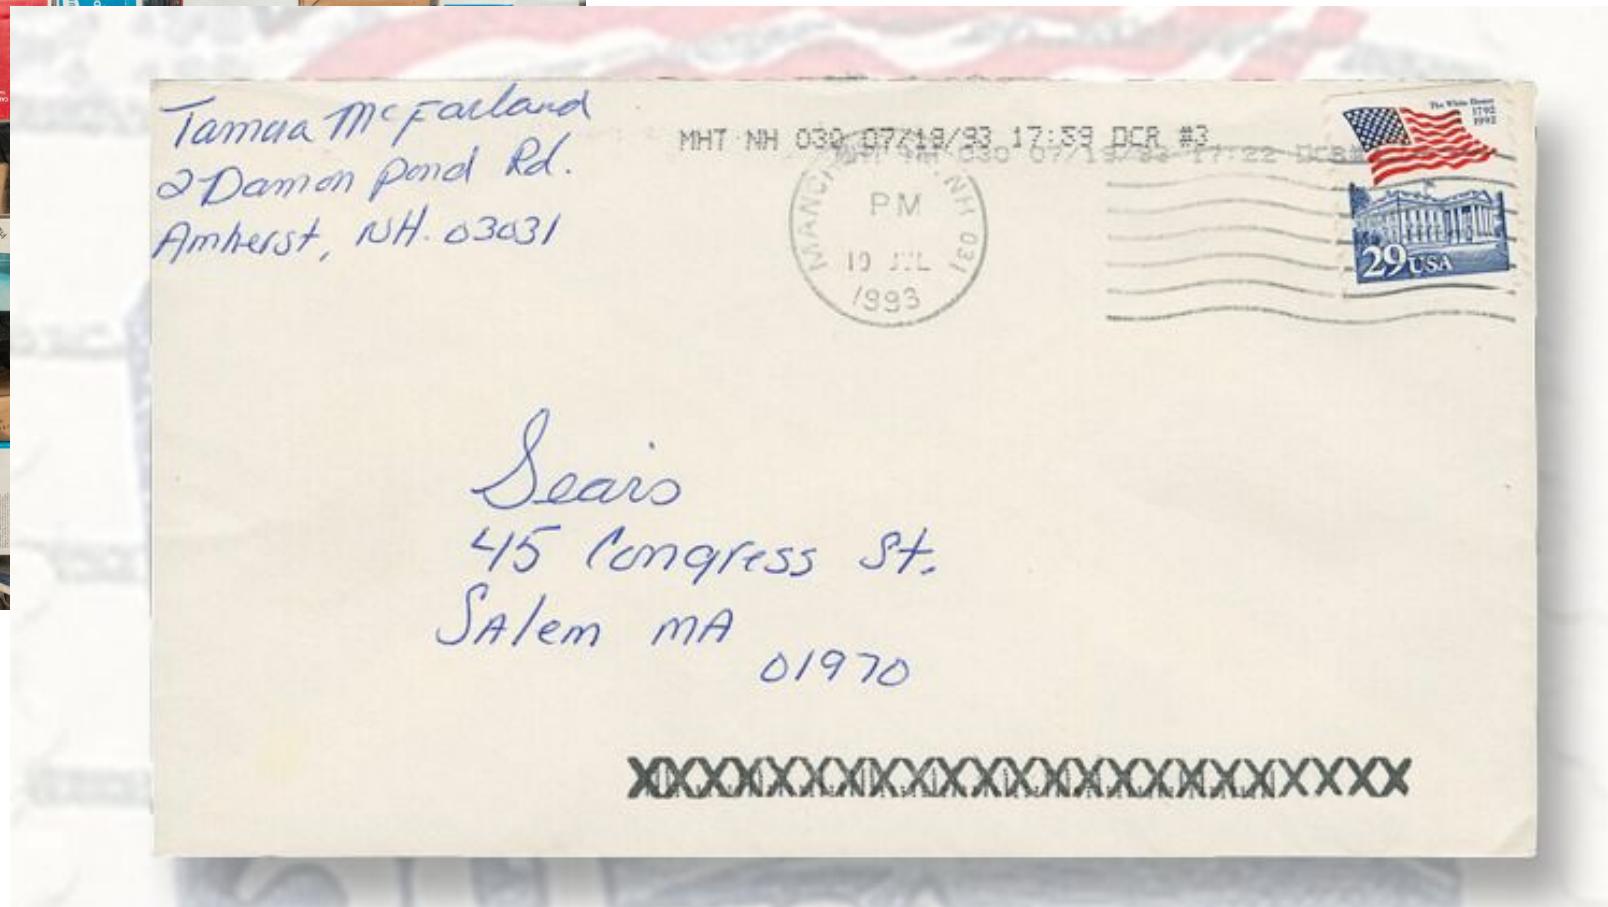


Deep L-Layers



Higher layers, higher levels of abstraction

AI – ML – Deep Learning – The US Postal Challenge



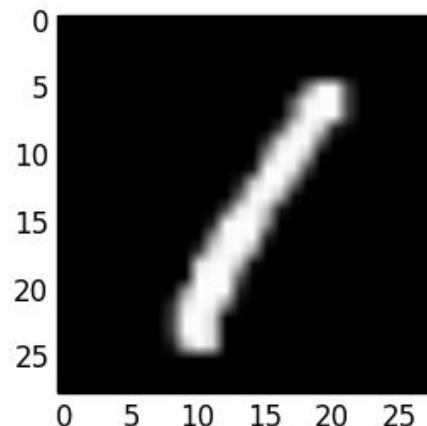
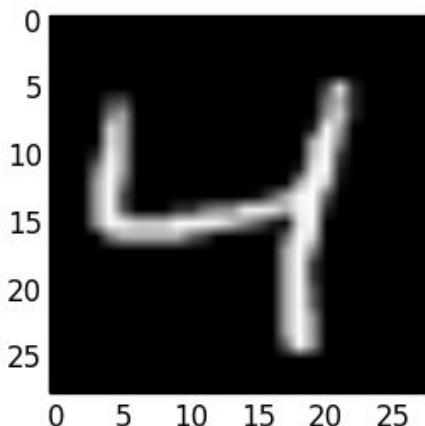
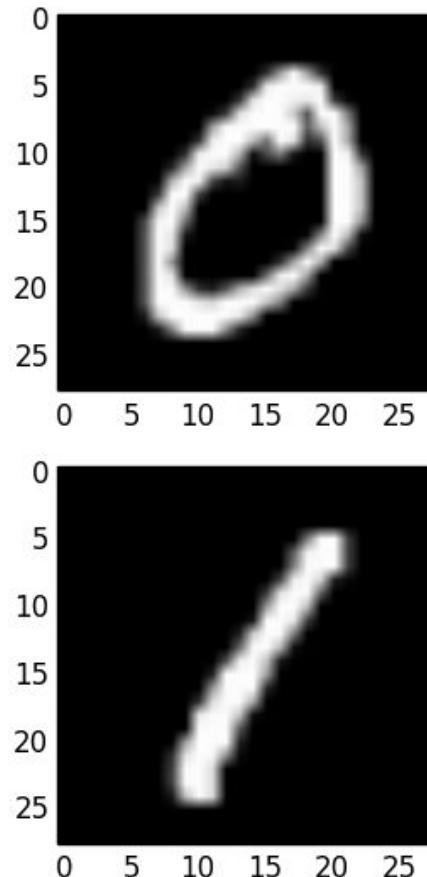
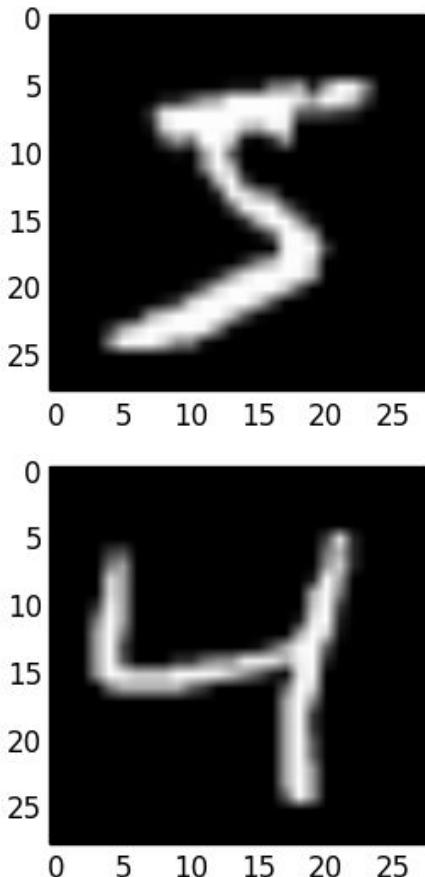
ANNs: The first challenge: MNIST

The MNIST database of handwritten digits has a training set of 60,000 examples, and a test set of 10,000 examples.

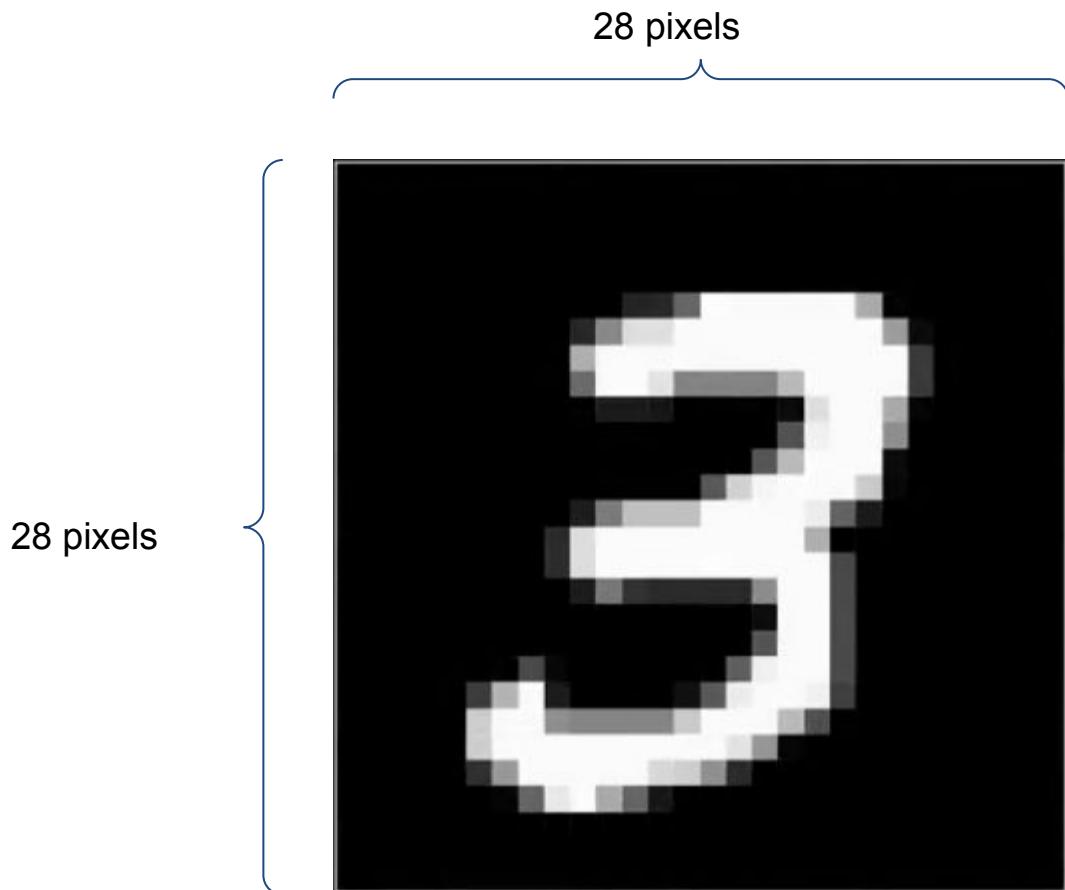
It is a very well-known benchmark for machine learning techniques

Each image is formed by 28x28 grayscale pixels

ANNs: The first challenge: MNIST



Deep L-Layers

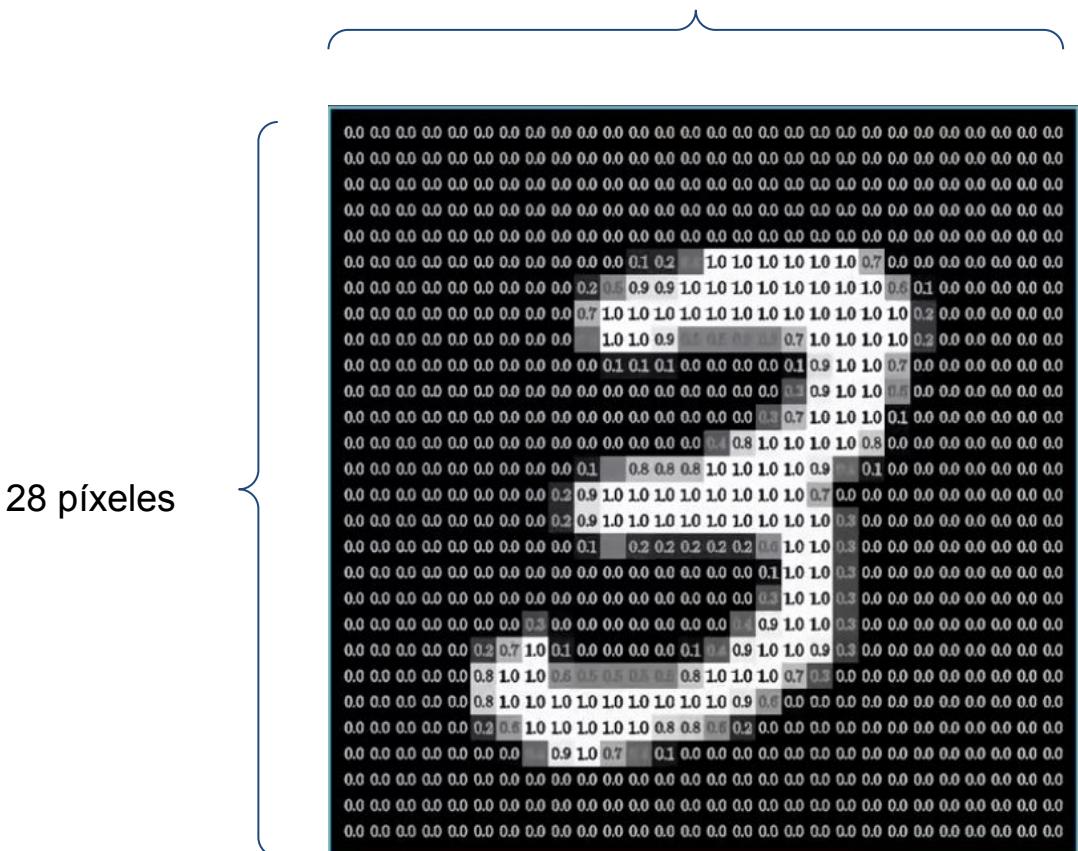


Let's get started with a real task:

- MNIST
- This is a 3... or a pixel array, depending on how you look at it
- Easy task for the brain but... How would you code a script to recognize this three?

Deep L-Layers

28 píxeles



Let's get started with a real task:

- MNIST
 - This is a 3... or a pixel array, depending on how you look at it
 - Easy task for the brain but... How would you code a script to recognize this three?

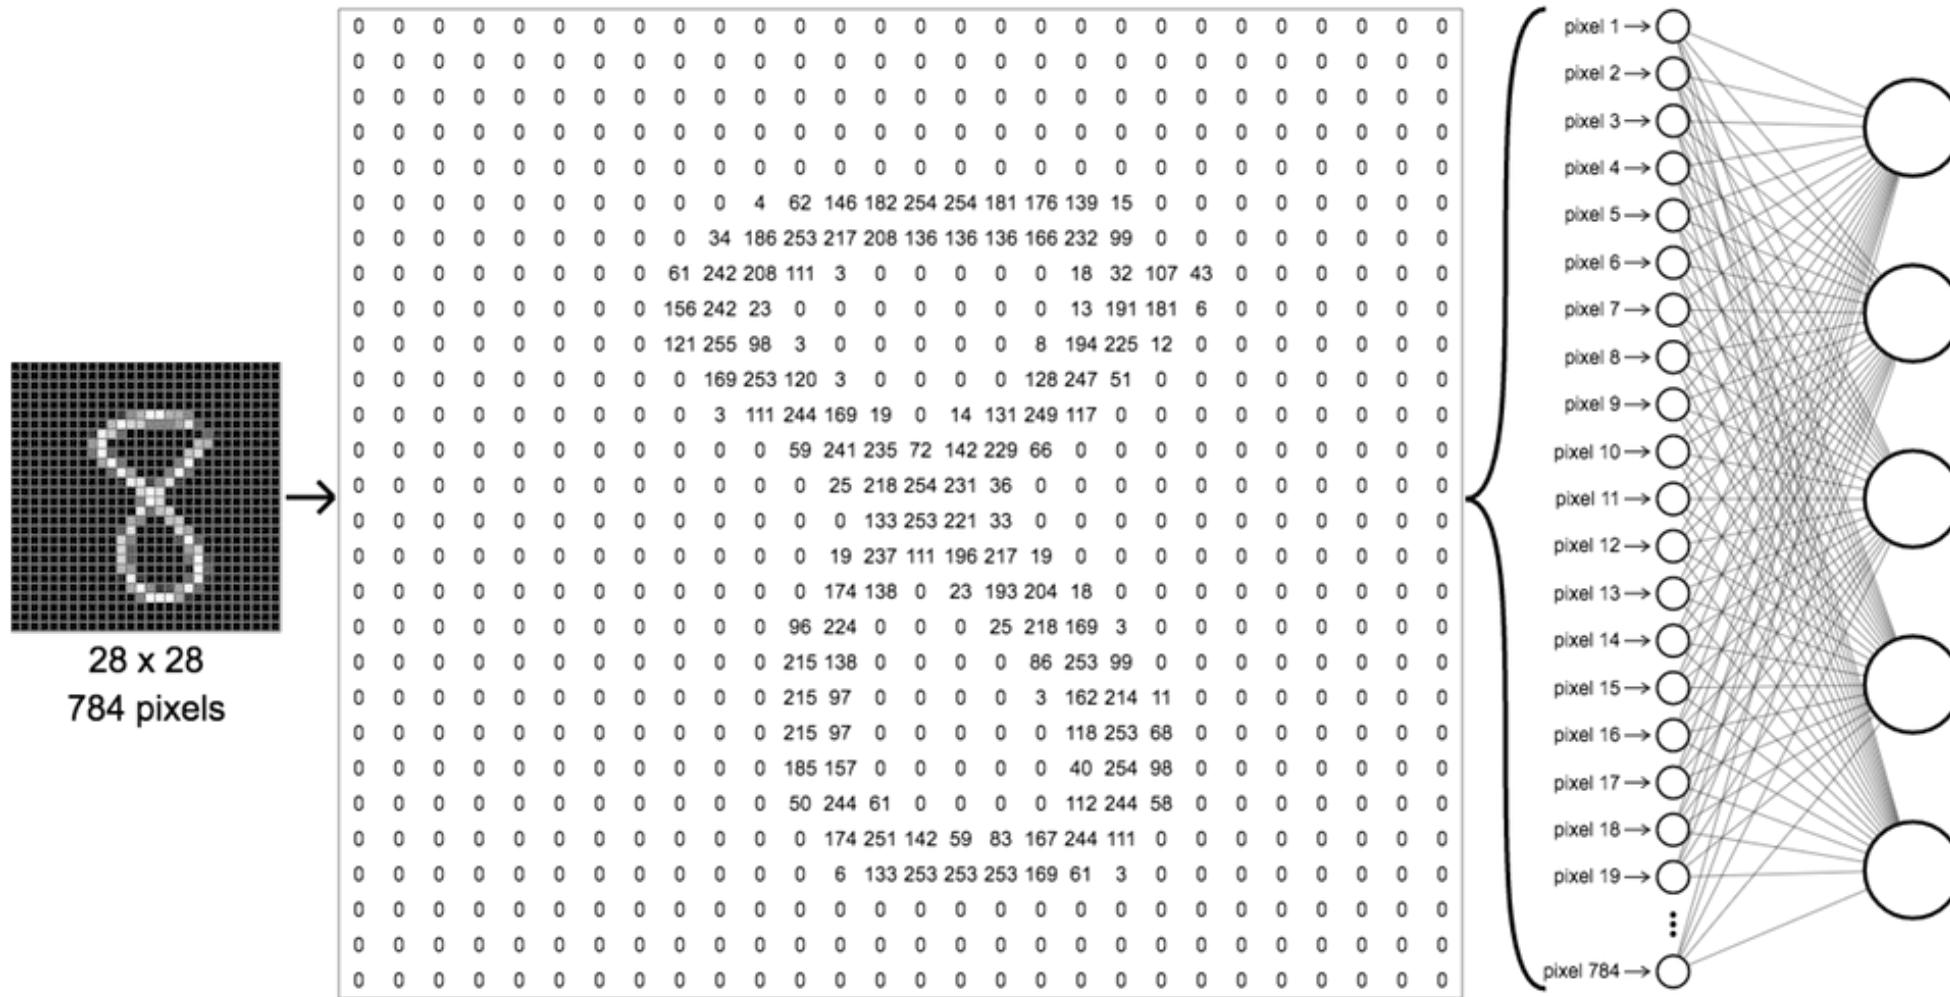
Deep L-Layers



Deep L-Layers

- How can we tackle this problem with a DNN?
- Let us begin here:
 - Standard DNN
 - Vanilla form, no rings nor bells
 - AKA Multilayer Perceptron

Artificial Neural Networks: Input MNIST



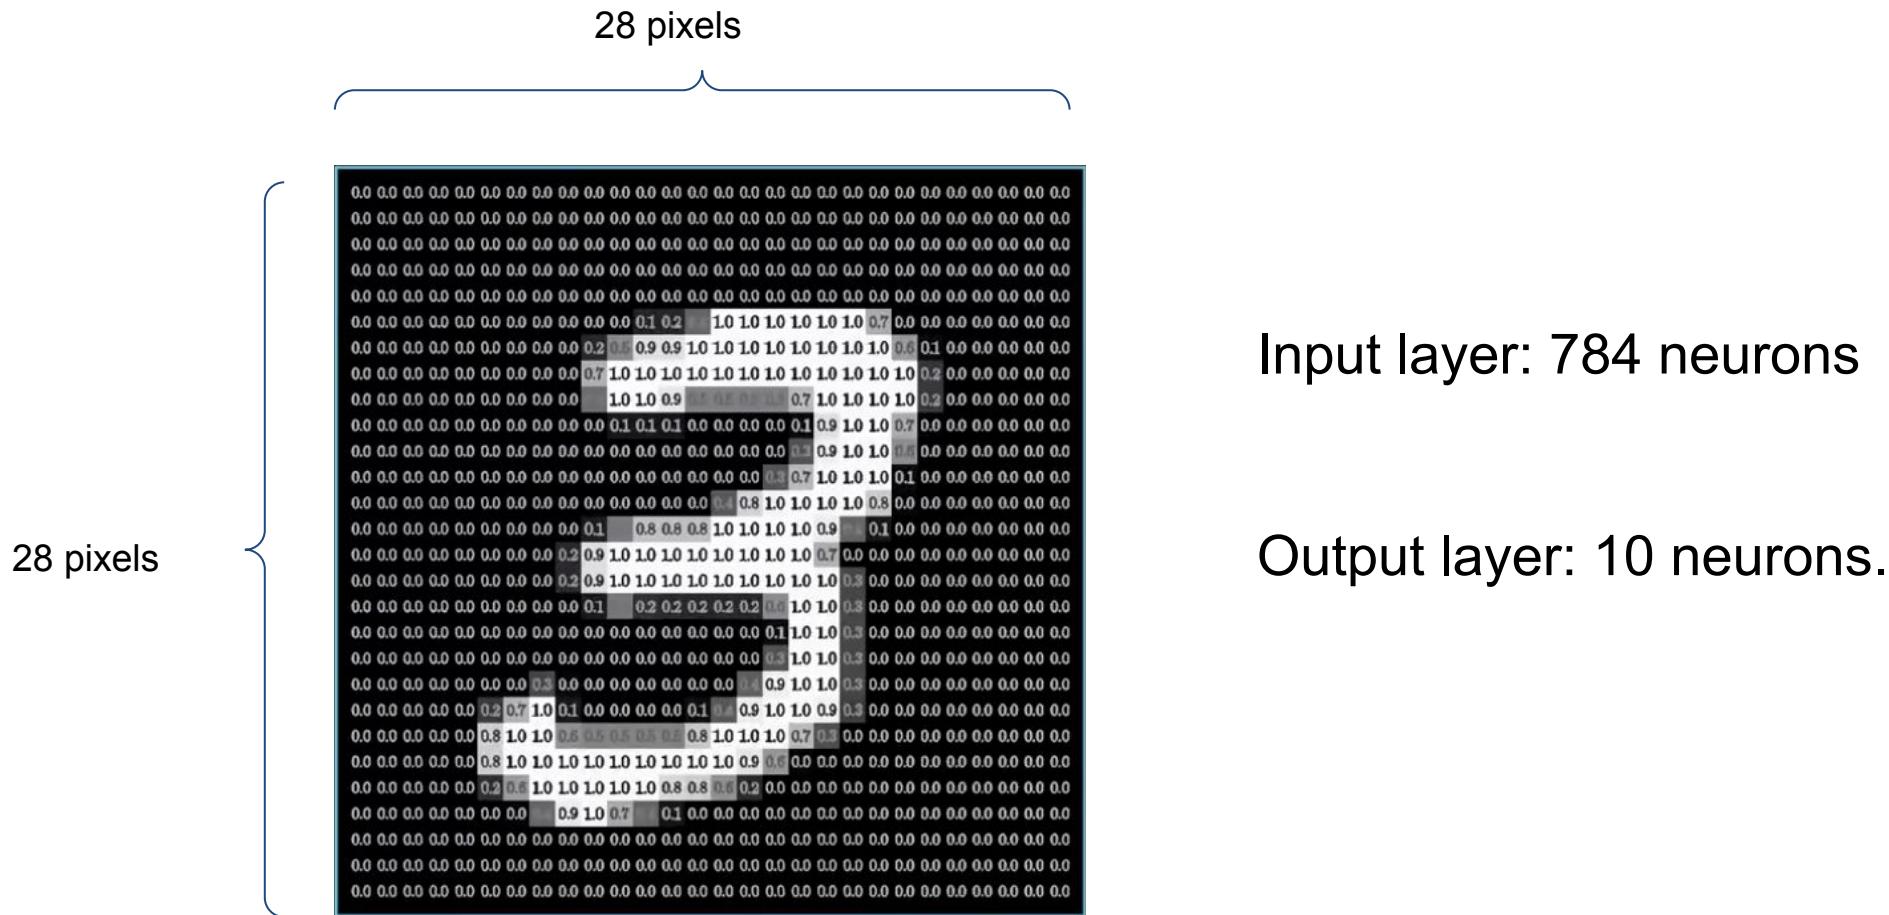
Deep L-Layers

- Step by step... Neural Networks
 - What is a neuron? So far, let's think of neurons as a container with some number.

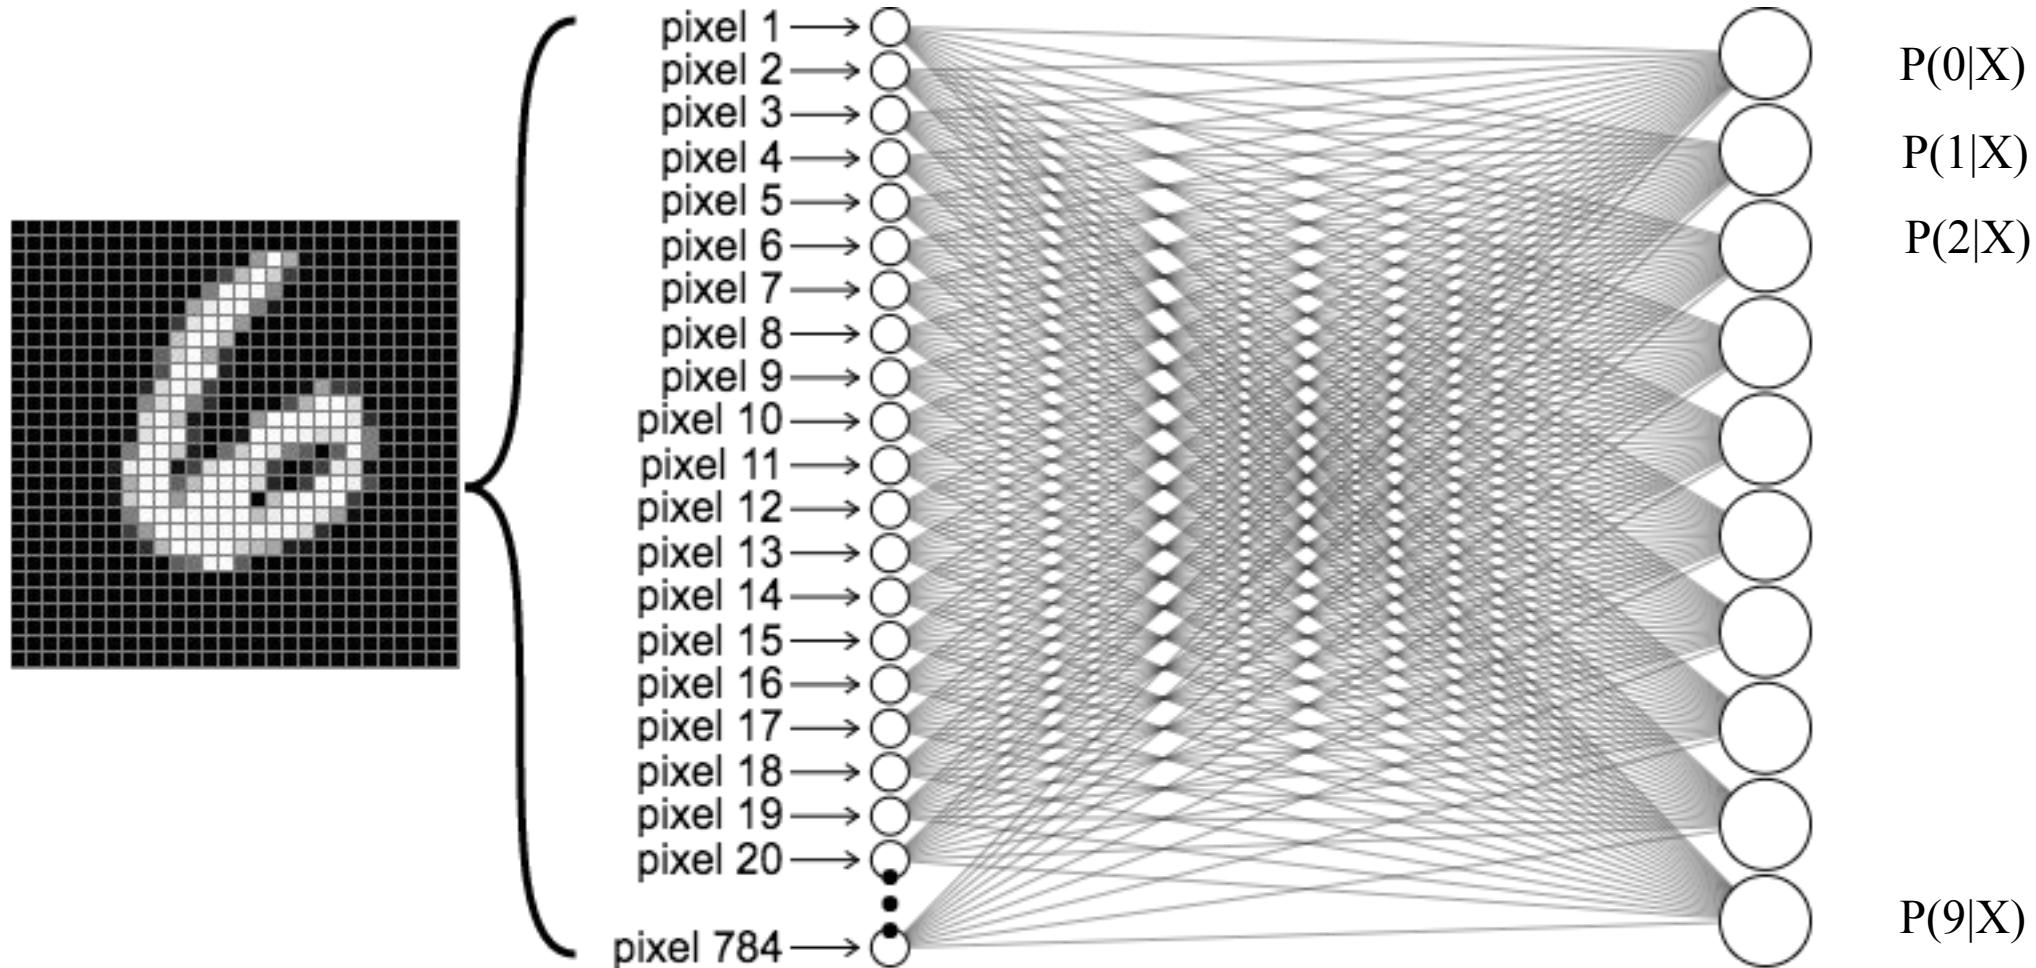


- How are they connected: Let's discover it!

Deep L-Layers



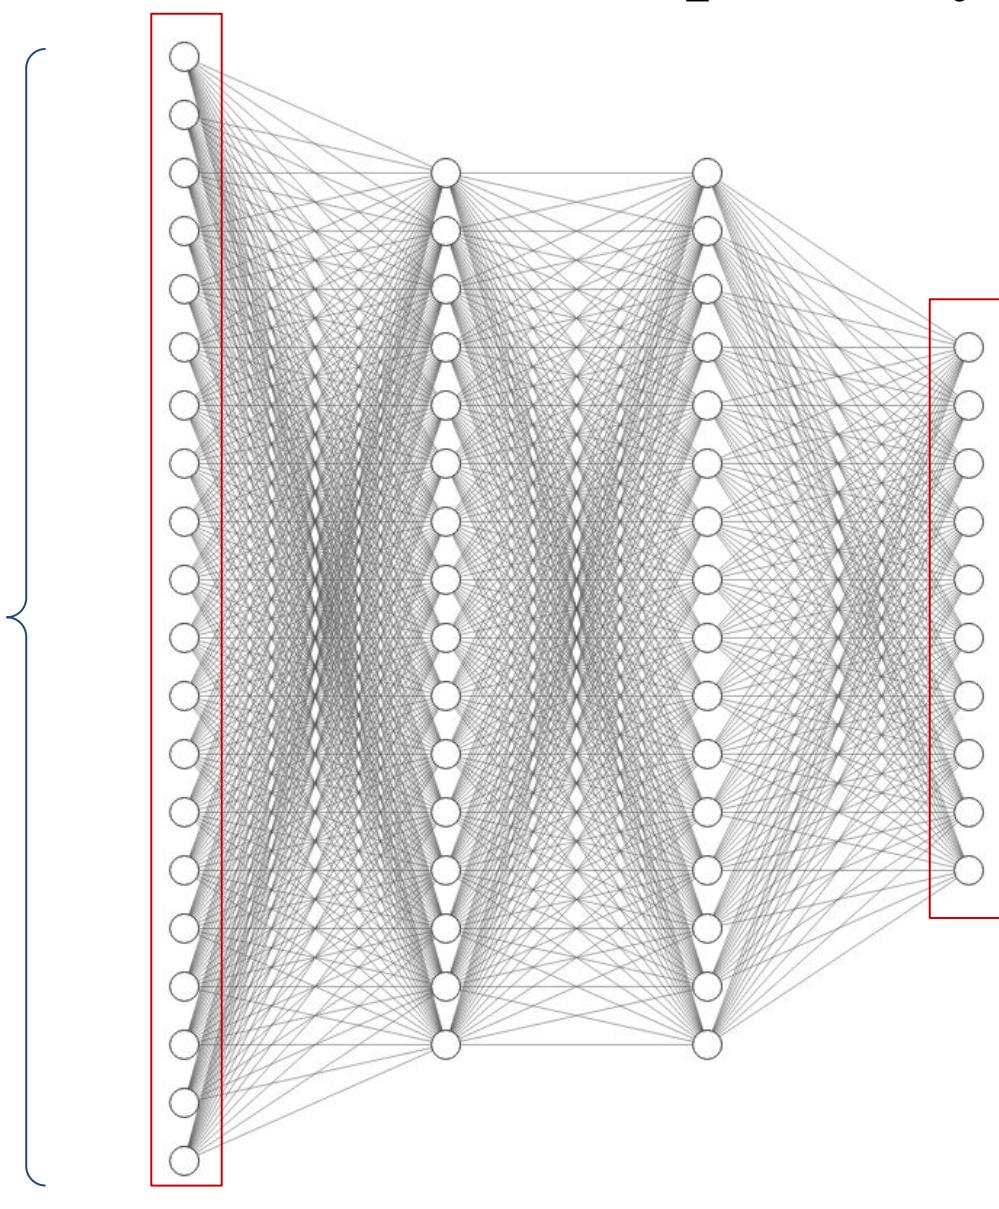
Artificial Neural Networks: Output MNIST



Deep L-Layers



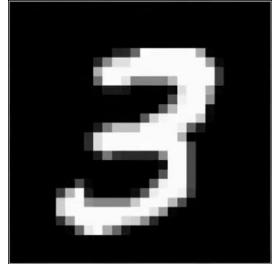
784 píxeles



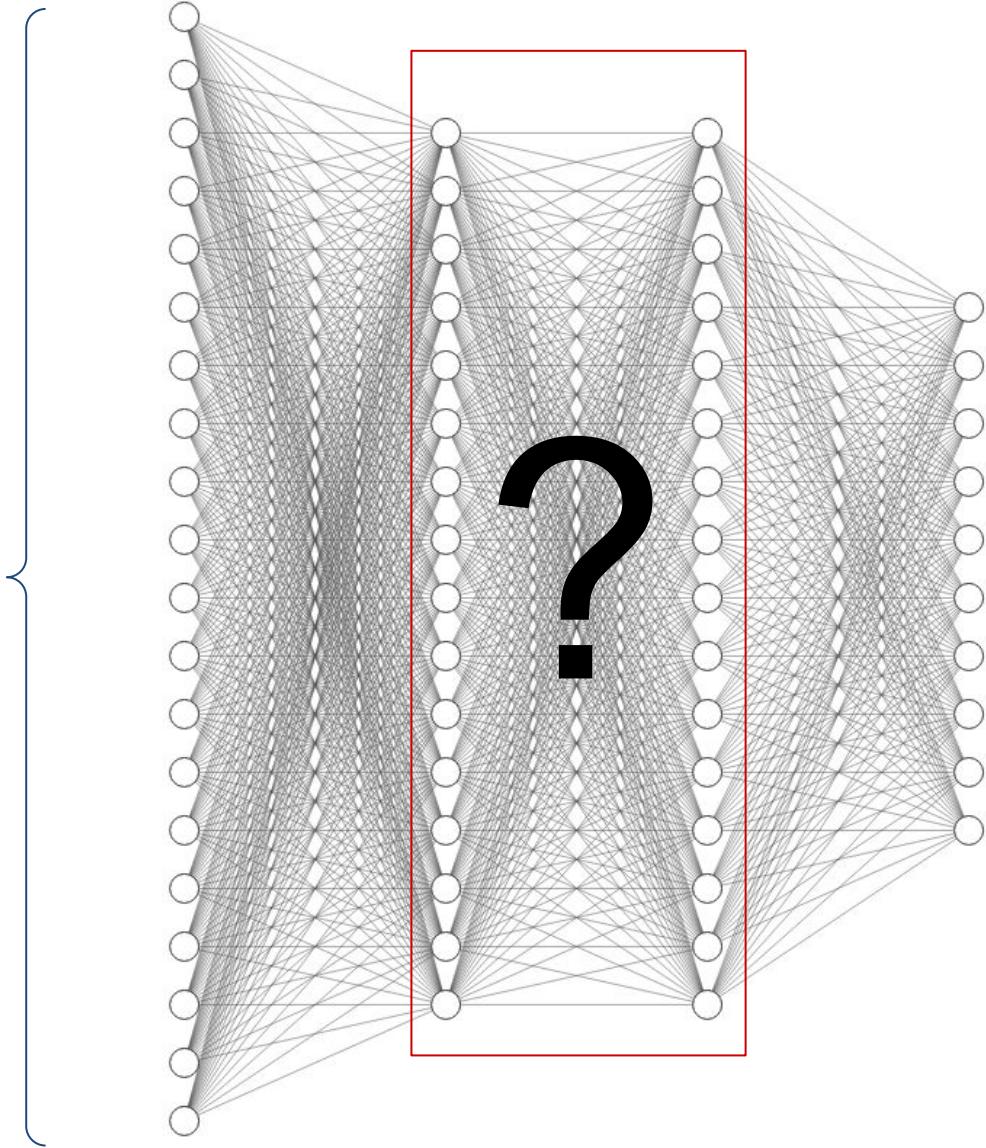
Input layer: 784 neurons

Output layer: 10 neurons.

Deep L-Layers



784 píxeles

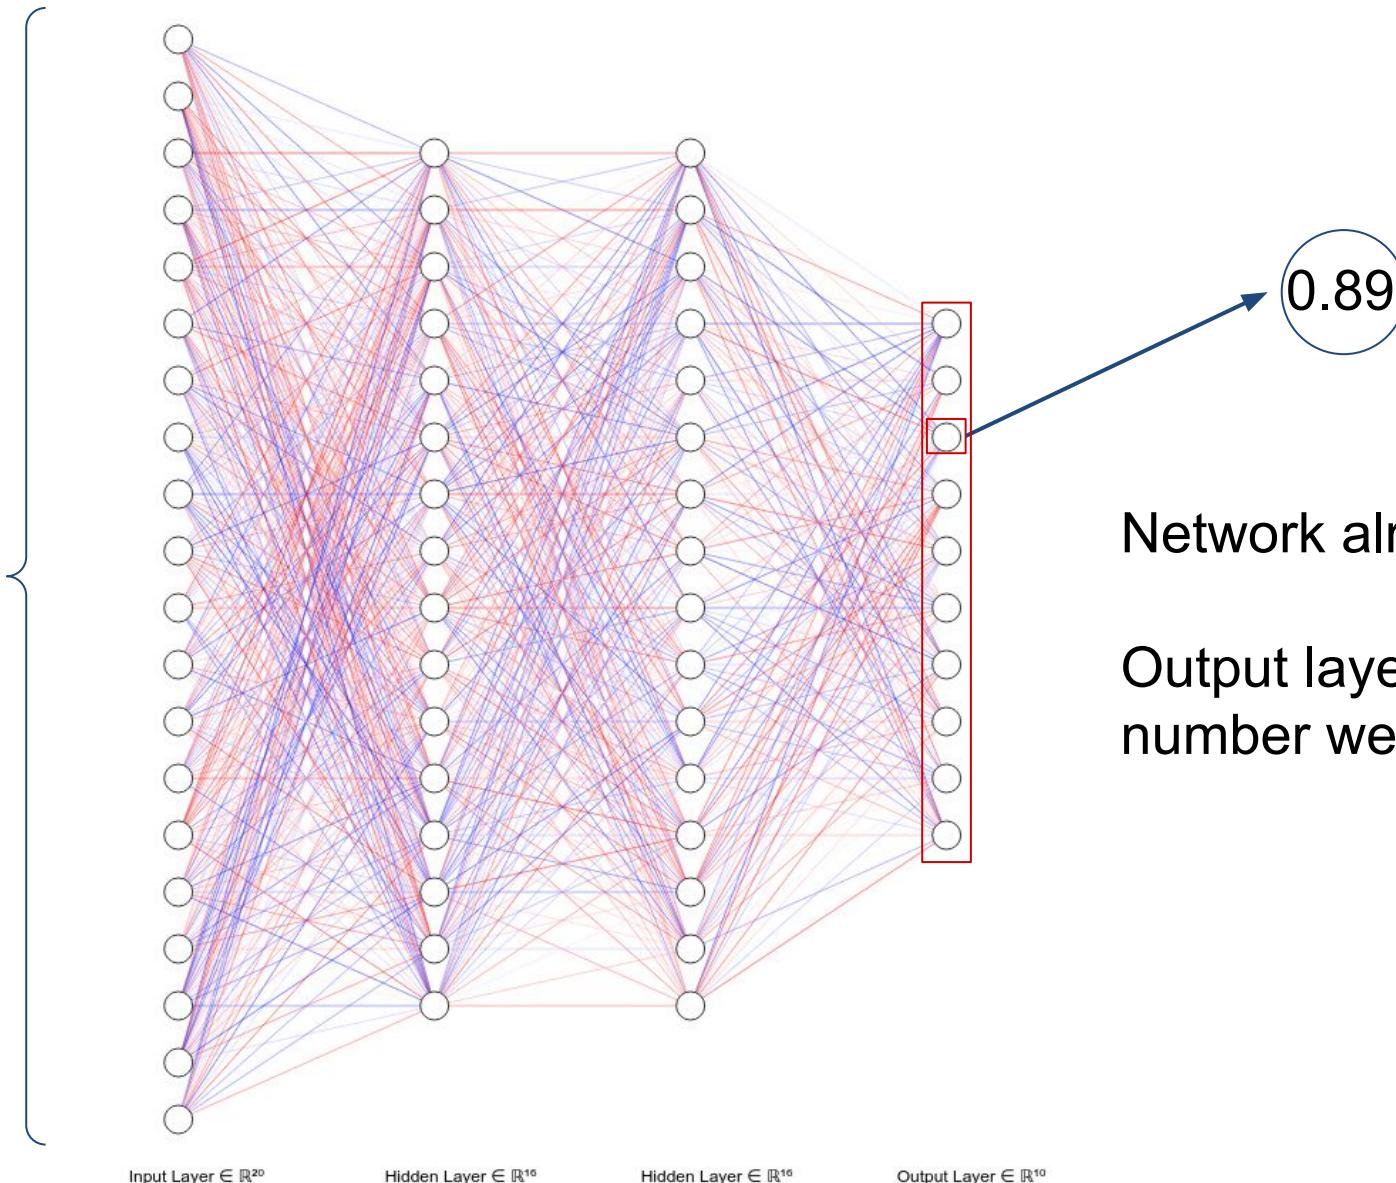


Hidden Layers: ¿?

Deep L-Layers



784 pixels

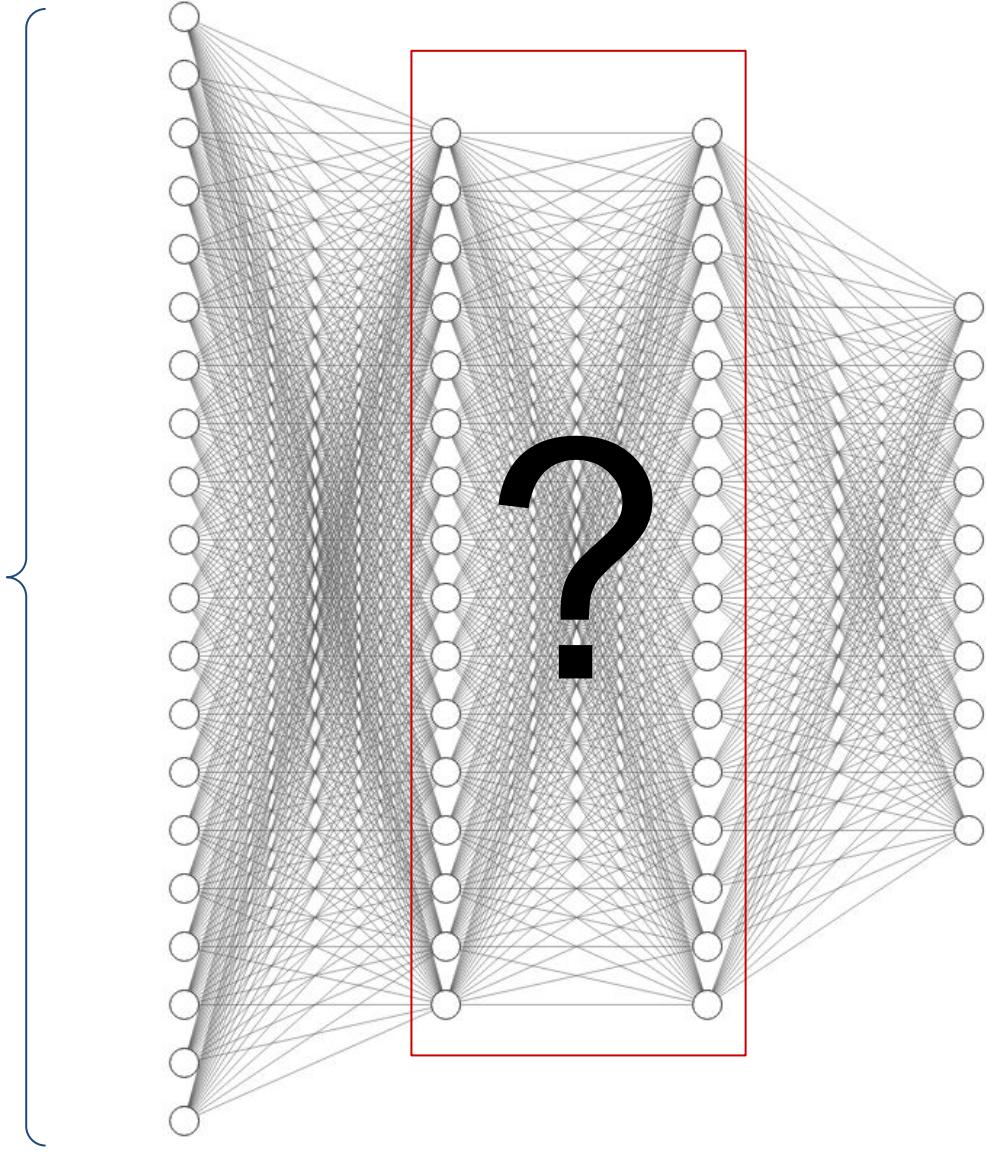


Network already trained.

Output layer tells us which number we have in the input



784 píxeles



Deep L-Layers

Remember:
¿Why do we use layers?

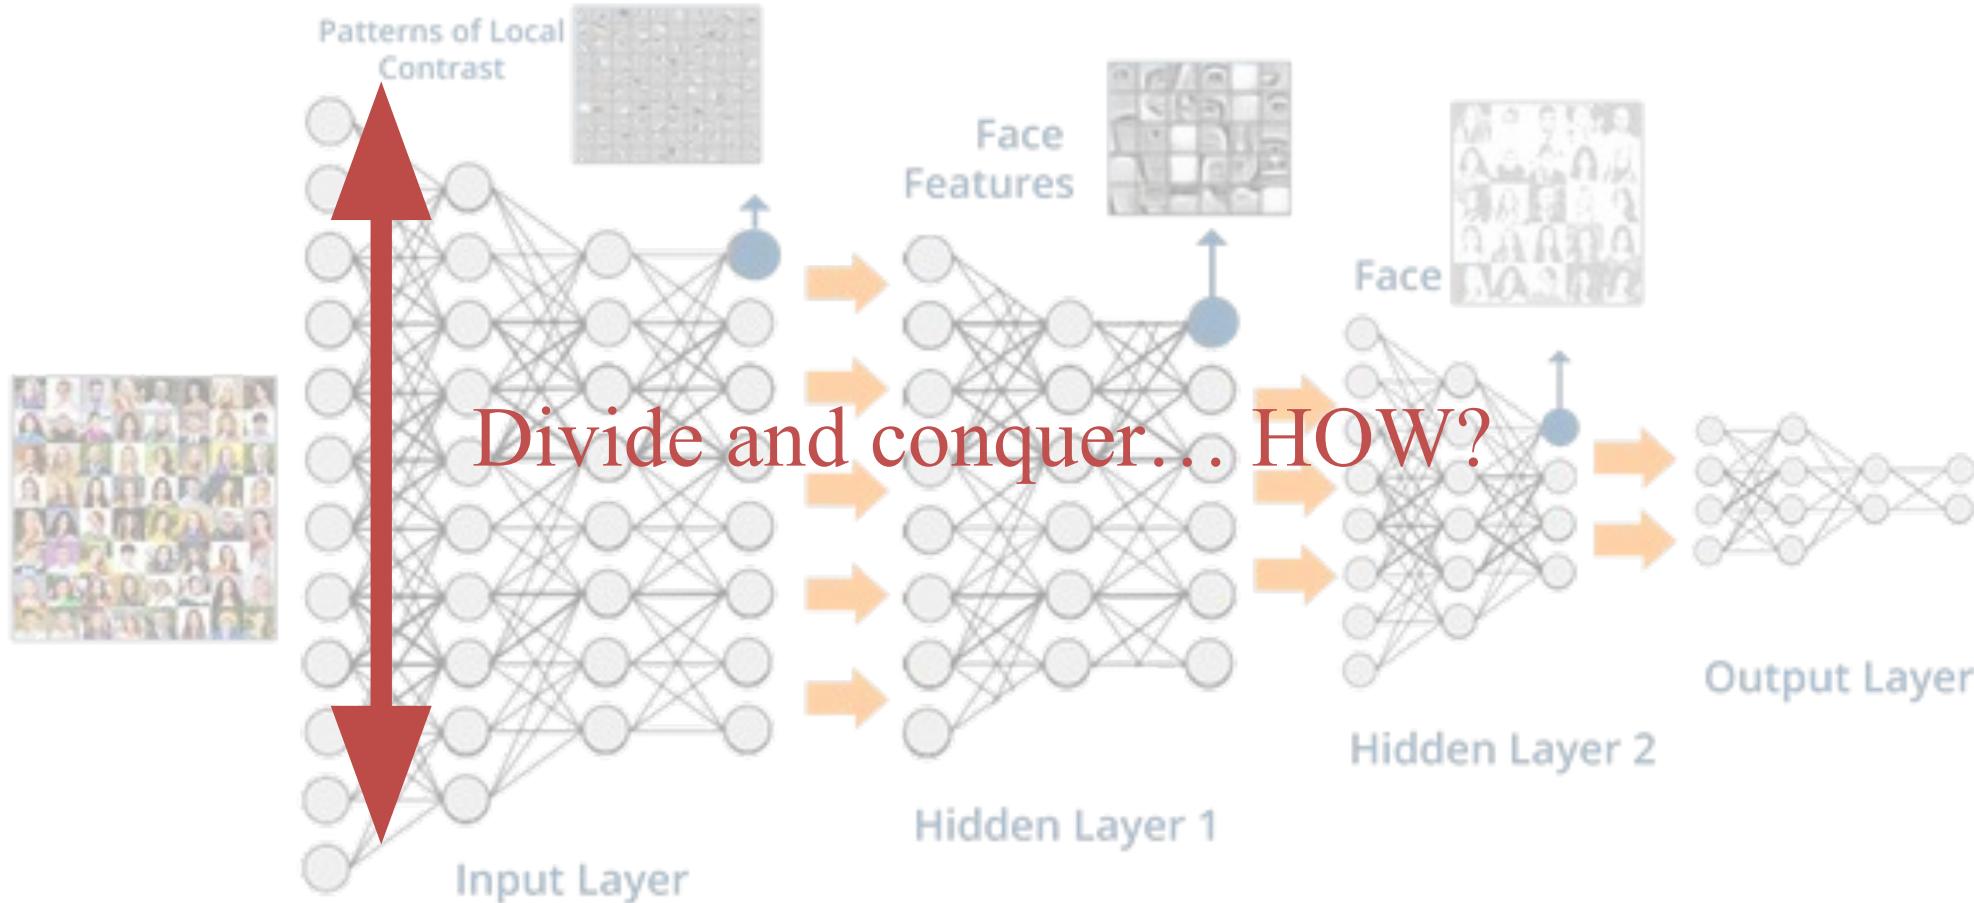
- XOR
- Circuit theory

Outline

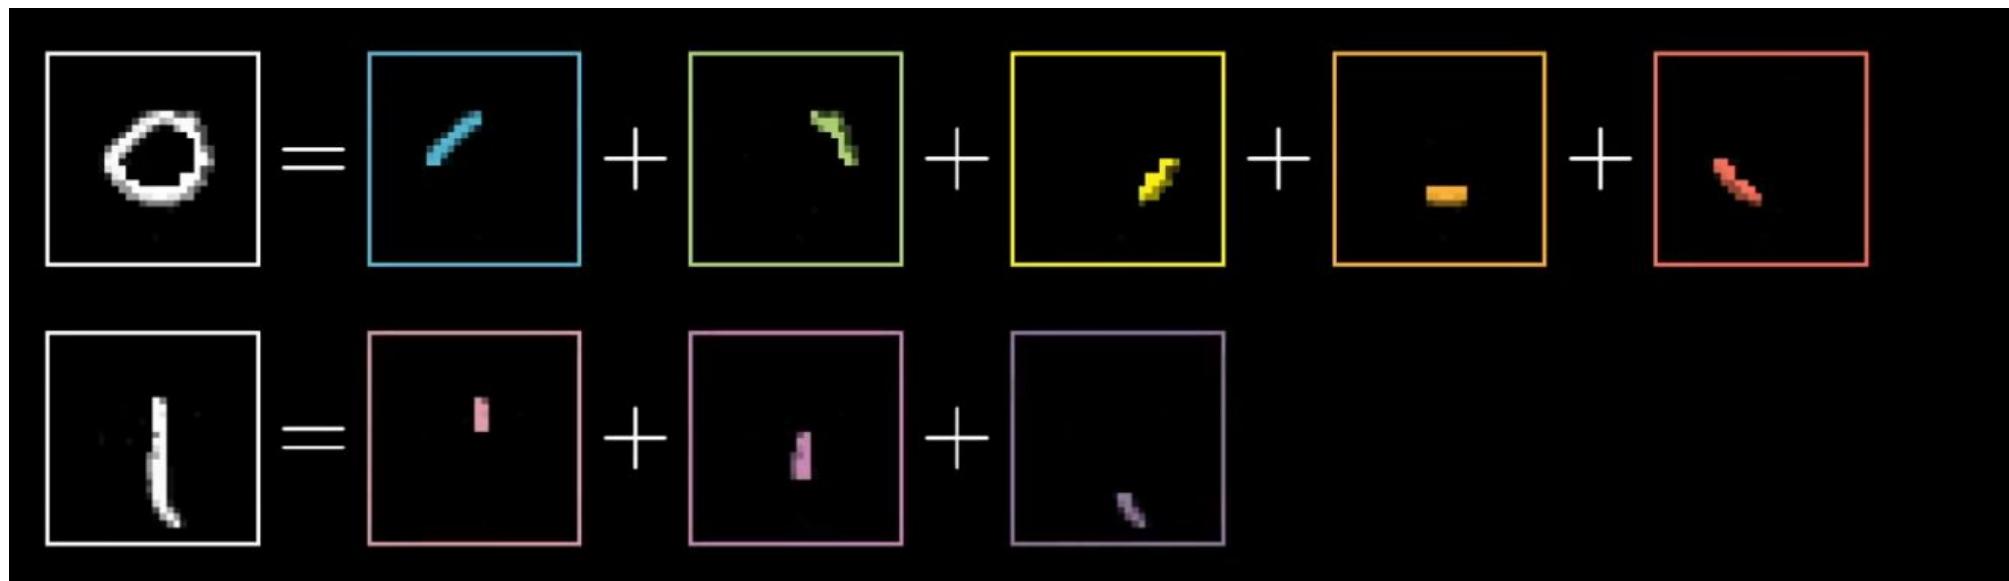
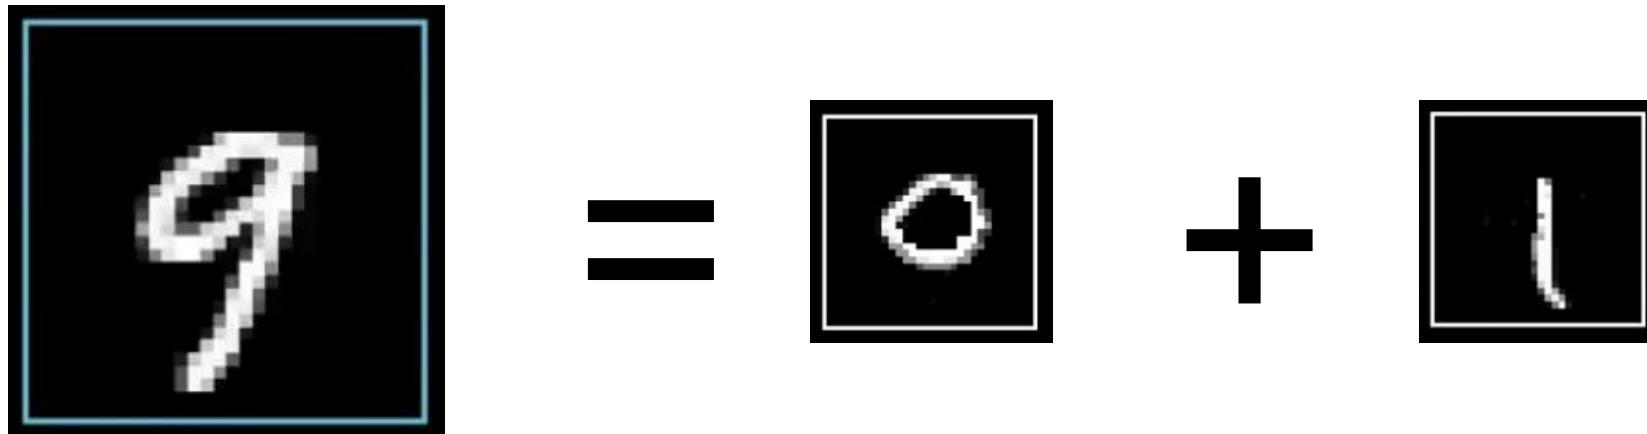


1. Recap
2. Deep L-Layers
3. Divide and Conquer
4. The Neuron: Pattern Recognition
5. DNNs

Divide and Conquer

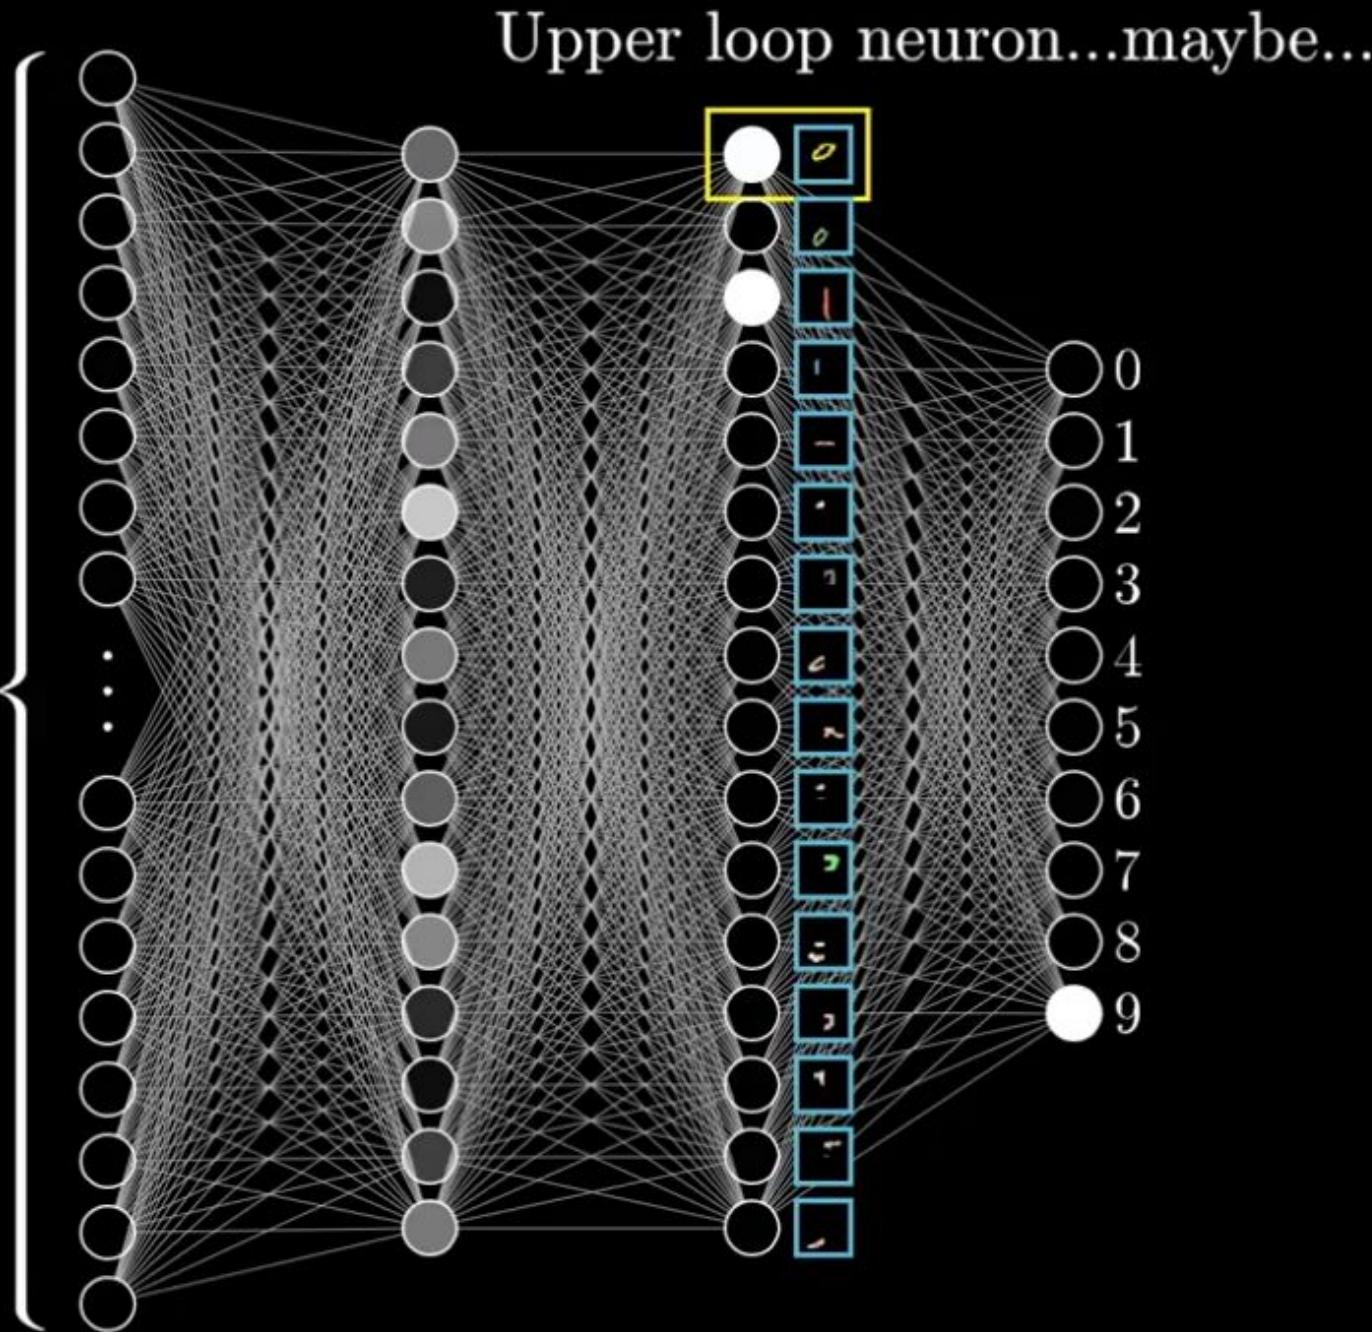


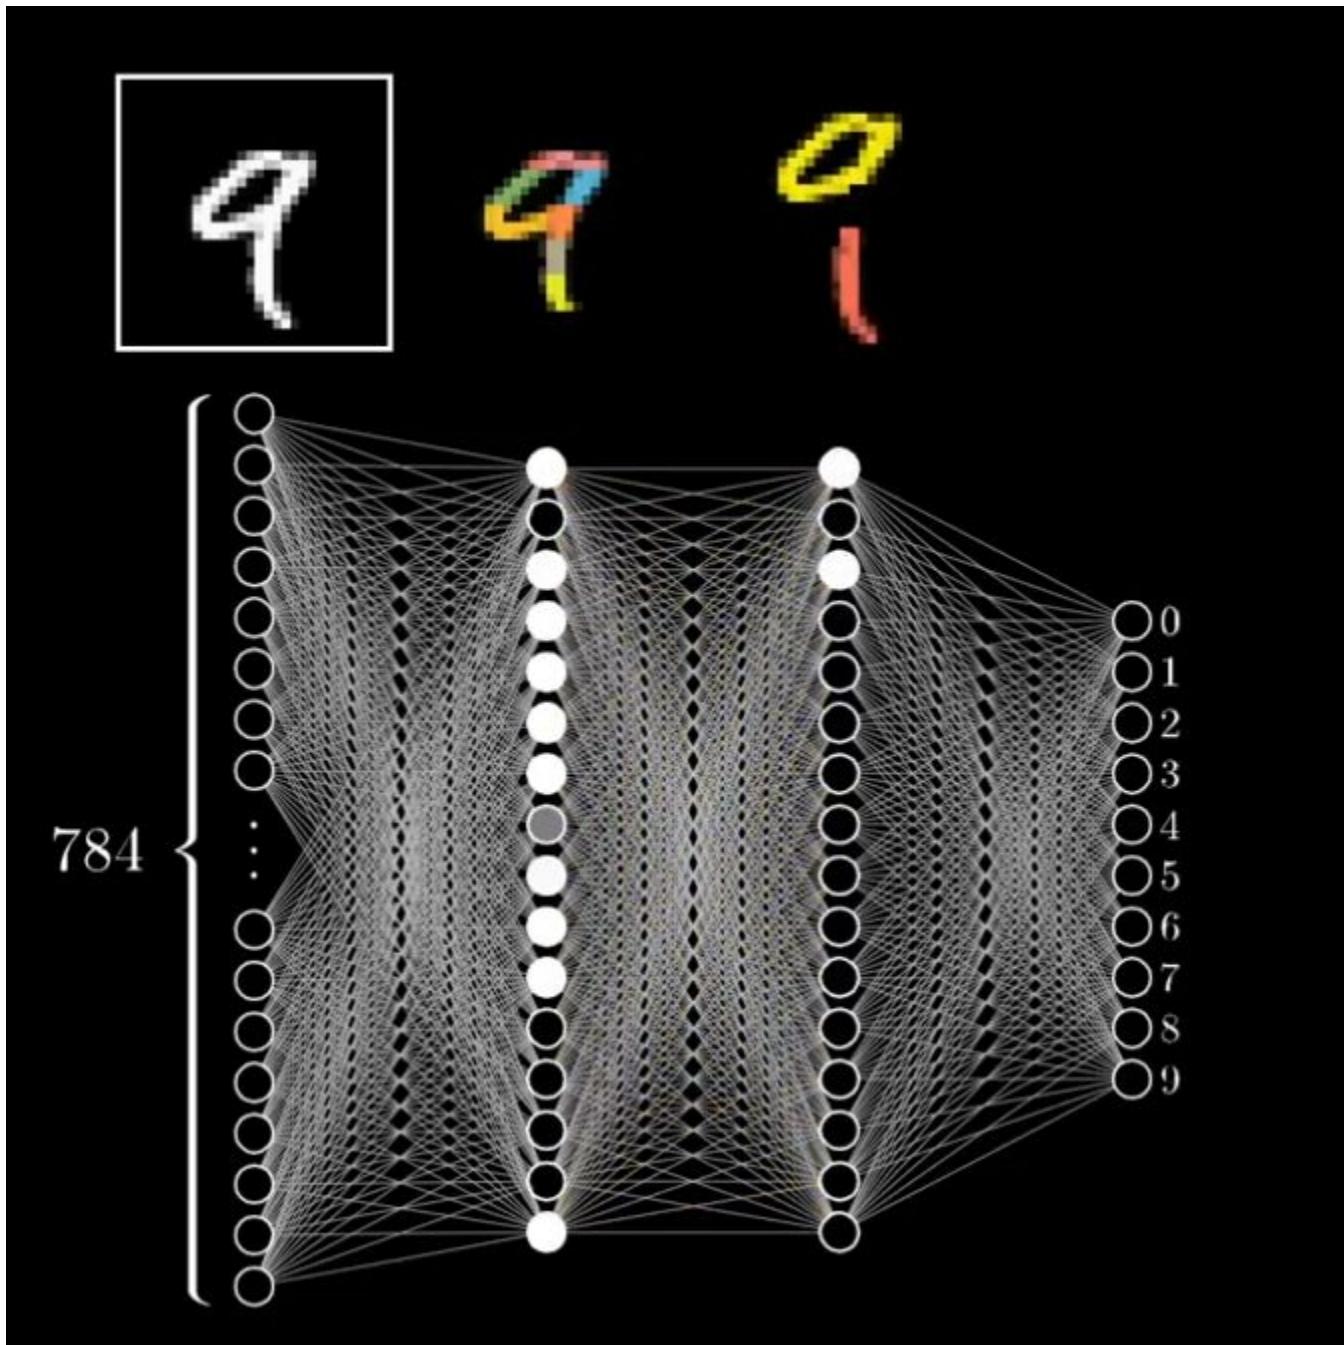
Divide and Conquer





784





Divide and Conquer

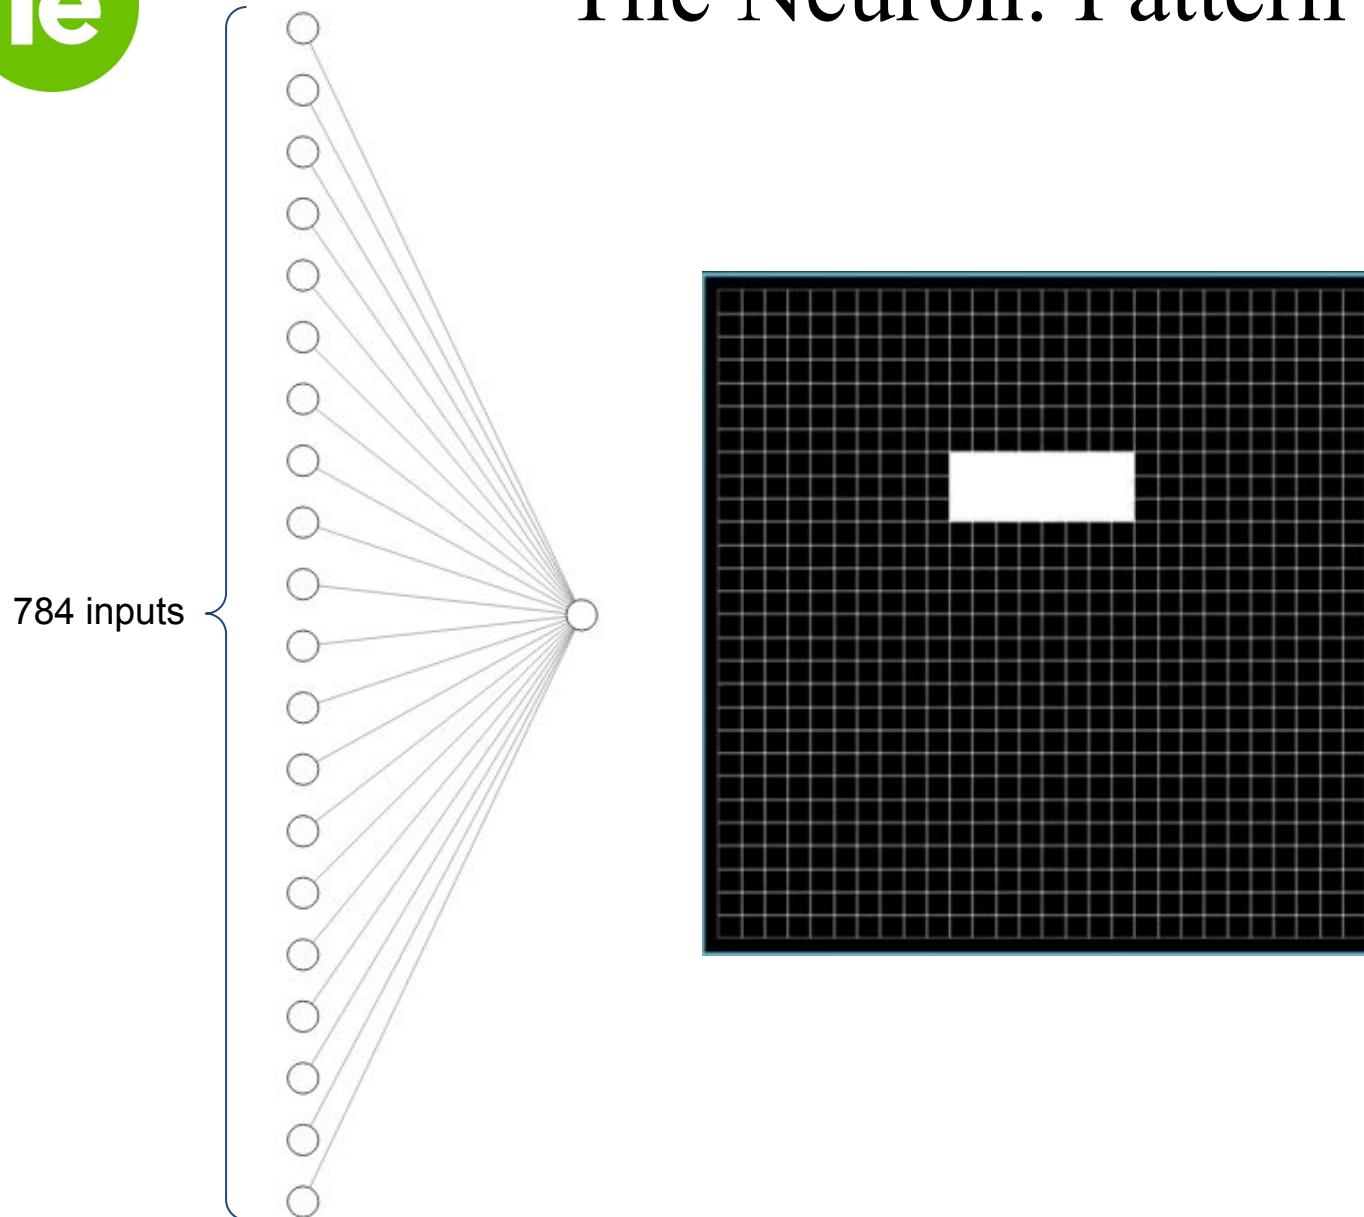


Outline



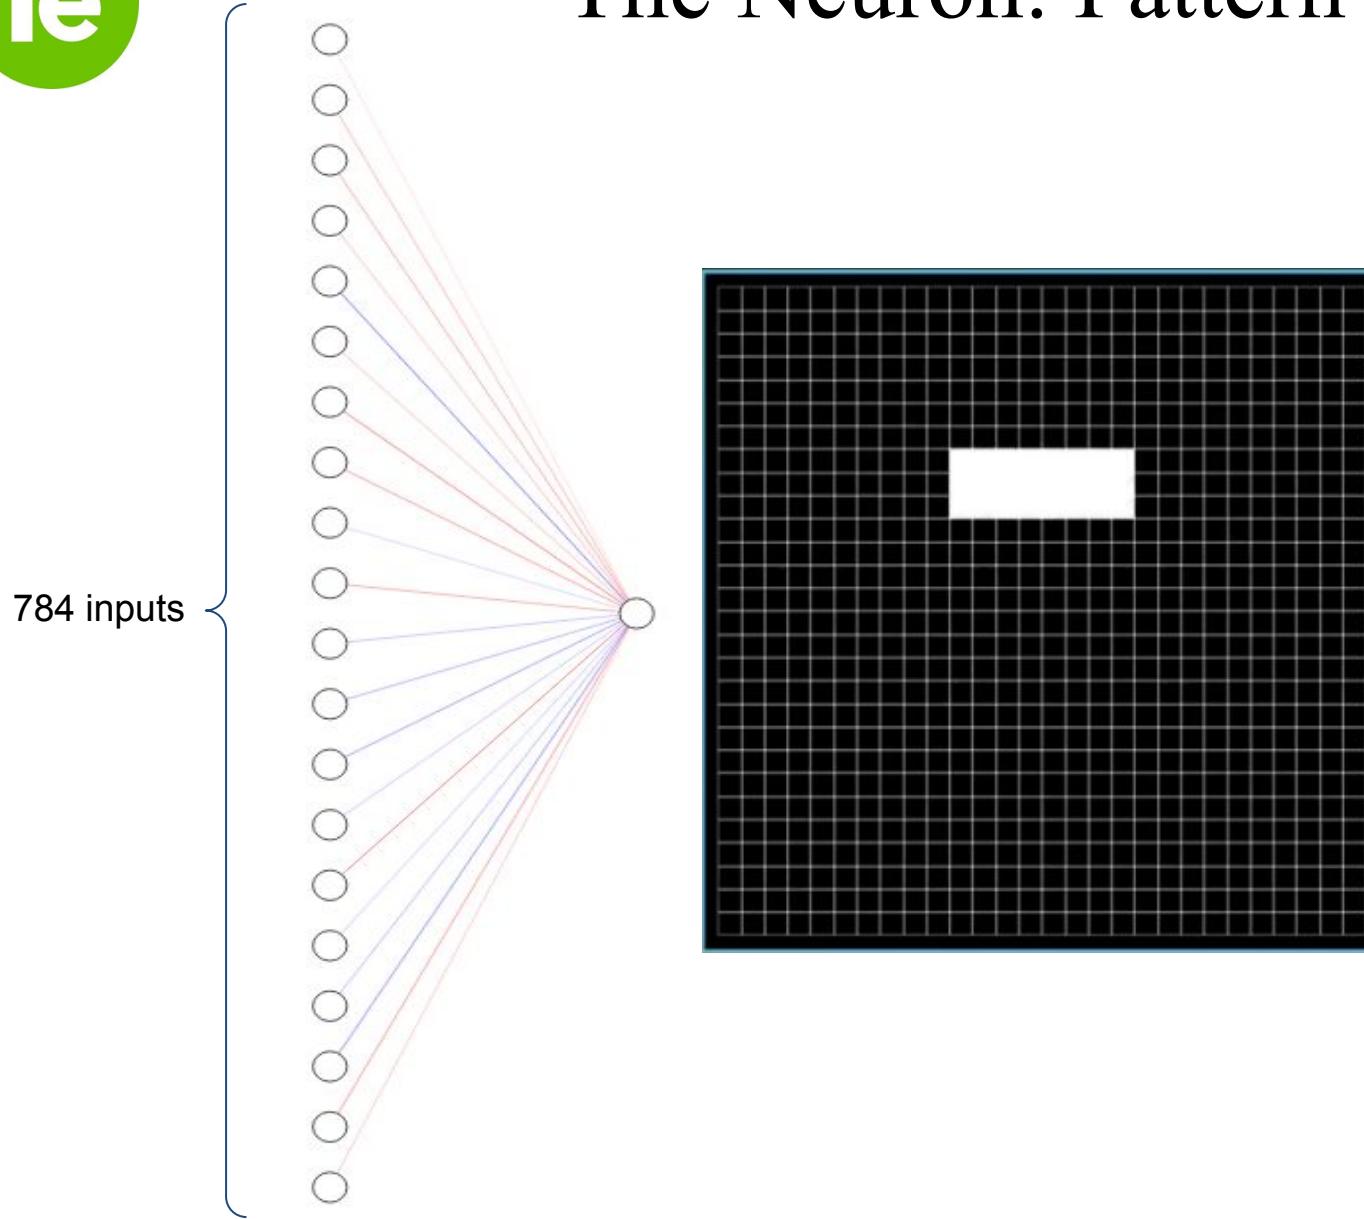
1. Recap
2. Deep L-Layers
3. Divide and Conquer
4. The Neuron: Pattern Recognition
5. DNNs

The Neuron: Pattern Recognition



How can a neuron detect
that pattern?

The Neuron: Pattern Recognition

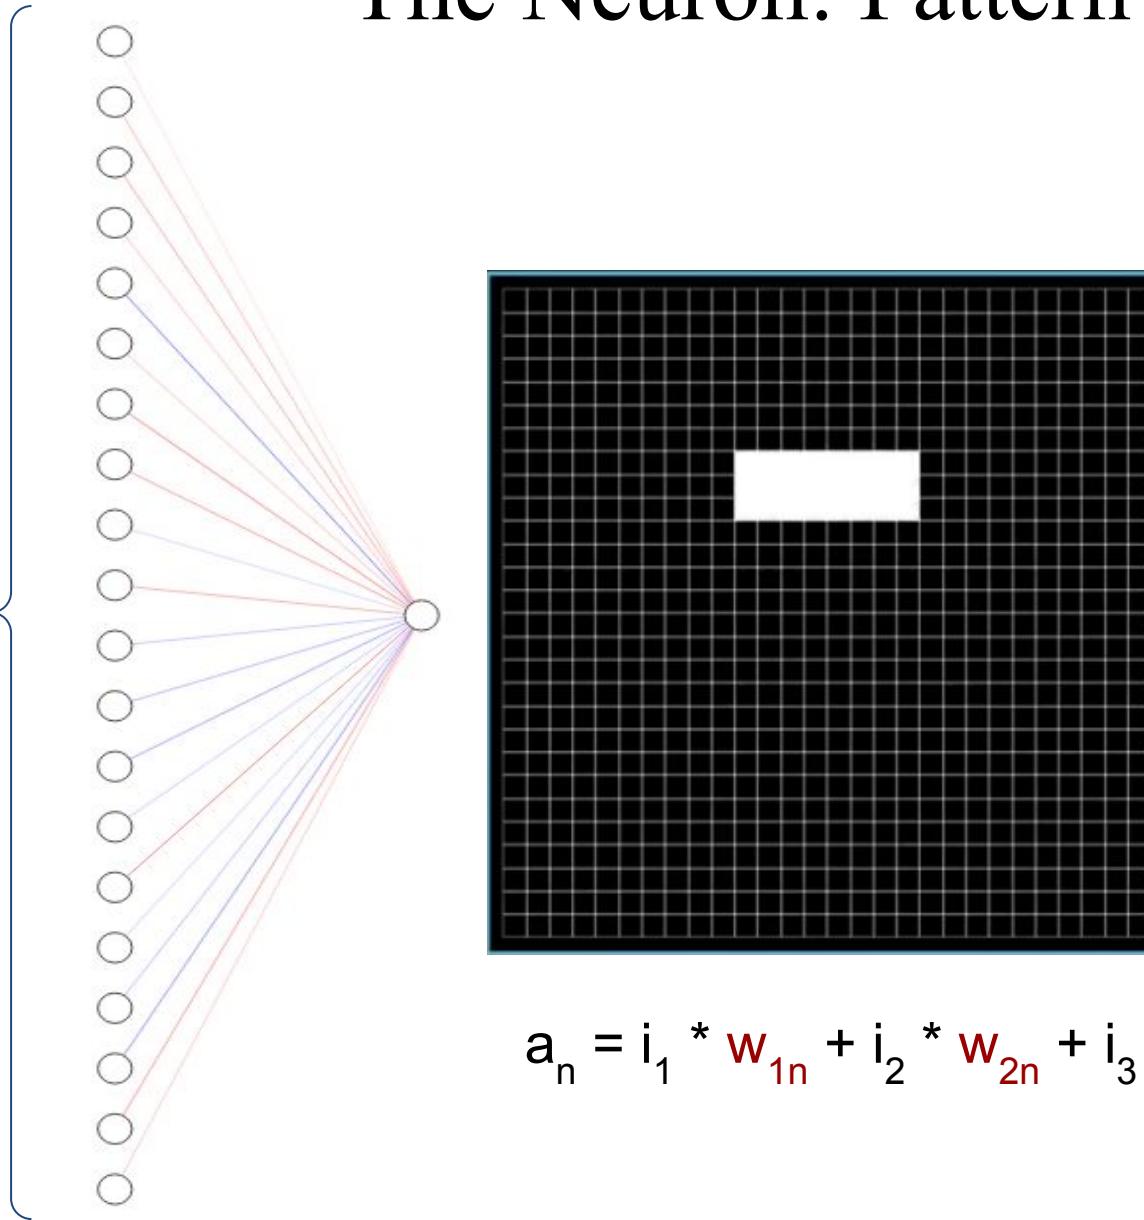


How can a neuron detect
that pattern?

- Every pixel is an input
- What's the output?

The Neuron: Pattern Recognition

784 inputs



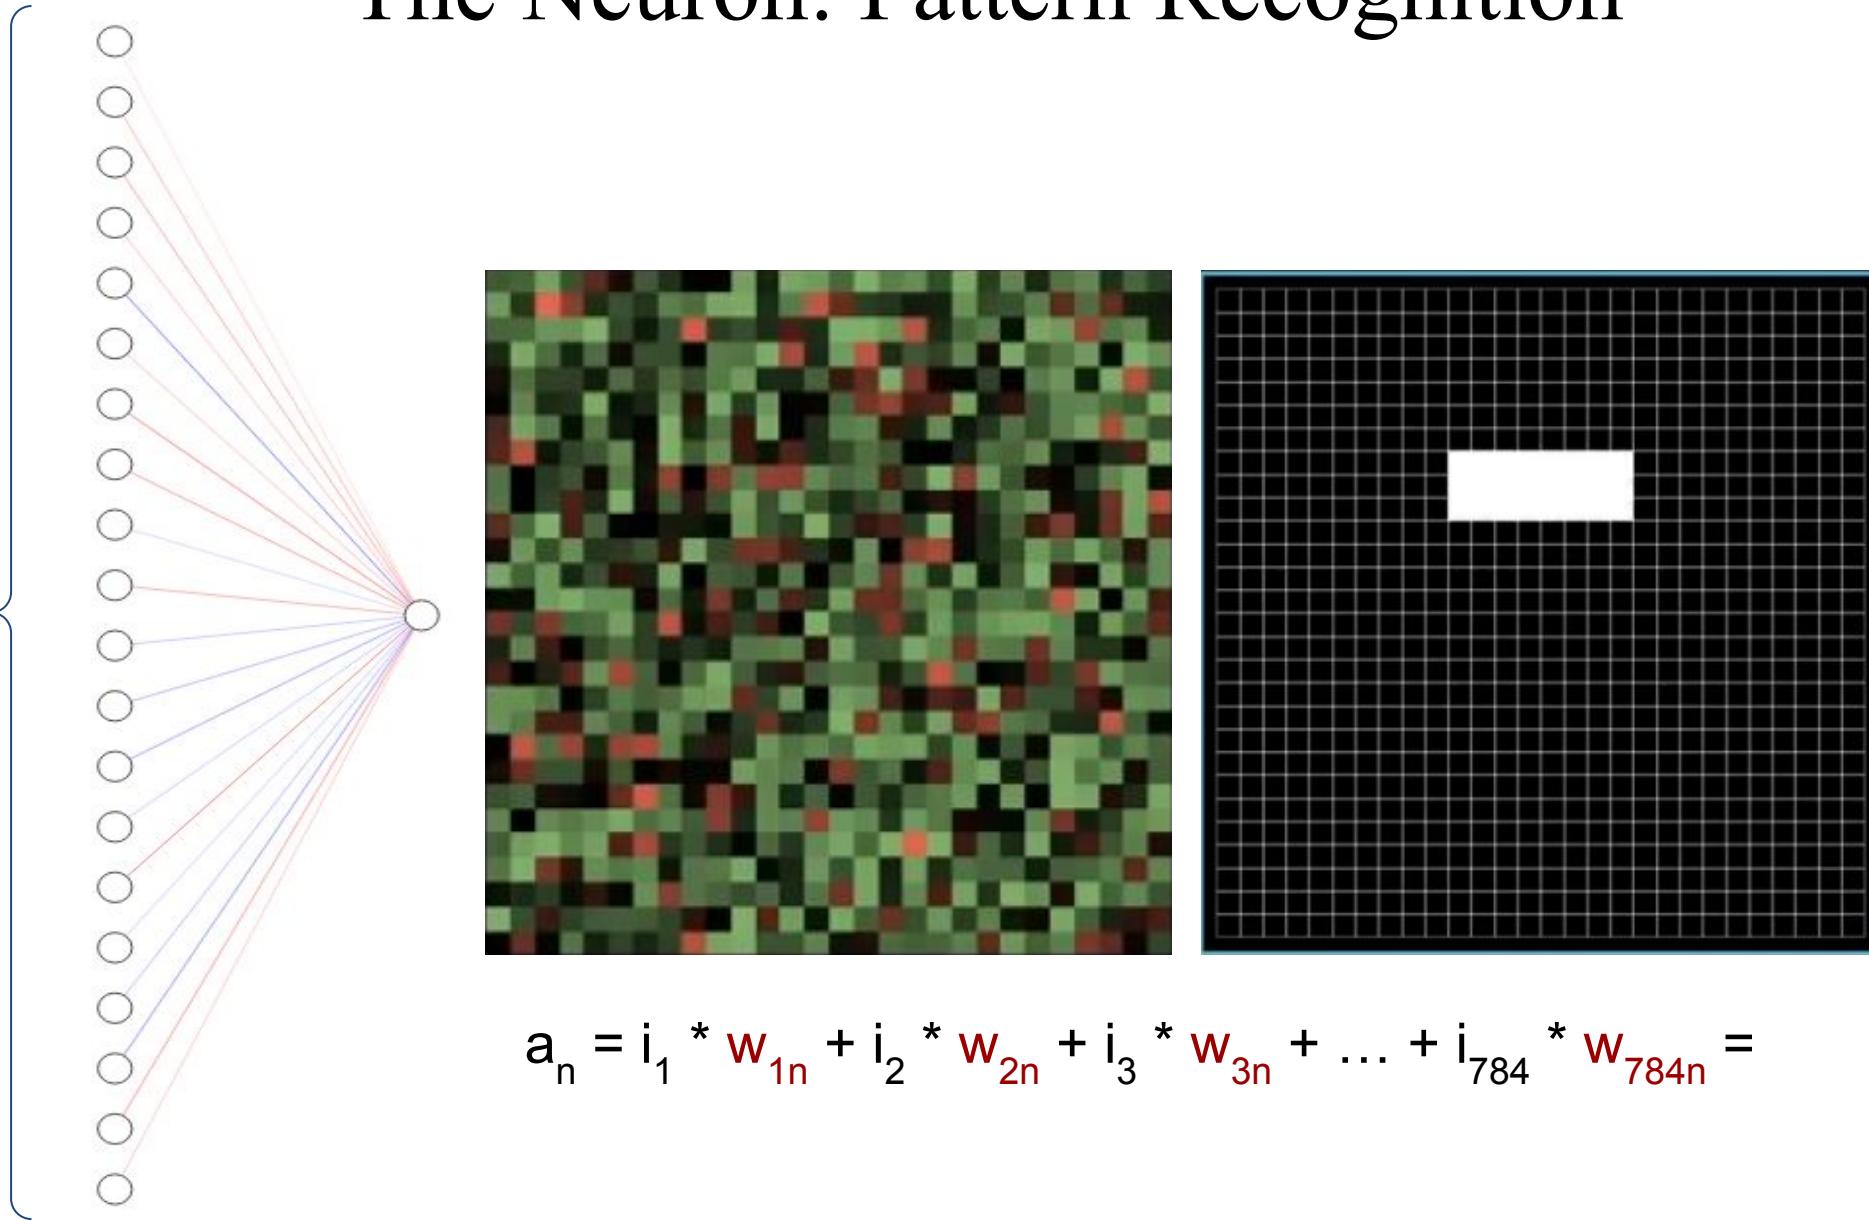
$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

How can a neuron detect that pattern?

- Every pixel is an input
- What's the output?

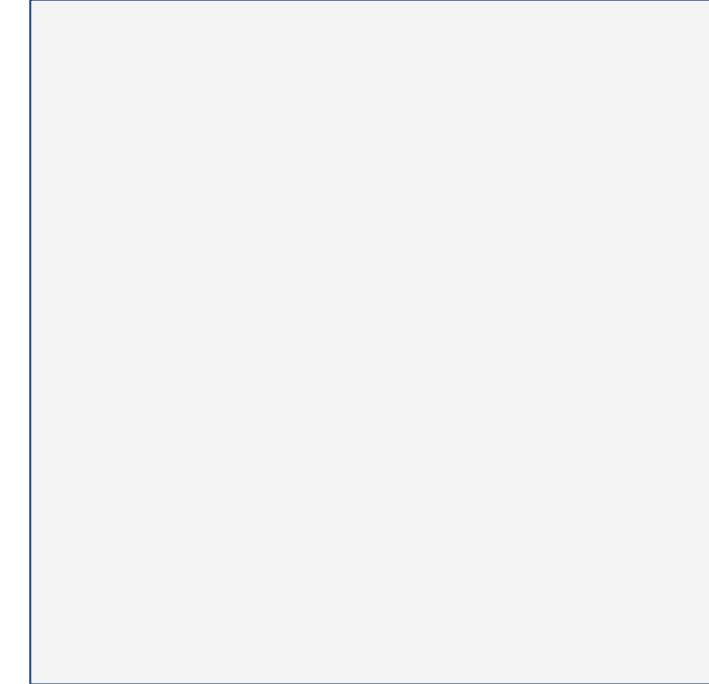
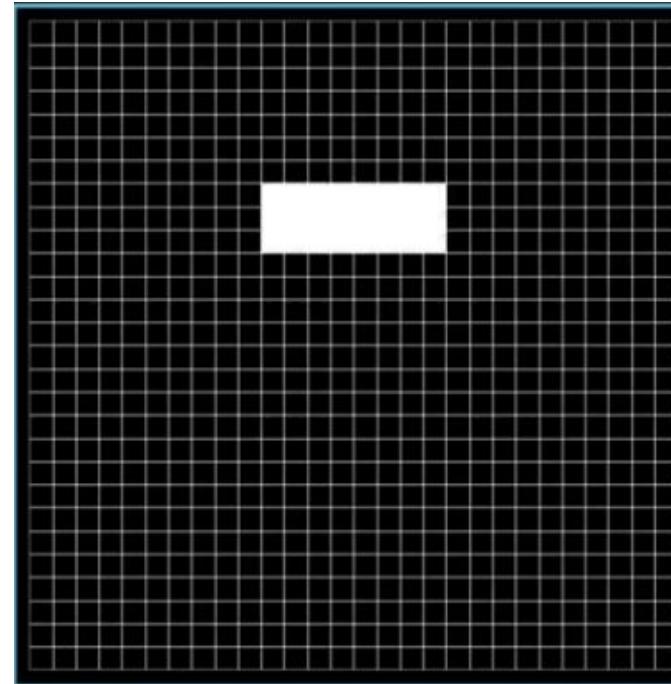
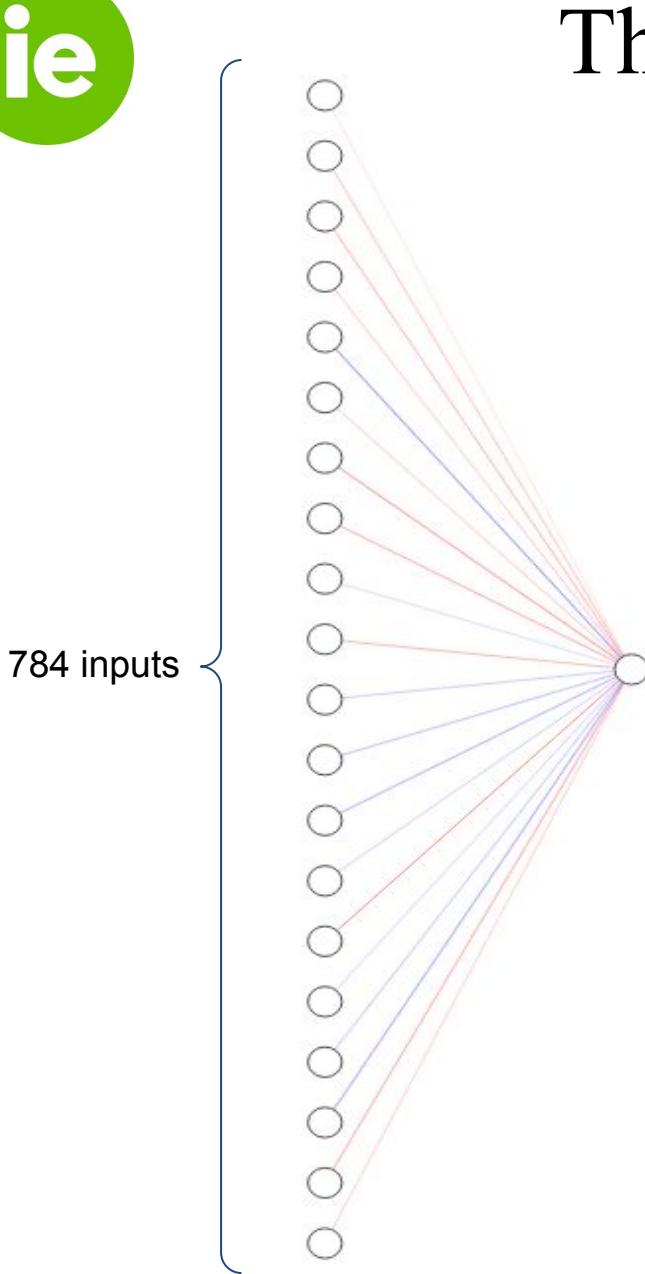
The Neuron: Pattern Recognition

784 inputs



The Neuron: Pattern Recognition

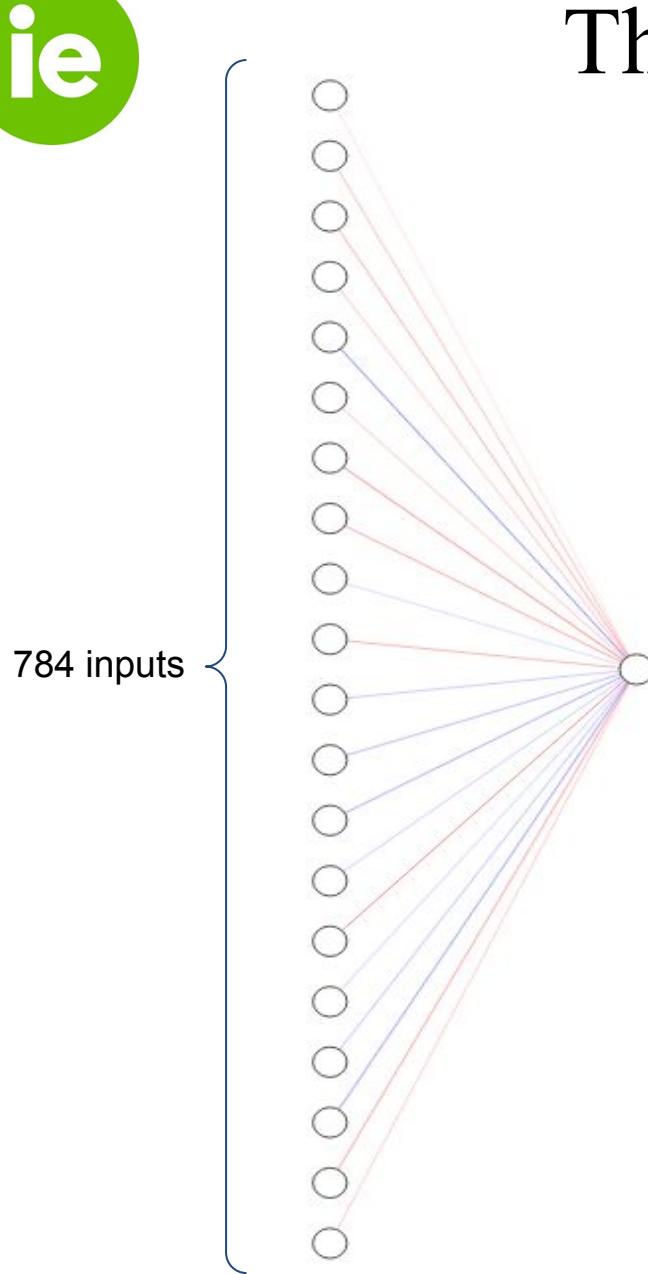
- Watch out! There is a mistake here!



$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

The Neuron: Pattern Recognition

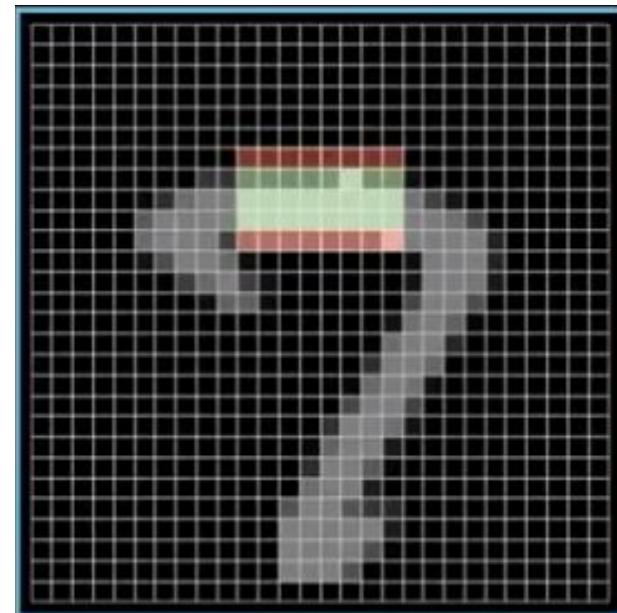
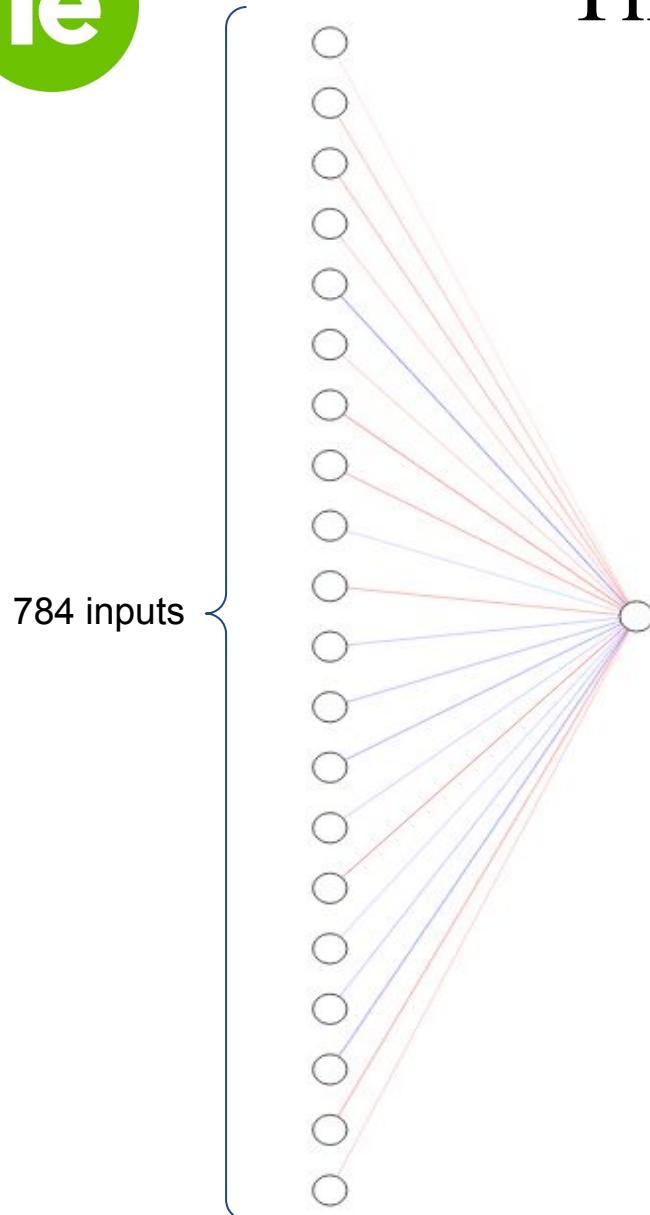
- Watch out! There is a mistake here!



$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

The Neuron: Pattern Recognition

- Watch out! There is a mistake here!

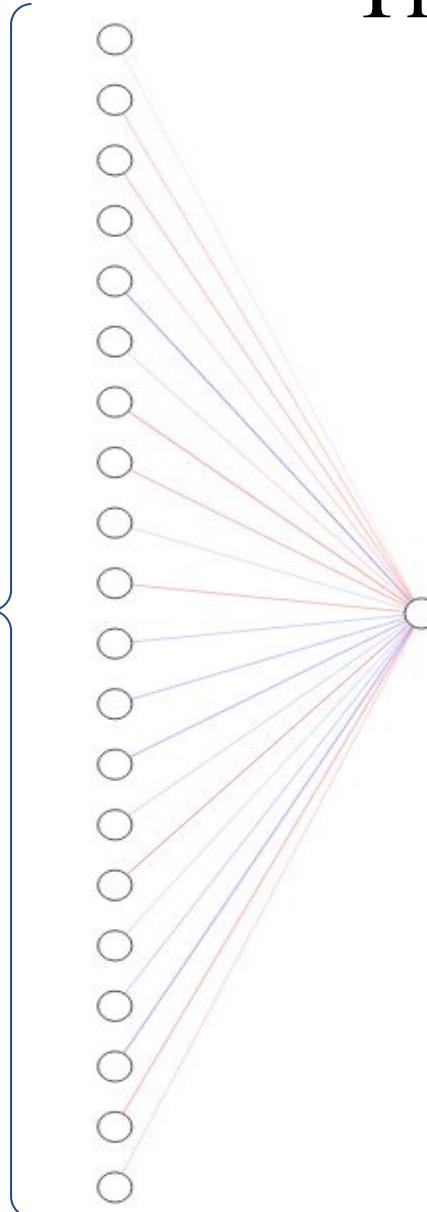


$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

The Neuron: Pattern Recognition

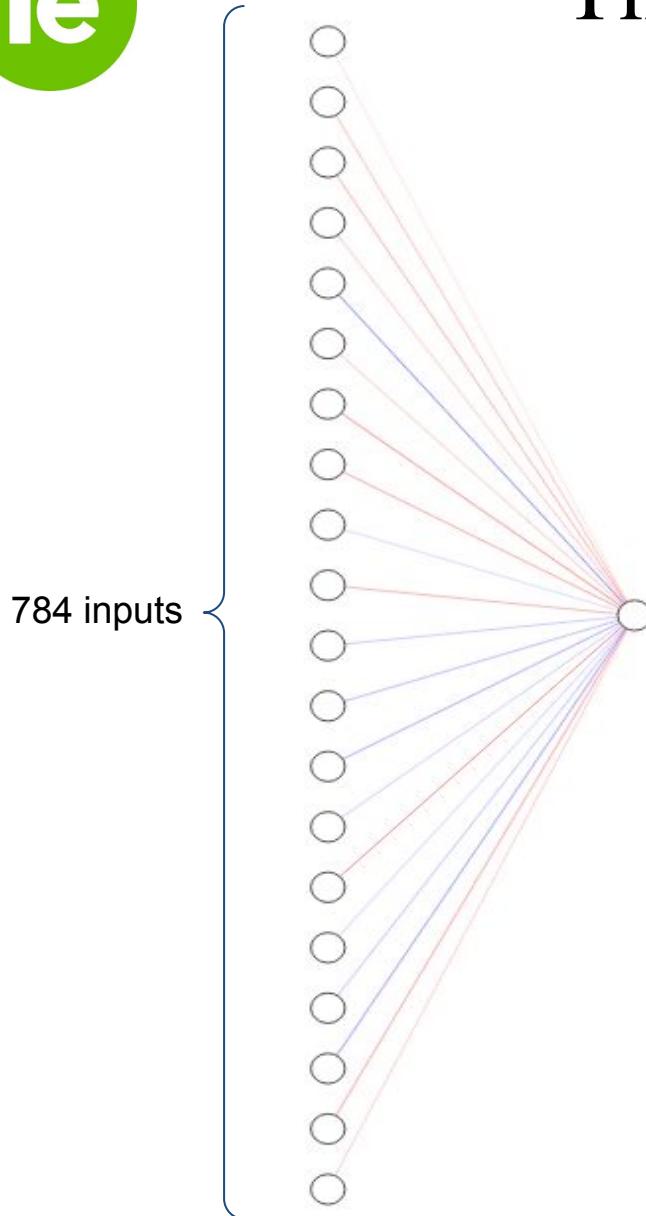
- What is our neuron output range?

784 inputs



$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

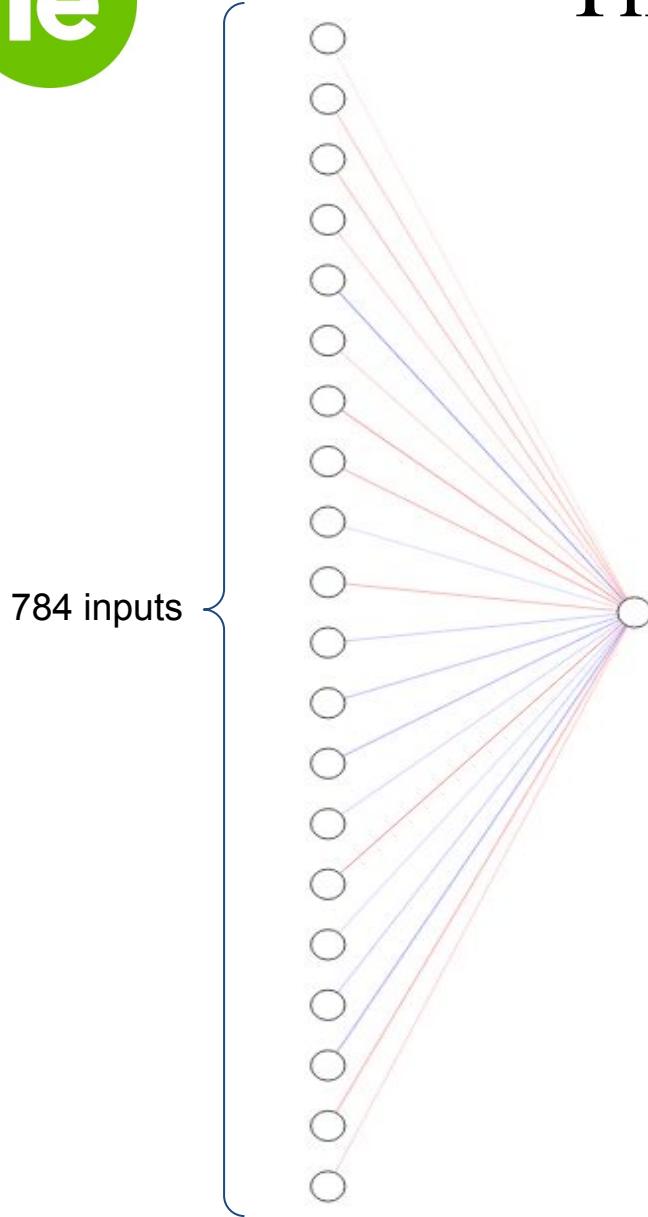
The Neuron: Pattern Recognition



- What is our neuron output range?
 - Real numbers

$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

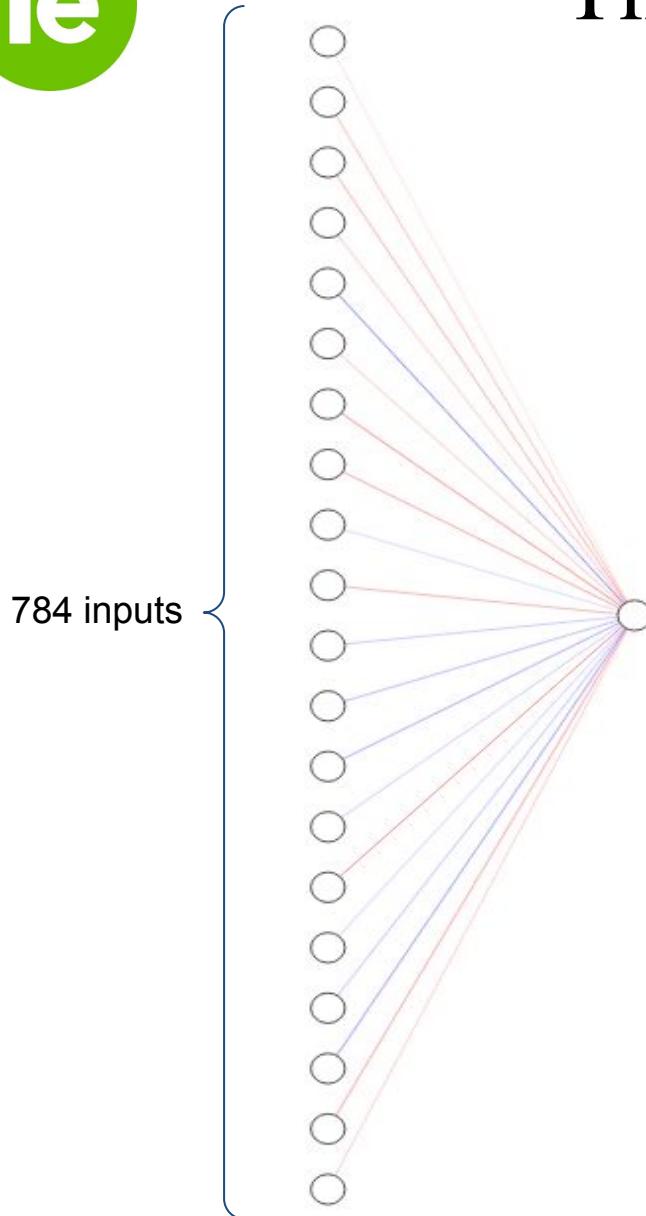
The Neuron: Pattern Recognition



- What is our neuron output range?
 - Real numbers
- What output range would we like?
 - Controlled range: (0,1) (-1,1).

$$a_n = i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} =$$

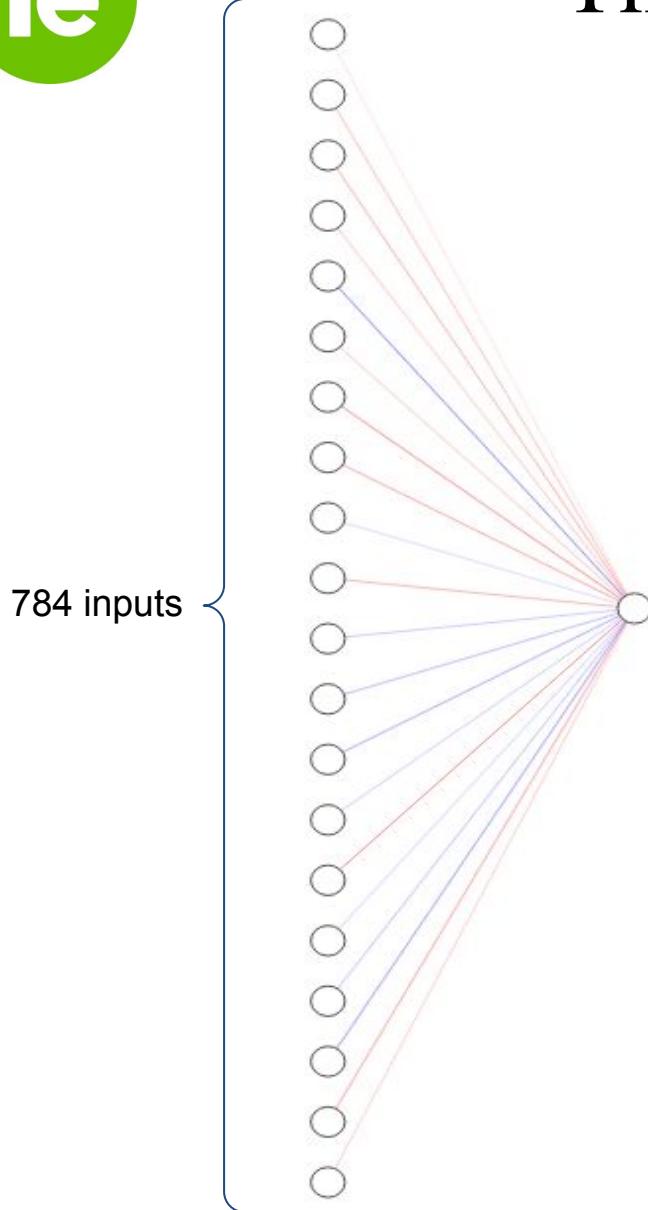
The Neuron: Pattern Recognition



- What is our neuron output range?
 - Real numbers
- What output range would we like?
 - Controlled range: (0,1) (-1,1).
 - Activation function (sigmoid, etc.)
- What if I want to be very sure that I have the exact pattern?

$$a_n = f(i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n}) =$$

The Neuron: Pattern Recognition



- What is our neuron output range?
 - Real numbers
- What output range would we like?
 - Controlled range: (0,1) (-1,1).
 - Activation function (sigmoid, etc.)
- What if I want to be very sure that I have the exact pattern?
 - Bias!

$$a_n = f(i_1 * w_{1n} + i_2 * w_{2n} + i_3 * w_{3n} + \dots + i_{784} * w_{784n} + b_n)$$

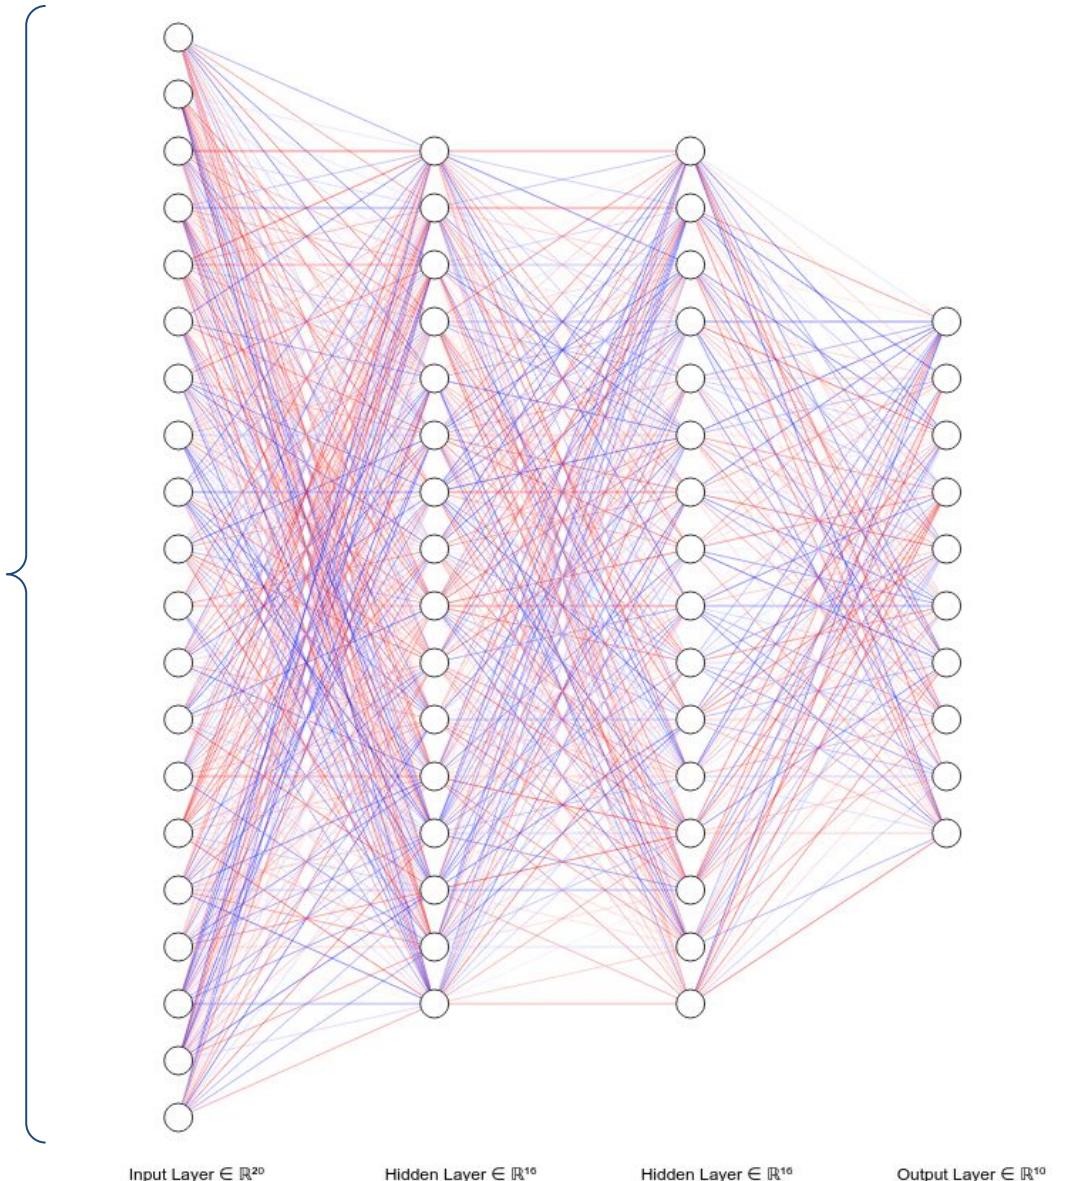
Outline



1. Recap
2. Deep L-Layers
3. Divide and Conquer
4. The Neuron: Pattern Recognition
5. DNNs



784 pixels



DNNs

Complete Network:

Weights:

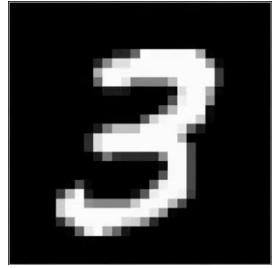
$$784 * 16 + 16 * 16 + 16 * 10$$

Bias:

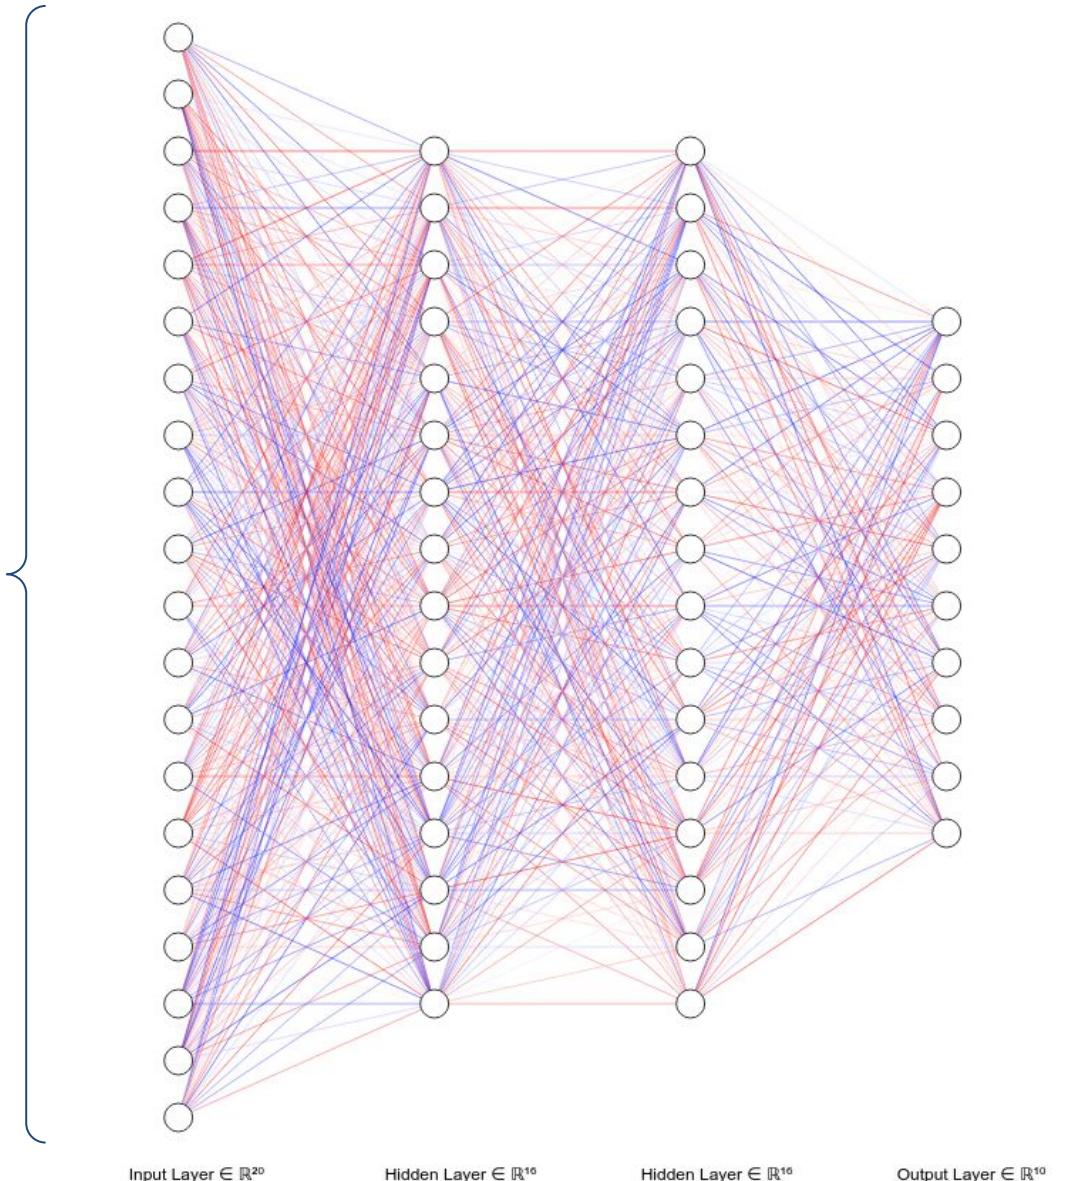
$$16 + 16 + 10$$

Trainable Params:

$$13002$$



784 pixels



DNNs

Complete Network:

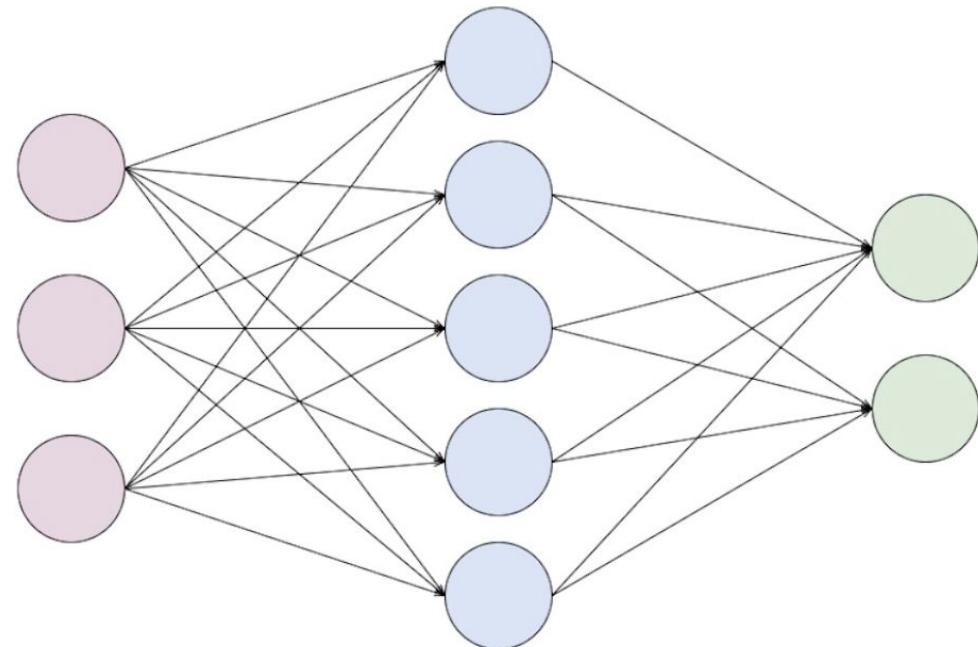
Trainable Parameters:
13002

Training consist of finding
the value for each trainable
parameters that minimize
the cost function on the
training dataset

DNNs



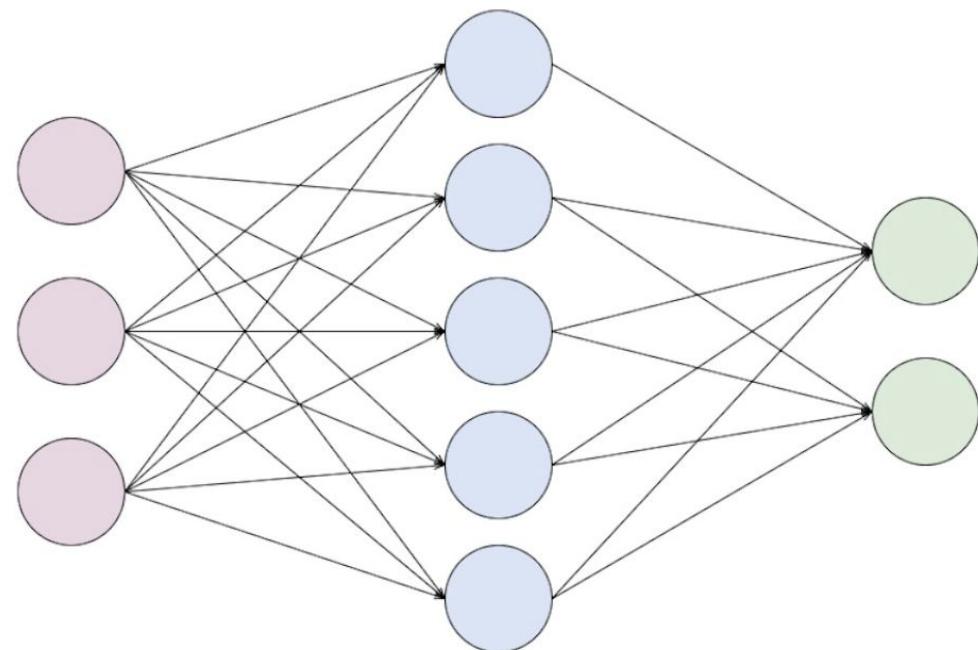
A net of information



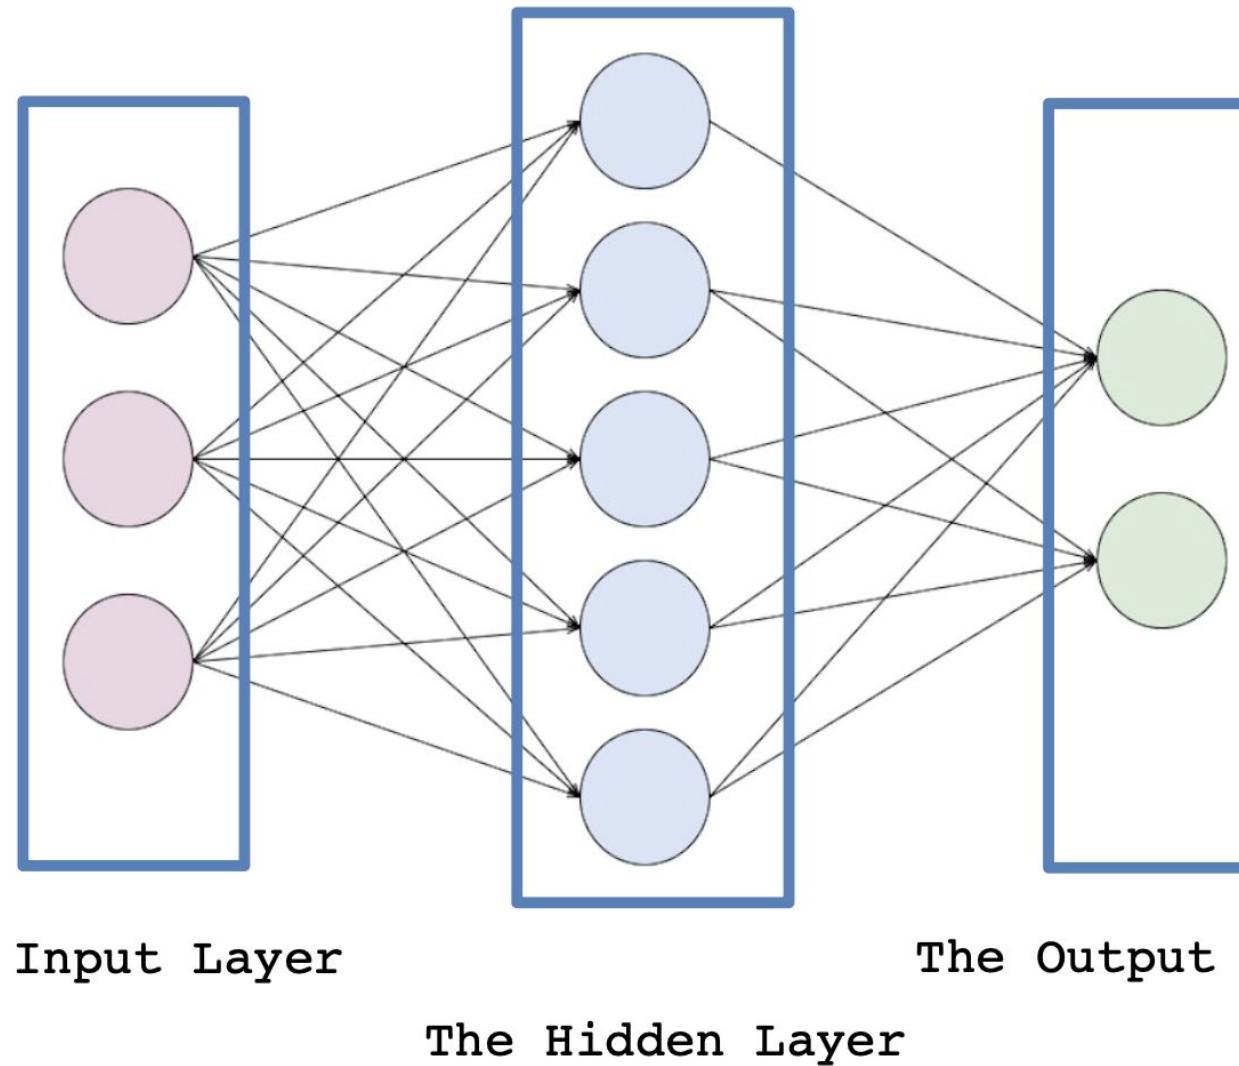
Disclaimer:

Several names:

- Dense
- Feed-forward
- NNs
- MLP
- ANNs
- DNNs
- ...

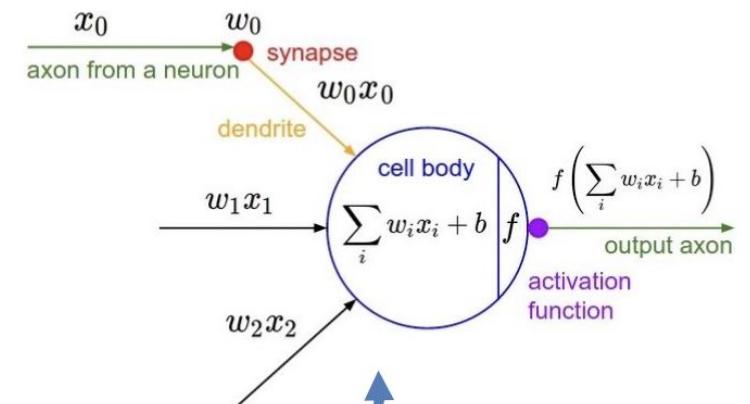
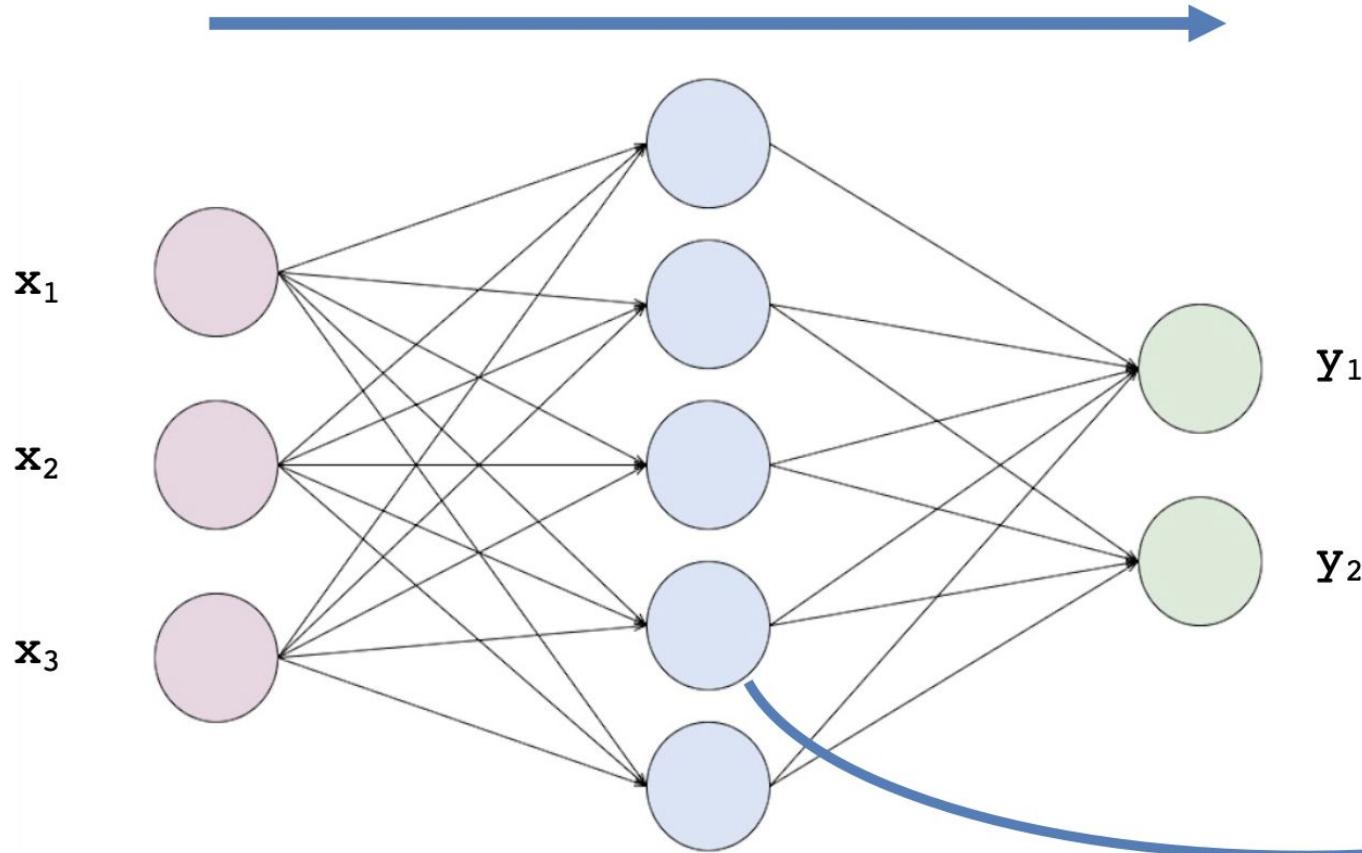


DNNs



DNNs

$$Y = F(x)$$



ANNs. Keras



```
my_first_neural_net = tf.keras.Sequential([
    layers.Input(shape=(3, ), name="input"),
    layers.Dense(units=5, activation="sigmoid", name="hidden"),
    layers.Dense(units=2, activation="softmax", name="output")
])
```

ANNs. Keras



```
my_first_neural_net.summary()
```

```
↳ Model: "sequential"
```

Layer (type)	Output Shape	Param #
<hr/>		
hidden (Dense)	(None, 5)	20
output (Dense)	(None, 2)	12

```
Total params: 32
```

```
Trainable params: 32
```

```
Non-trainable params: 0
```



ANNs. Activation Functions



- [Star 56,732](#)
- About Keras
- Getting started
- Developer guides
- Keras API reference
- Models API
- Layers API
- The base Layer class
- Layer activations
- Layer weight initializers
- Layer weight regularizers
- Layer weight constraints
- Core layers
- Convolution layers
- Pooling layers
- Recurrent layers

Search Keras documentation...

» Keras API reference / Layers API / Layer activation functions

Layer activation functions

Usage of activations

Activations can either be used through an `Activation` layer, or through the `activation` argument supported by all forward layers:

```
model.add(layers.Dense(64, activation=activations.relu))
```

This is equivalent to:

```
from tensorflow.keras import layers
from tensorflow.keras import activations

model.add(layers.Dense(64))
model.add(layers.Activation(activations.relu))
```

All built-in activations may also be passed via their string identifier:

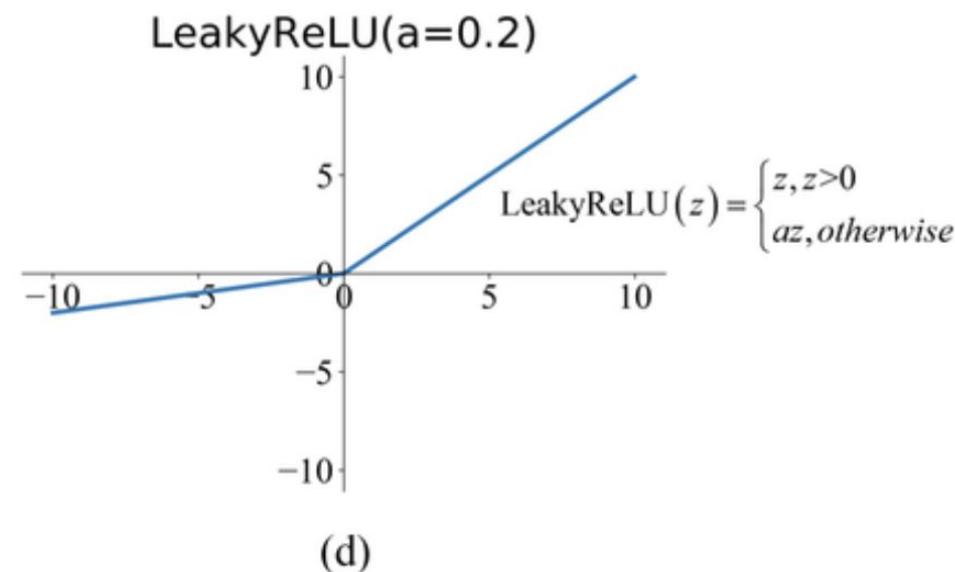
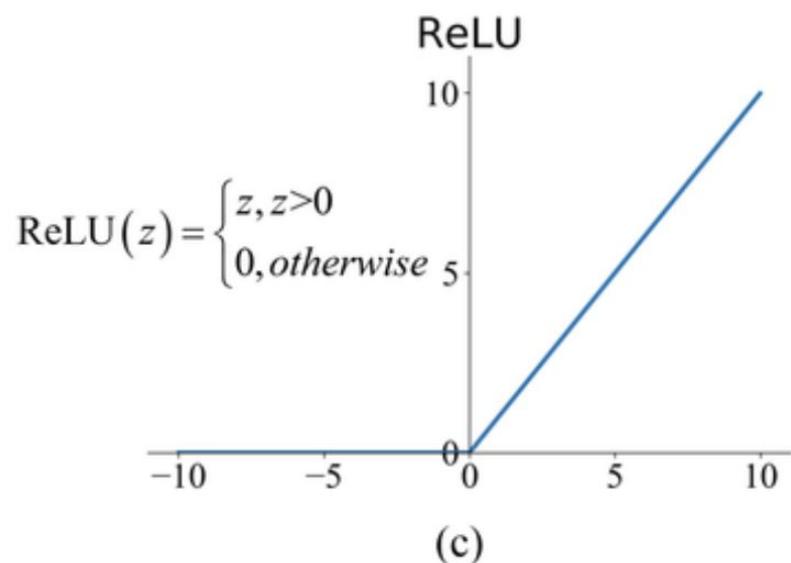
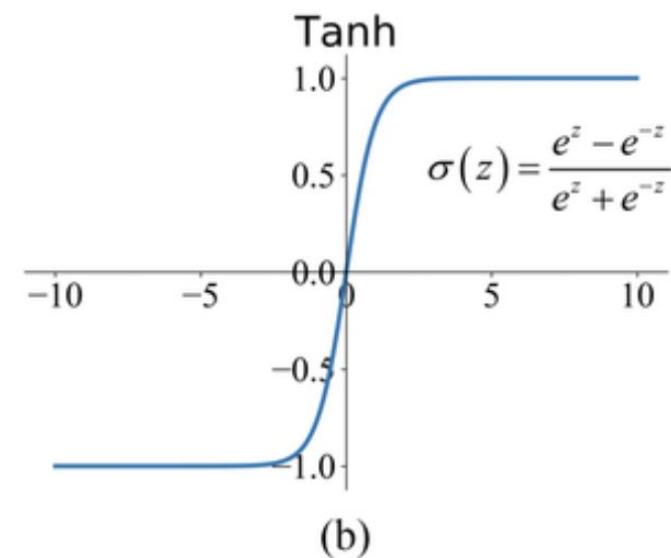
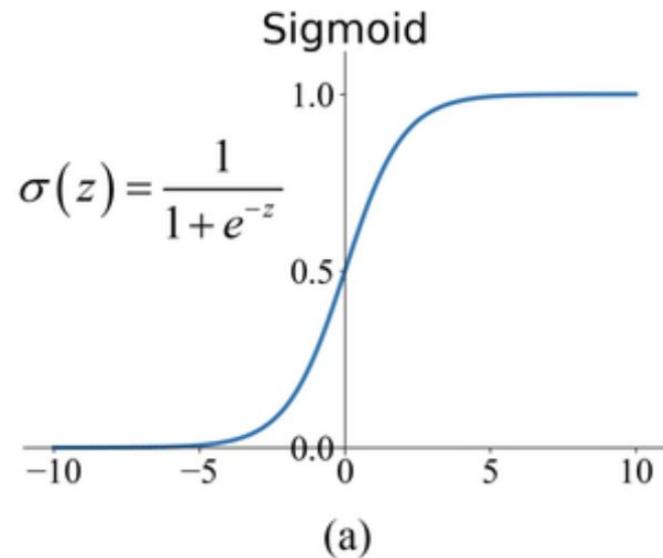
```
model.add(layers.Dense(64, activation='relu'))
```

- Layer activation functions**
 - ◆ Usage of activations
 - ◆ Available activations
 - relu function
 - sigmoid function
 - softmax function
 - softplus function
 - softsign function
 - tanh function
 - selu function
 - elu function
 - exponential function
 - ◆ Creating custom activations
 - ◆ About "advanced activation" layers

ANNs. Activation Functions

Name	Plot	Equation
Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) ^[2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) ^[3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$

ANNs. Activation Functions



ANNs. Activation Functions

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
Arctan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) ^[2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) ^[3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

ANNs. Activation Functions

- The combination of linear functions is just another linear function. No matter the number of hidden layers, if we use linear functions as activations, we could just deal with outputs that are a linear combination of our inputs... not interesting.
- Non-linear activations functions are key in neural nets to represent non-linear mapping functions.
- **In the output layer**, at general: use sigmoid for binary classification, softmax for multiclass classification and linear for regression.
 - You can use softmax for binary classification using two outputs and in the case of a problem where data that can belong to multiple classes, then use sigmoid. [i.e., a picture of Audi belongs to the classes' car and vehicle]
- **In hidden layers:**
 - Tanh used to work better than sigmoid as it centers data.
 - ReLu is faster and it works very well too. It is the most used.